

Convolutional Neural Networks: architectures, interpretation

LSSDS 2024

Mauricio Cerda, Eng. PhD

Programa de Biología Integrativa
I.C.B.M., Facultad de Medicina, Universidad de Chile

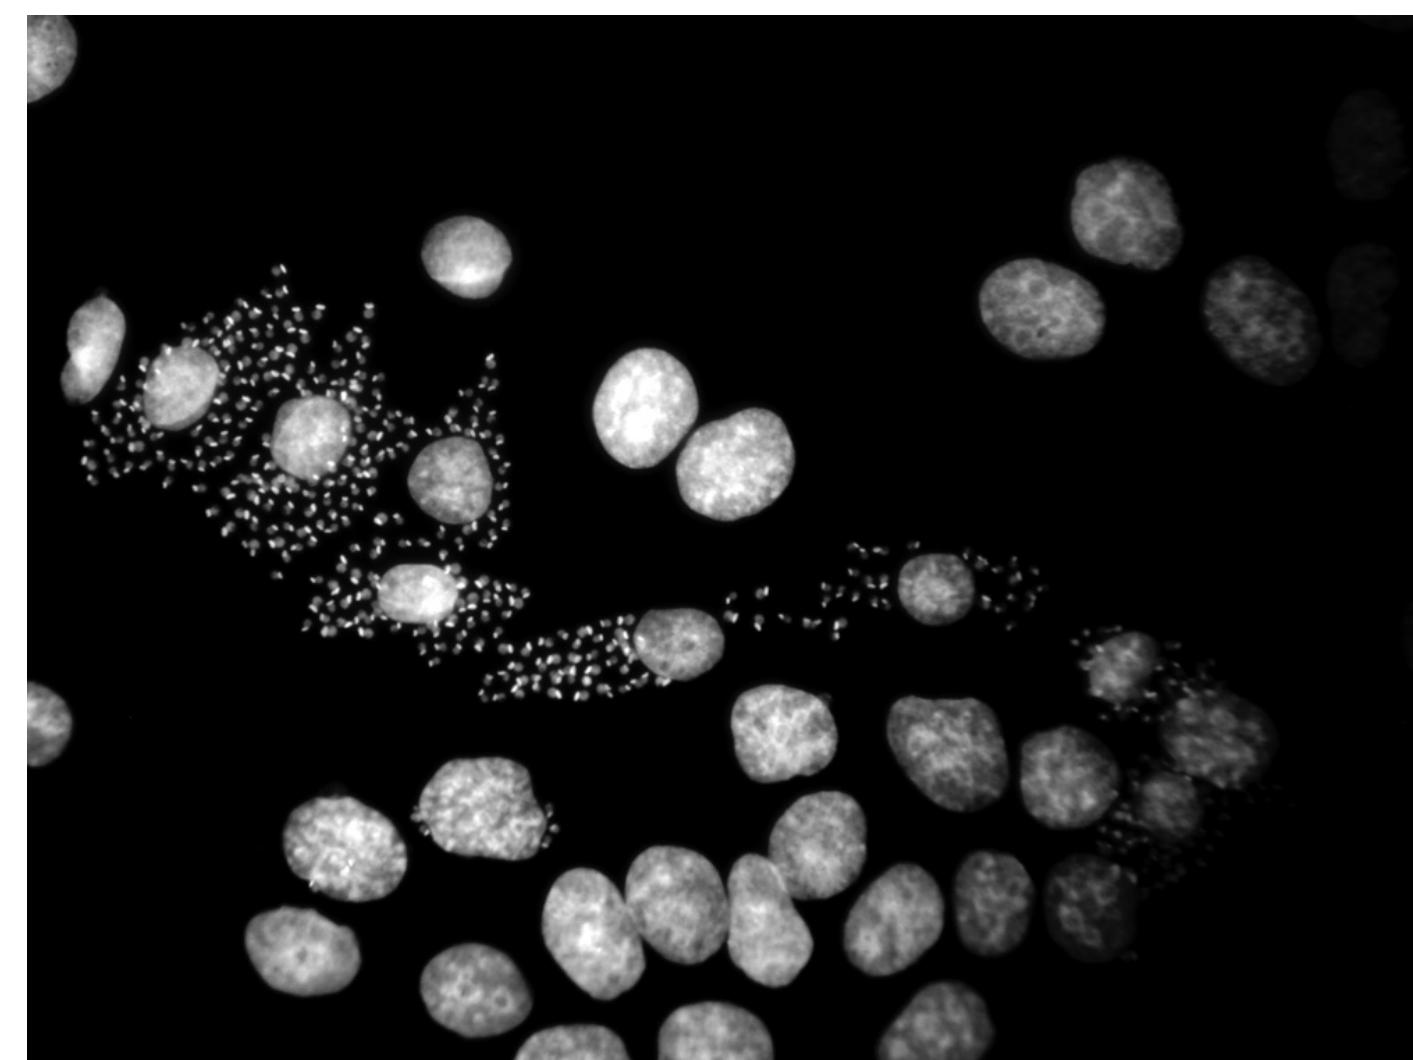


Outline



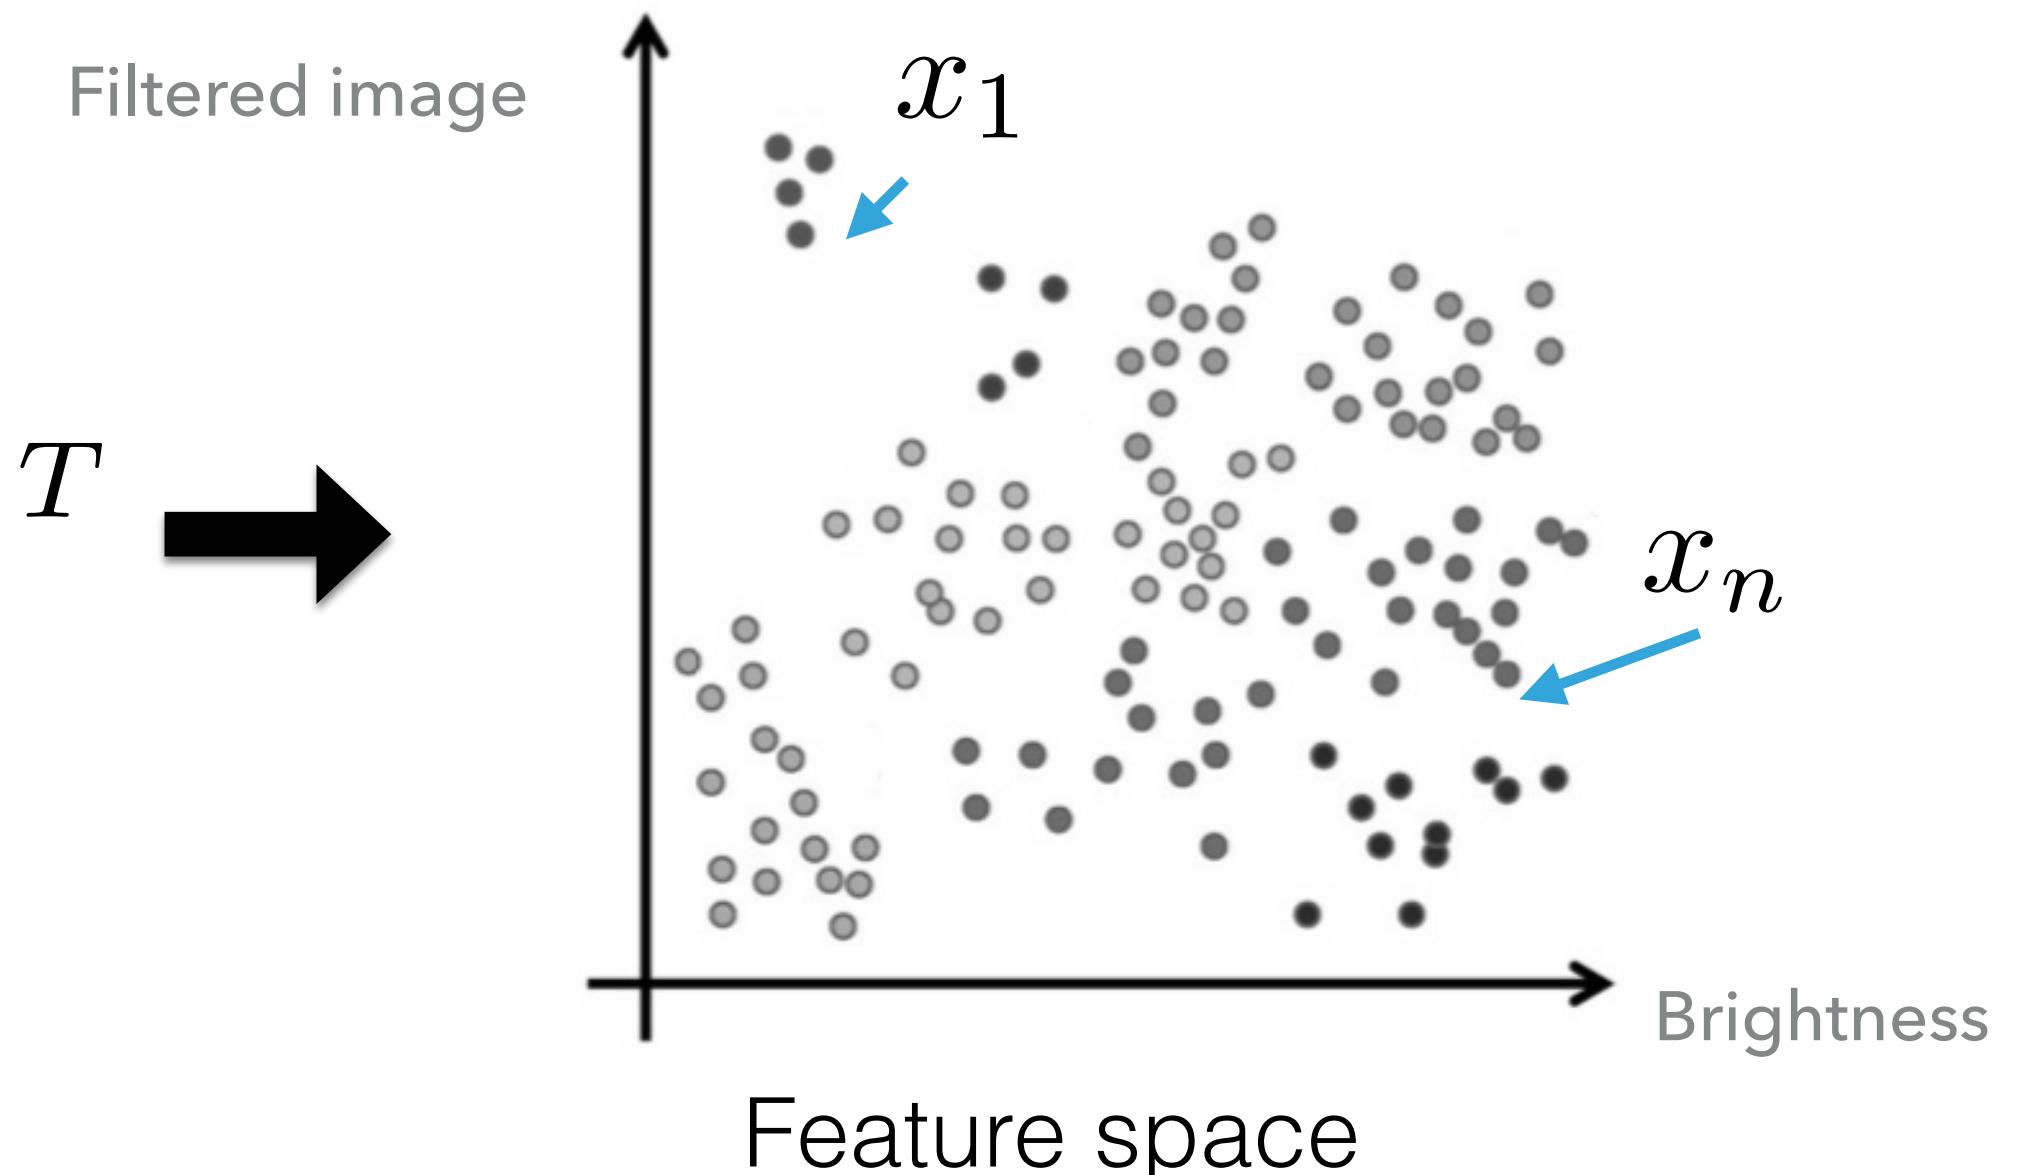
- How to select the right image features to classify/find clusters?
- Convolutional Neural Networks
- CNN architectures (some)
- CNN visualization

How to select the right features?

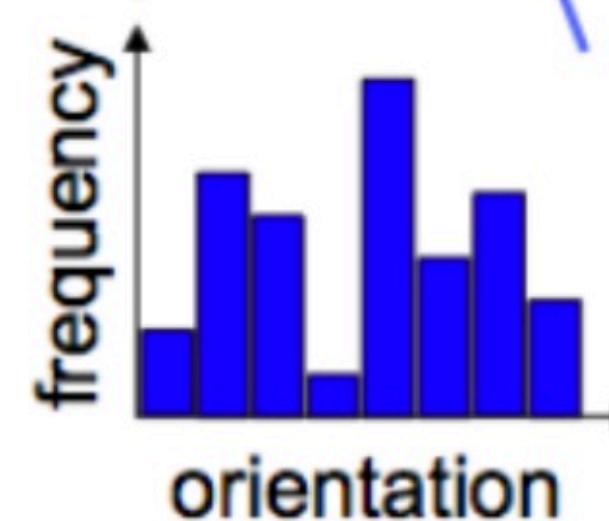
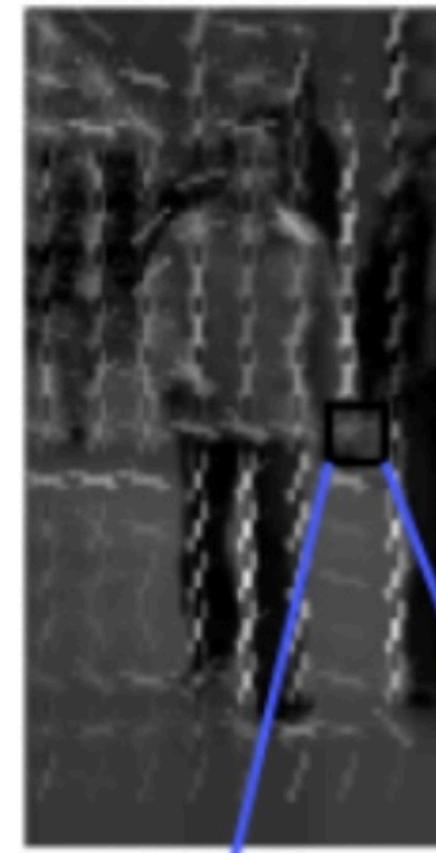


Images

→ $[x_1, x_2, \dots x_d]^T$ →
Features



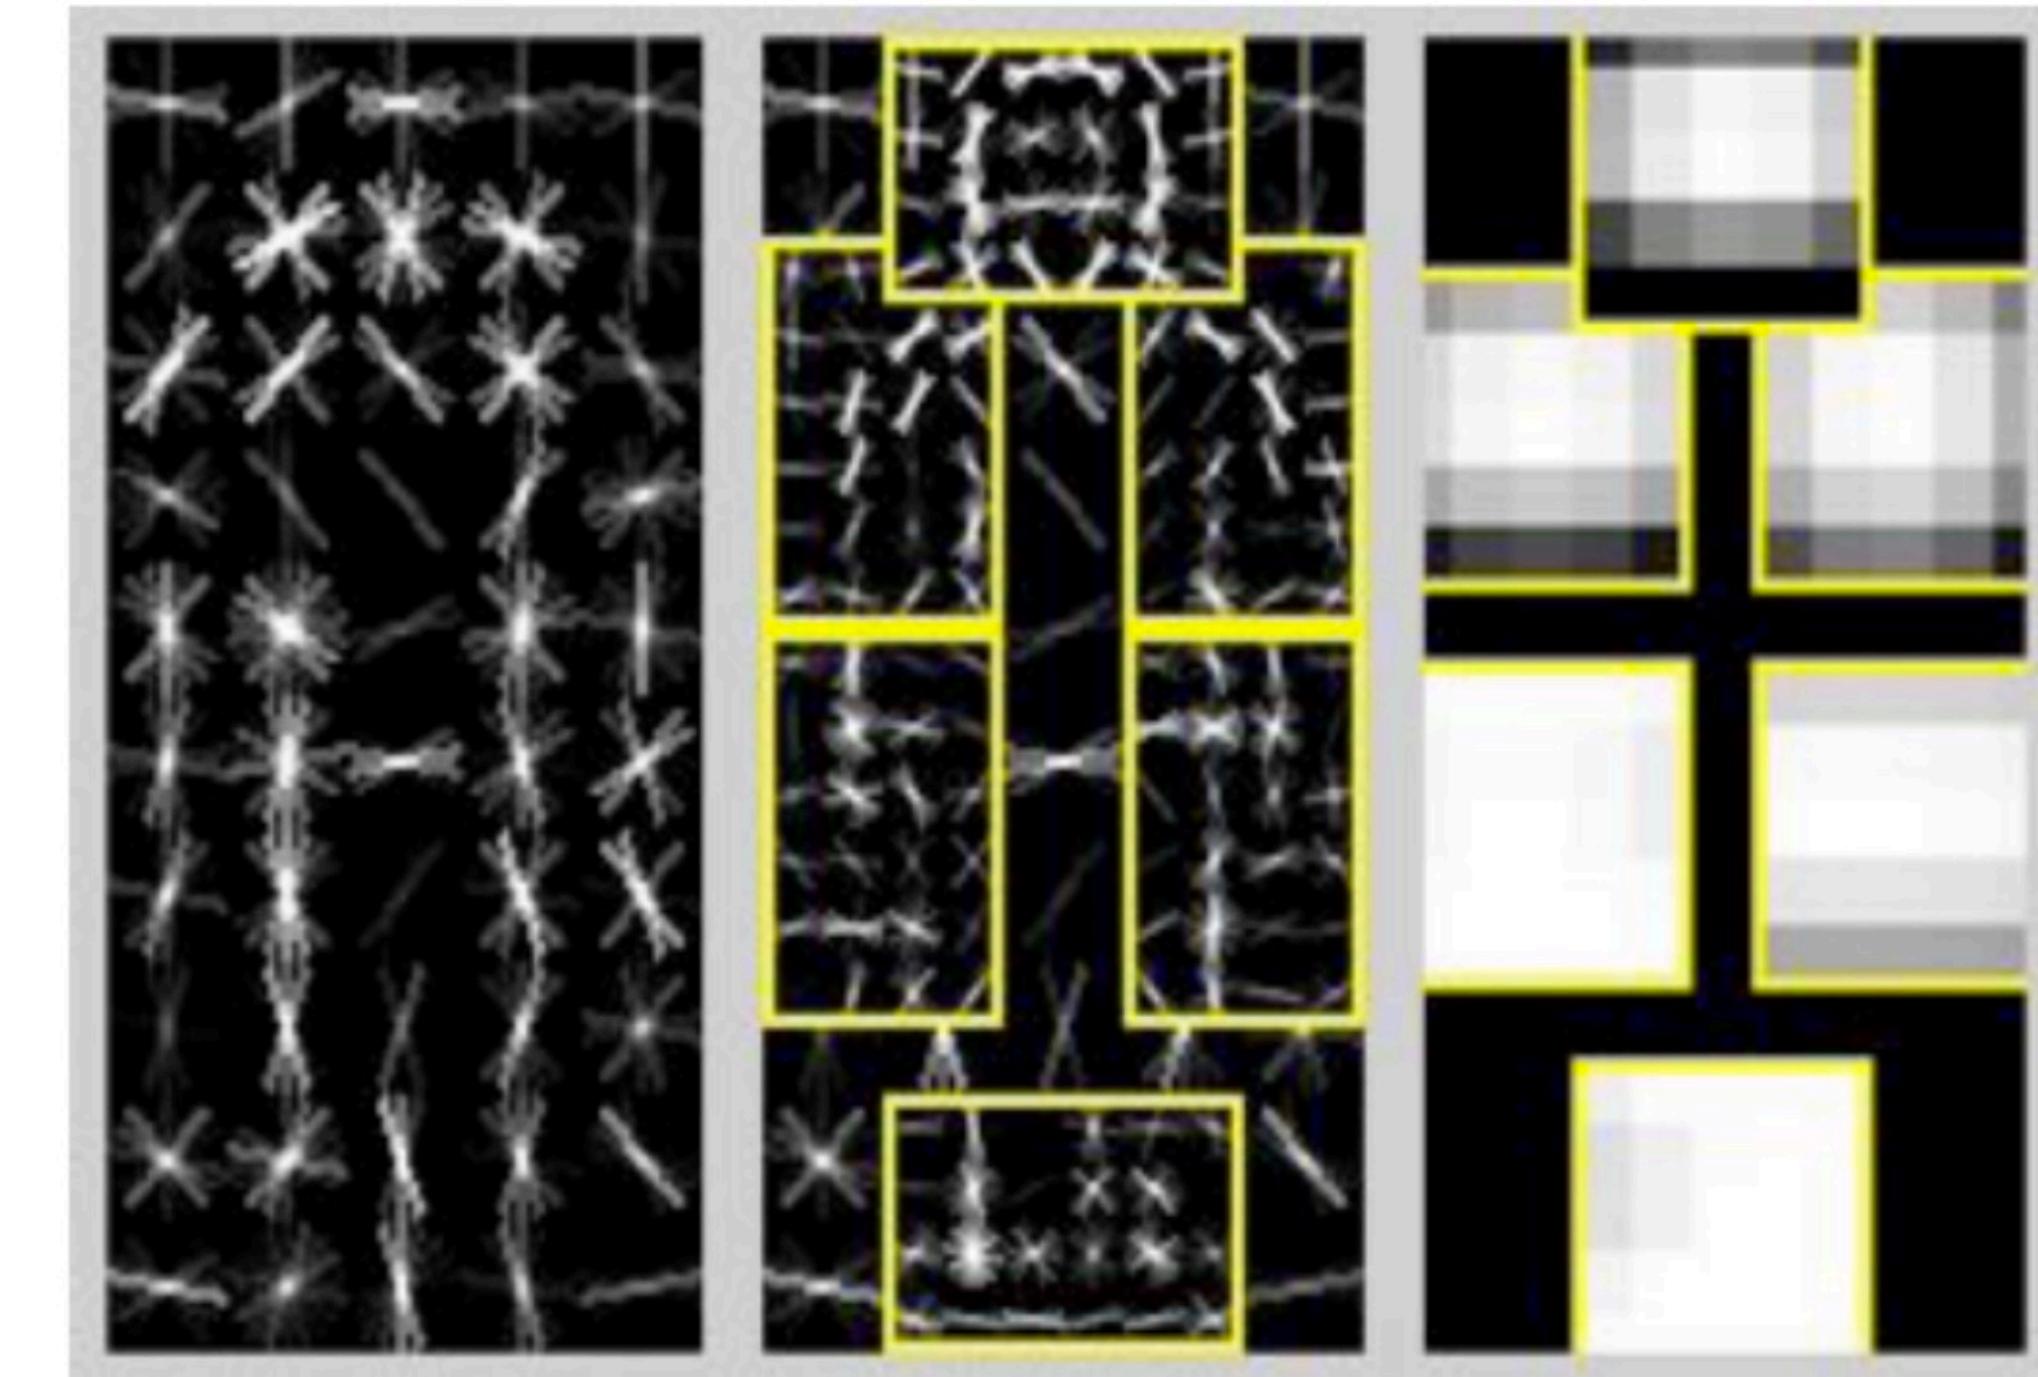
What about image filters?



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$mag(\nabla f) = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Histogram of Gradients (HoG)
Dalal & Triggs, 2005



Deformable Part Model
Felzenswalb, McAllester, Ramanan, 2009

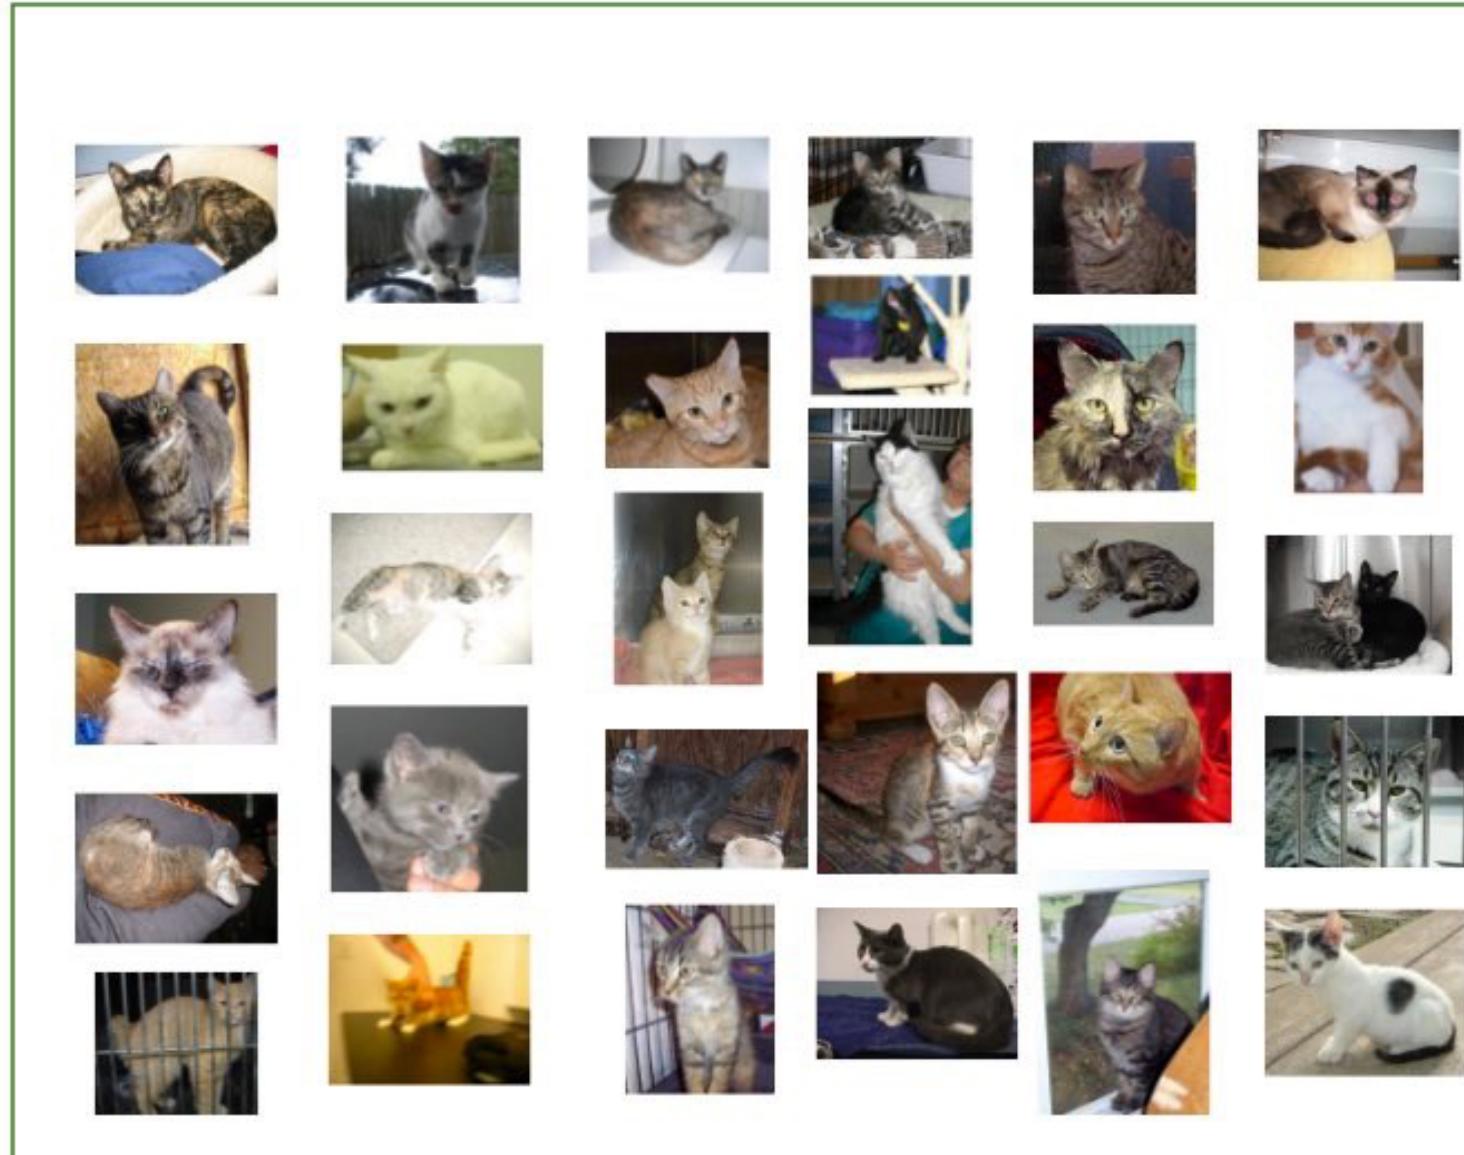
What about image filters?



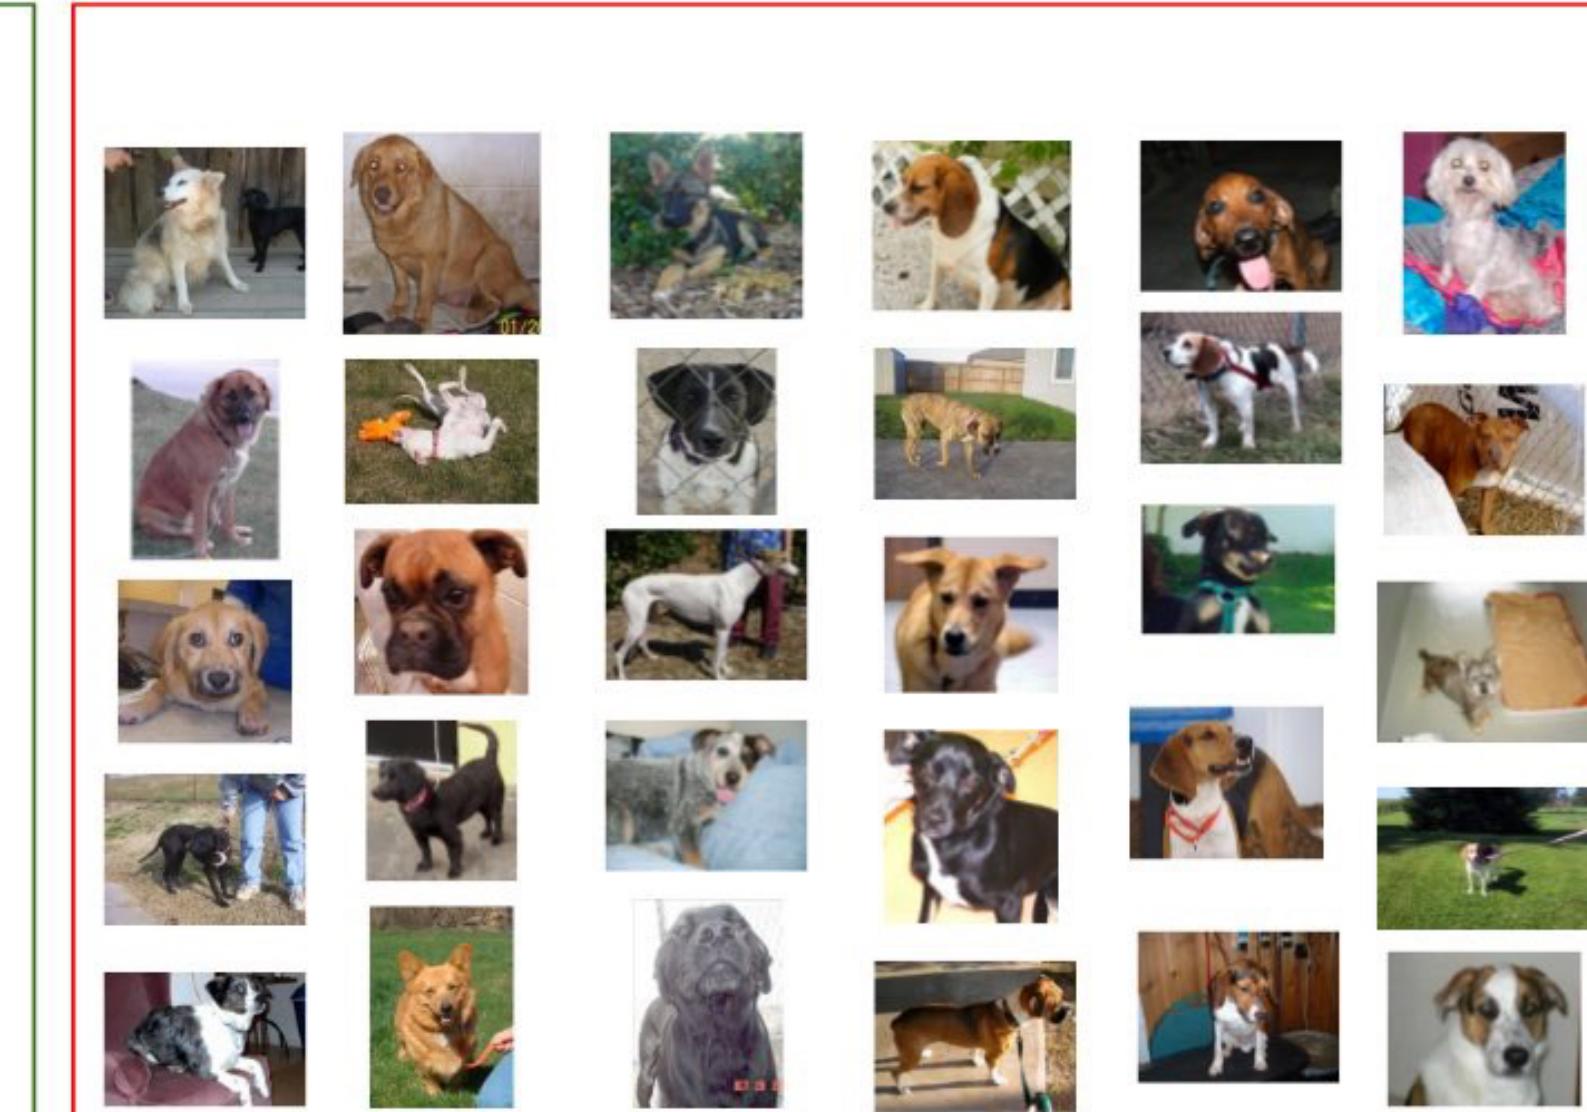
- Image features required **manual** feature selection.
- It was all about features!
- Many unsolved difficult problems.

Deep Learning key ideas

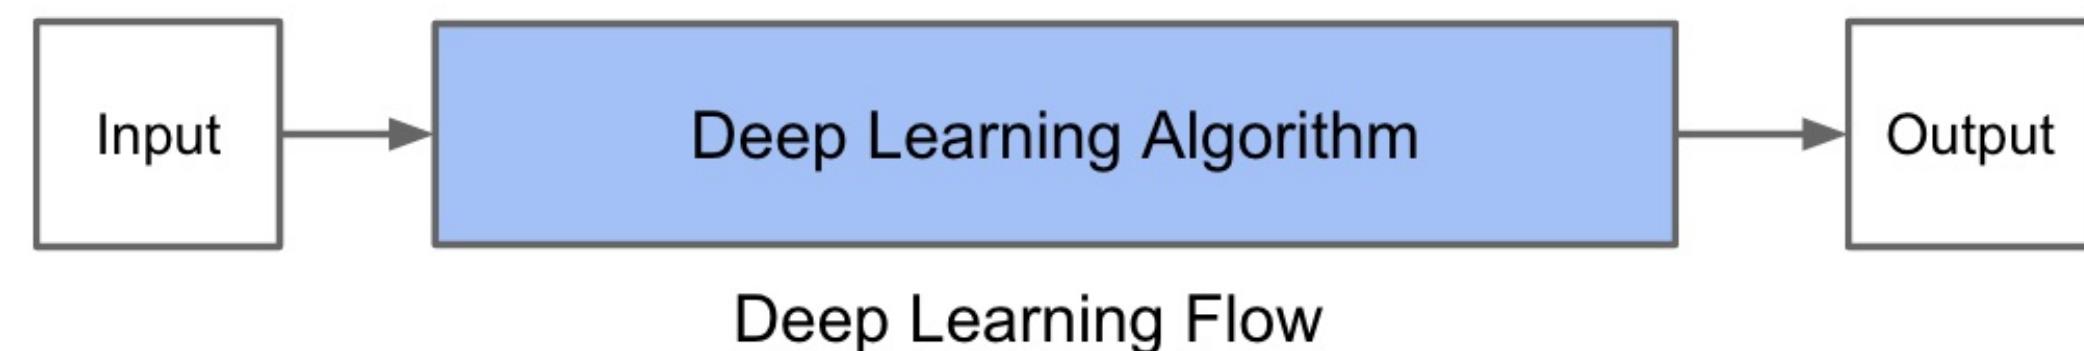
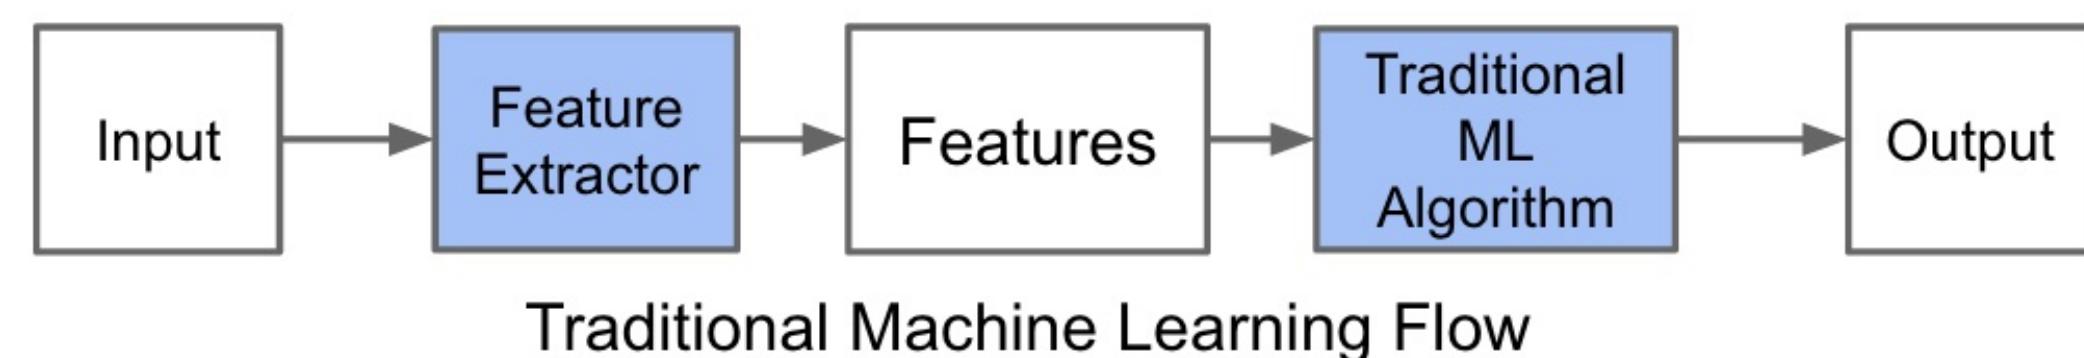
Cats



Dogs

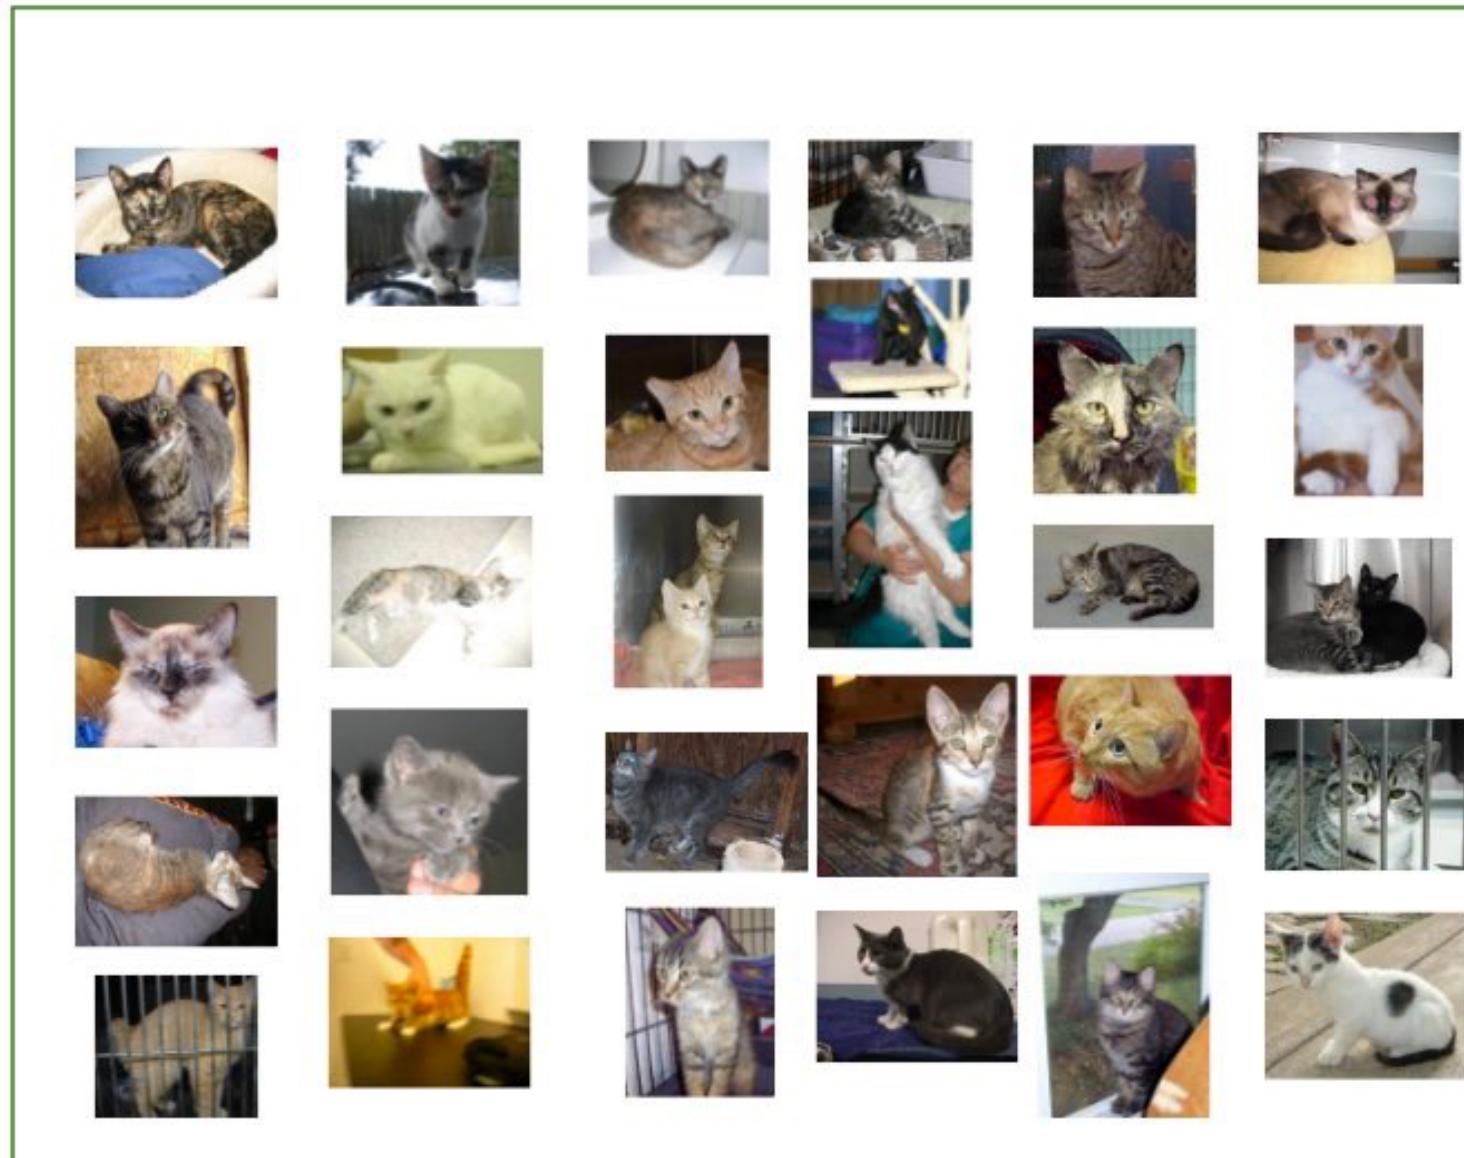


Sample of cats & dogs images from Kaggle Dataset

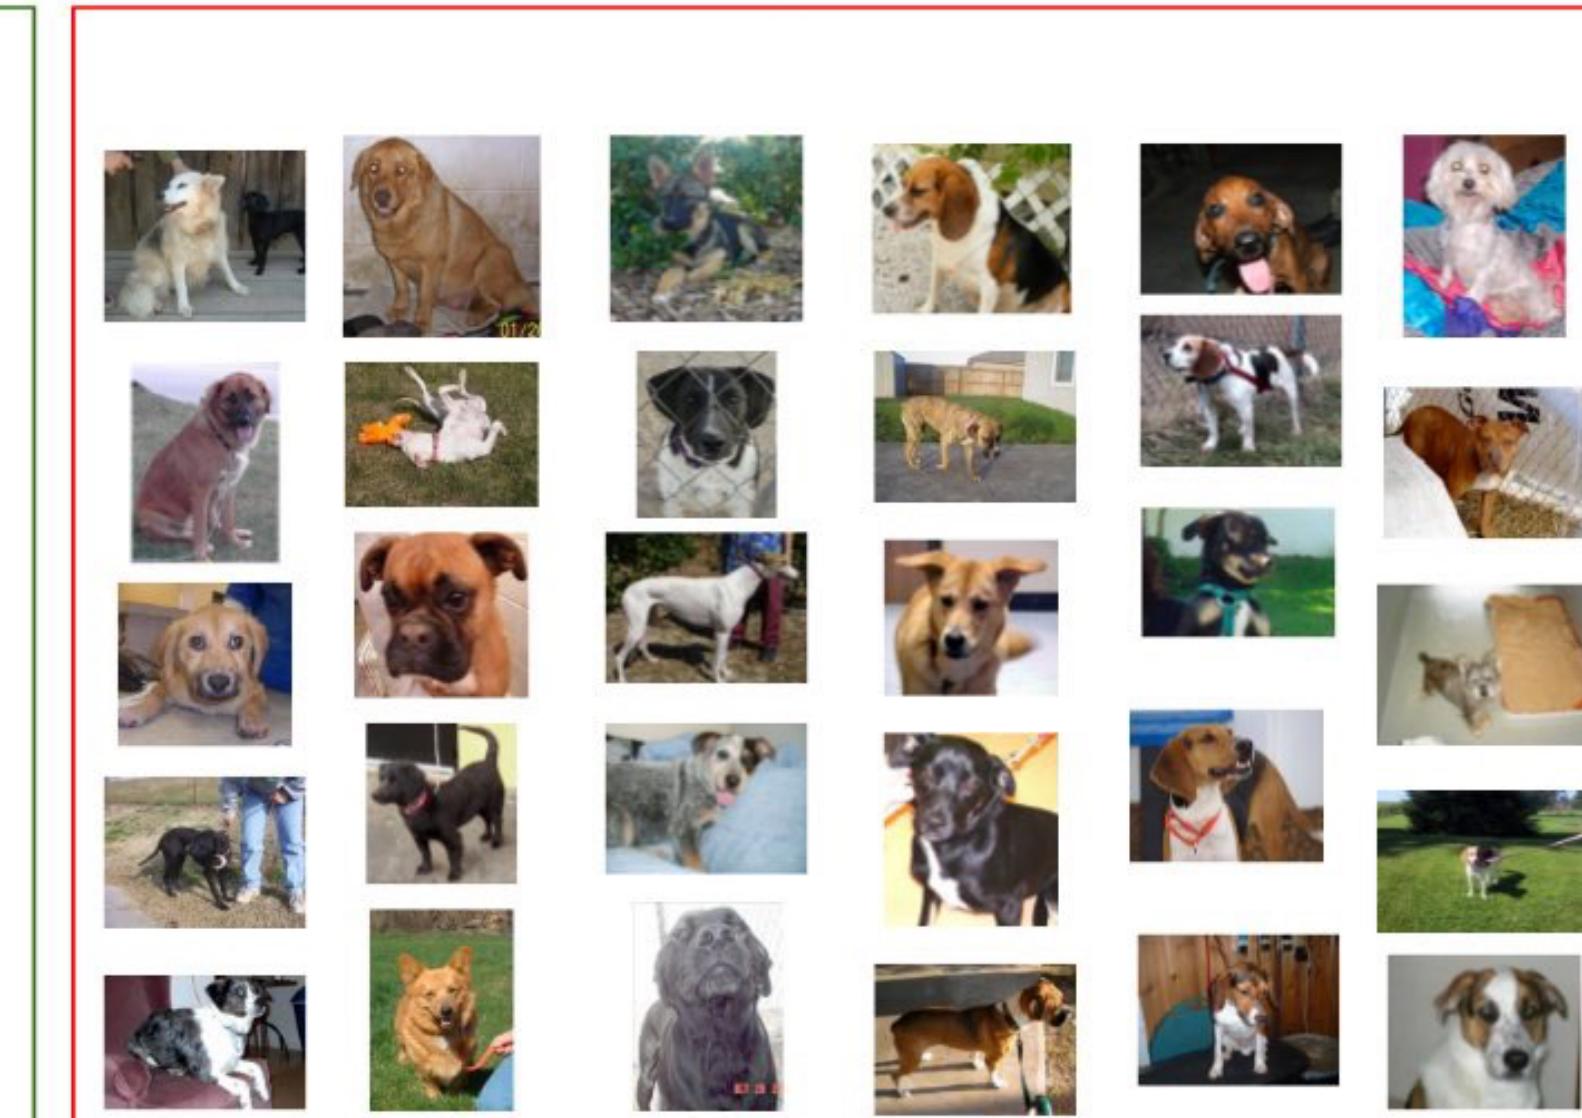


Deep Learning key ideas

Cats

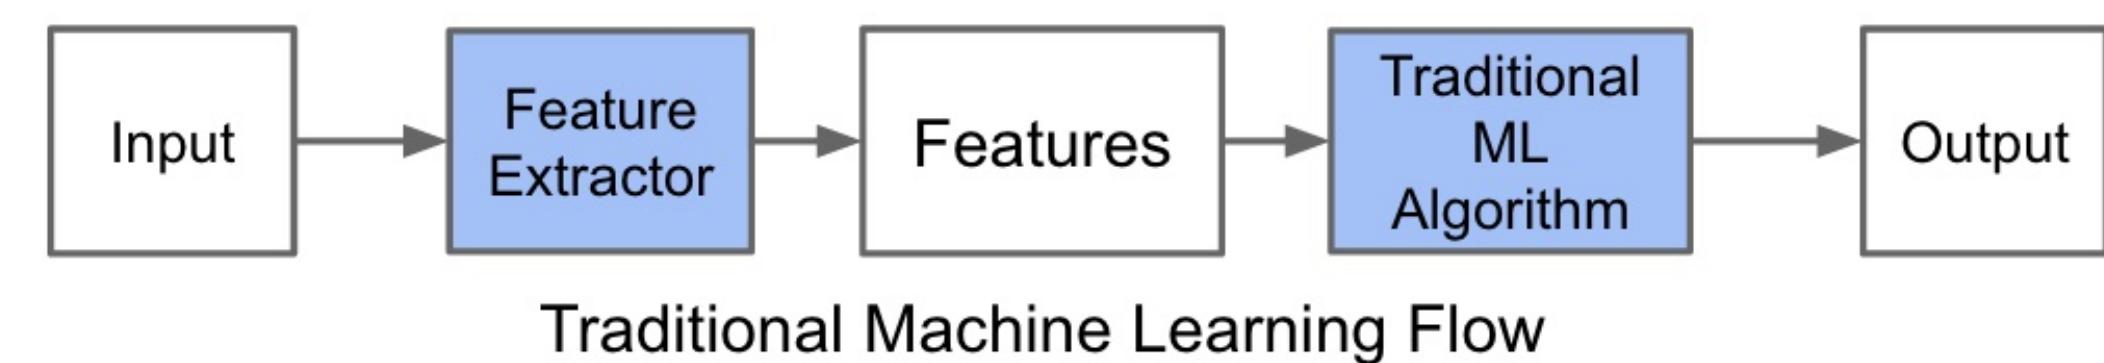


Dogs

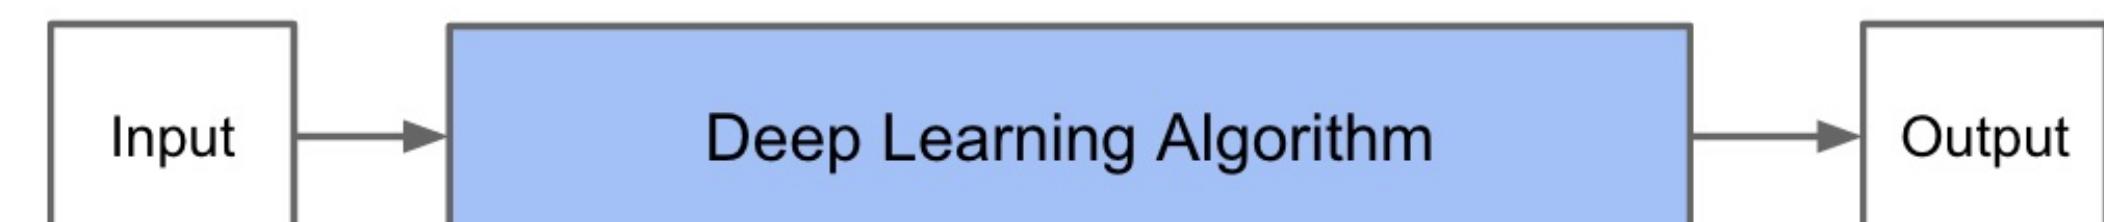


Sample of cats & dogs images from Kaggle Dataset

Promise:
Works better



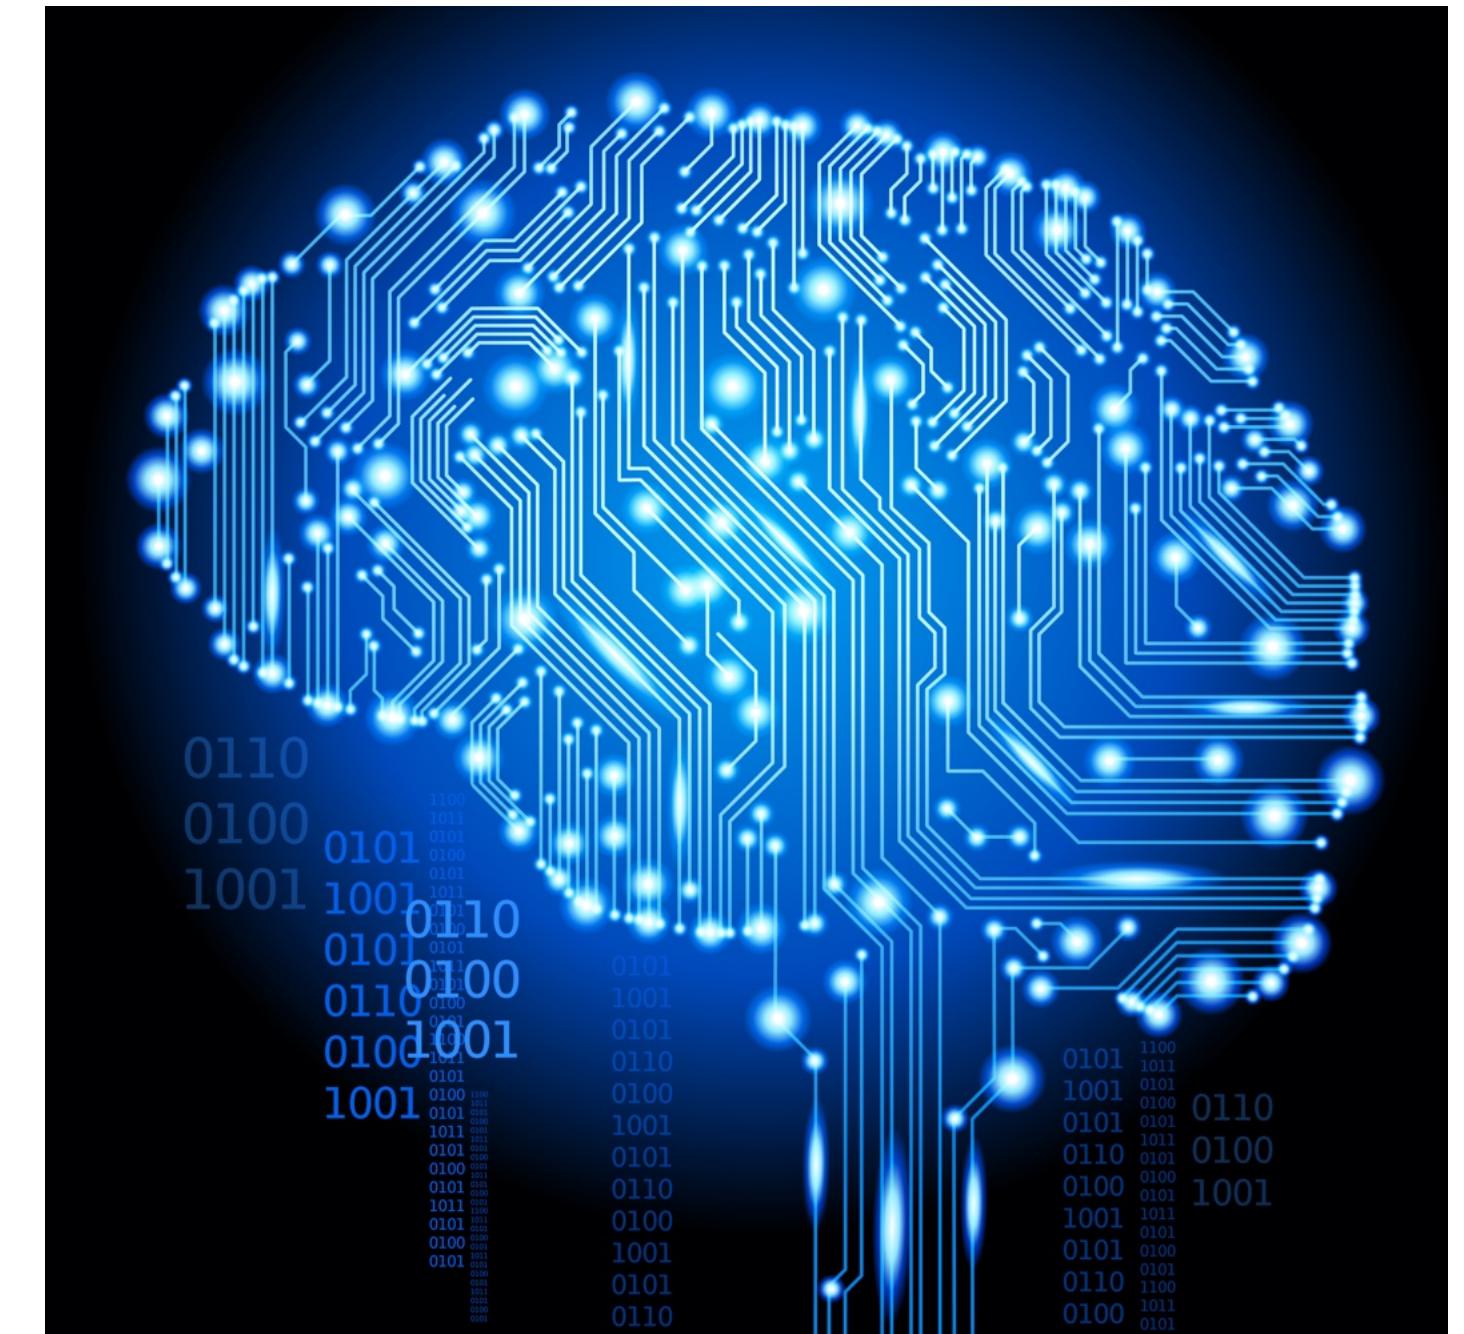
Pitfall:
Blacker box



Deep Learning Flow

Biological inspiration

- Idea: IA should mimic how the brain works.
- Because, the brain solves complex problems.
- But... brain it is complex, even 1 neuron is complex, and we don't fully understand how it works.



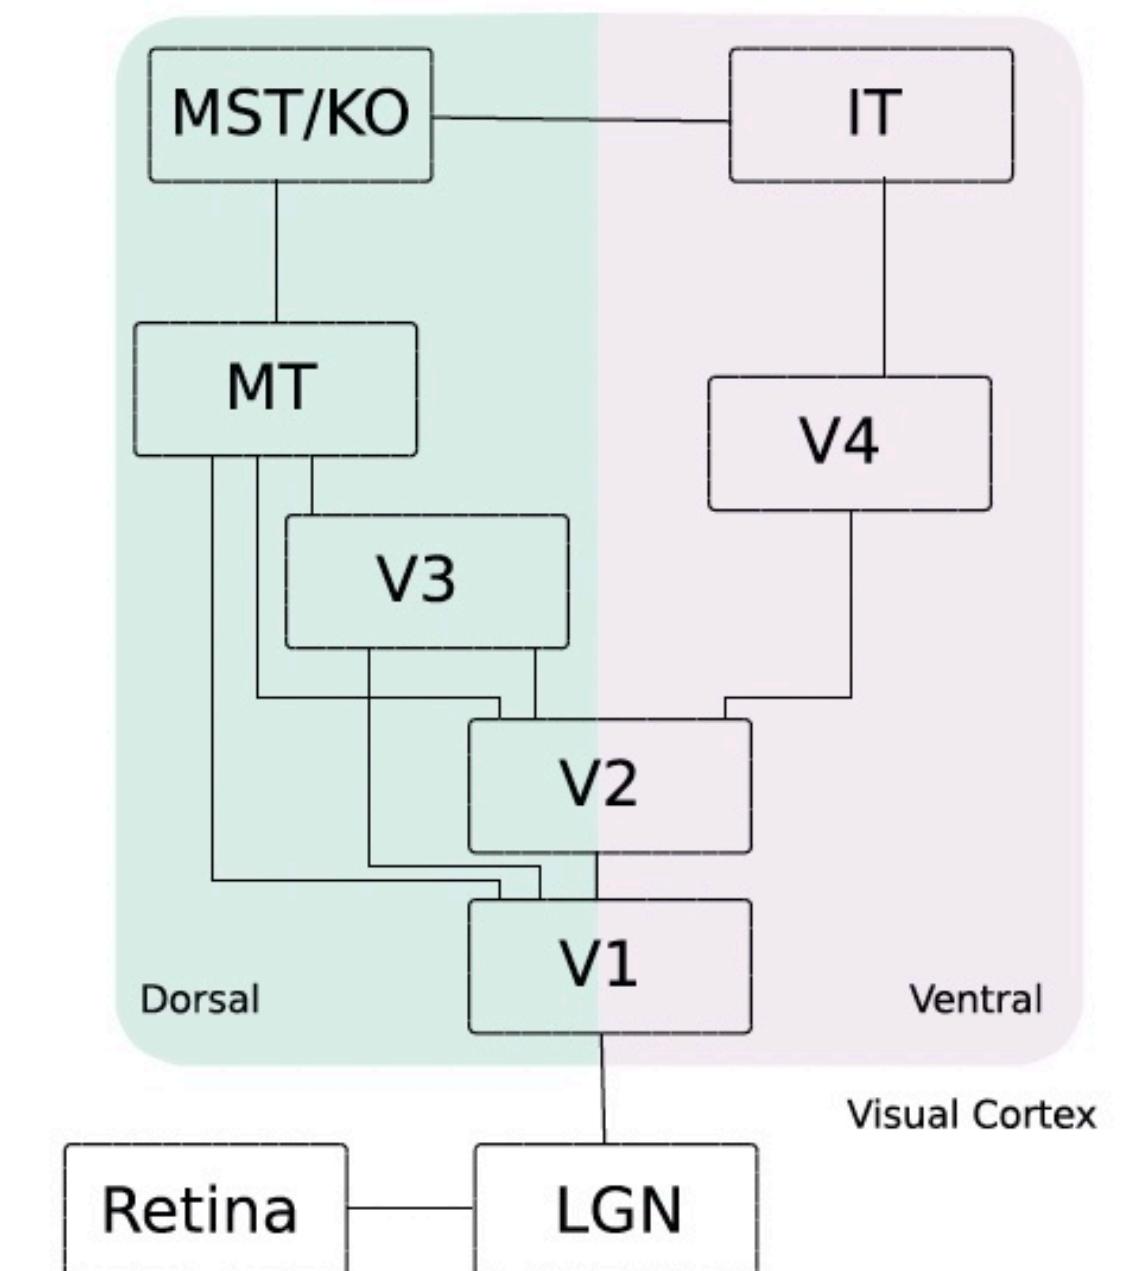
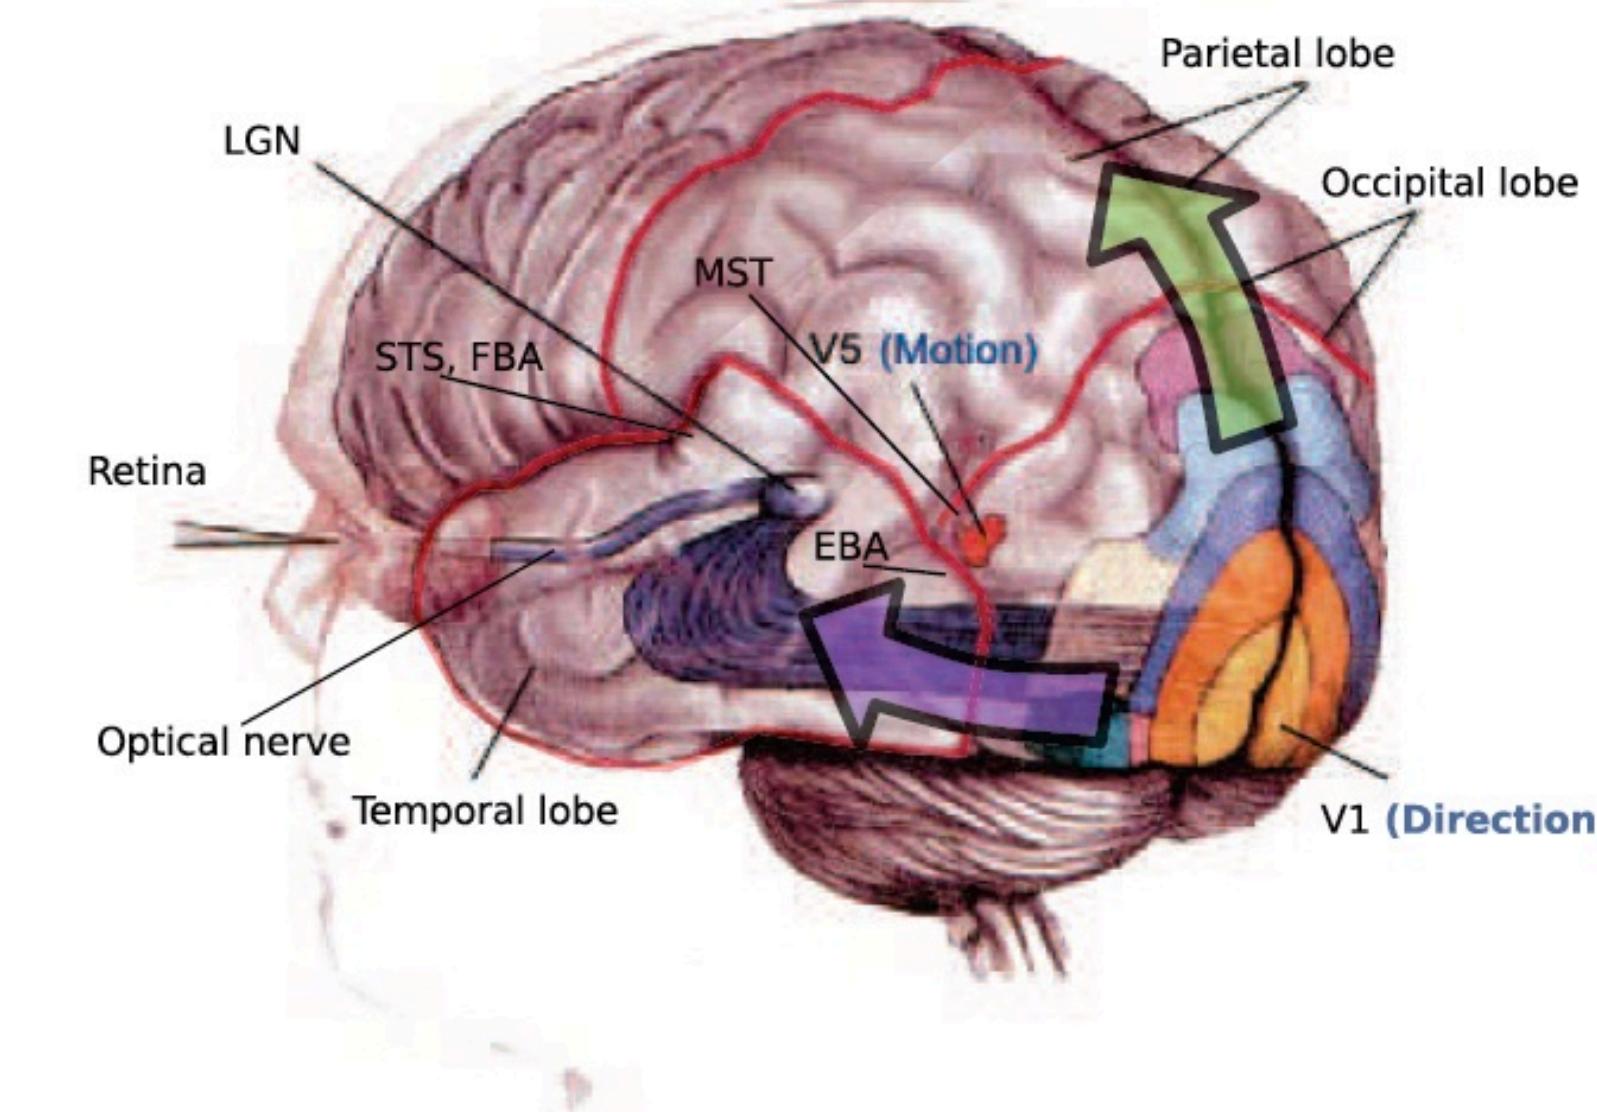
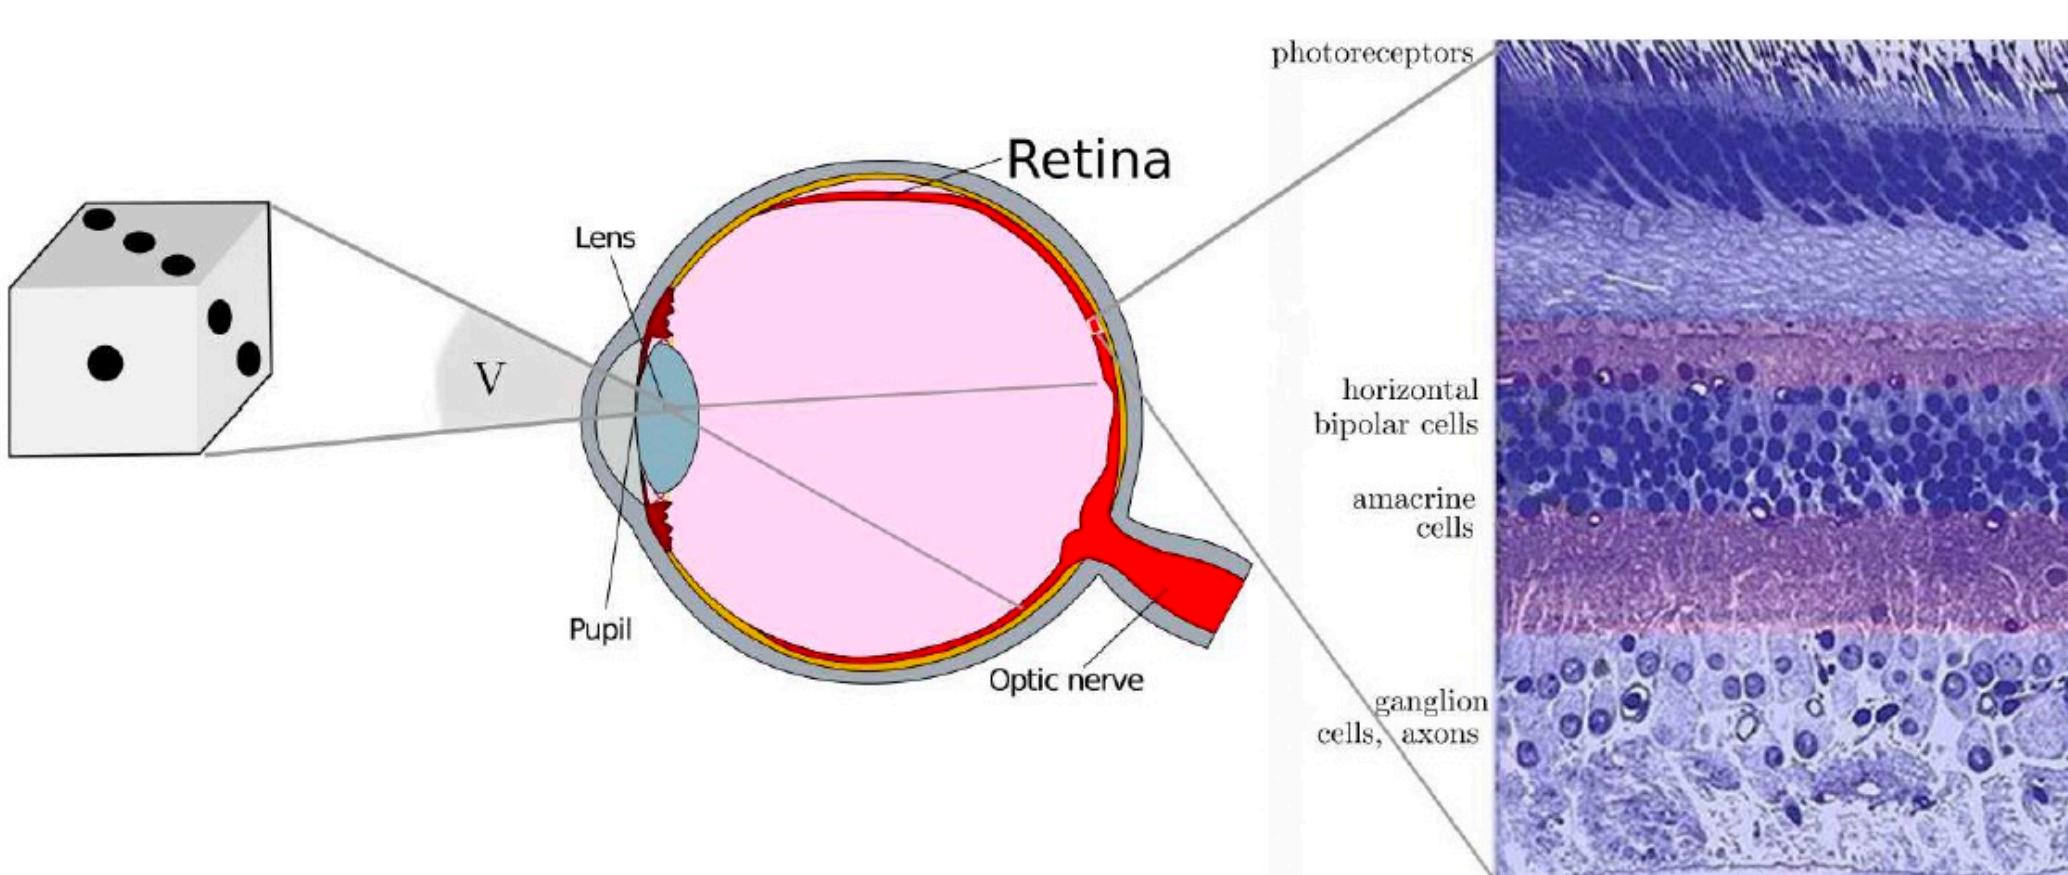
Biological inspiration

- 10^{10} Neurons!
- Response time > 1 [ms] (slow)
- Facial recognition in 100 [ms]
- In average, each neurons has thousands of connections.
- Hundred of operations per second
- High parallelism
- Distributed information representation
- Common failures (Never replaced)
- Compensate problems with massive parallelism.



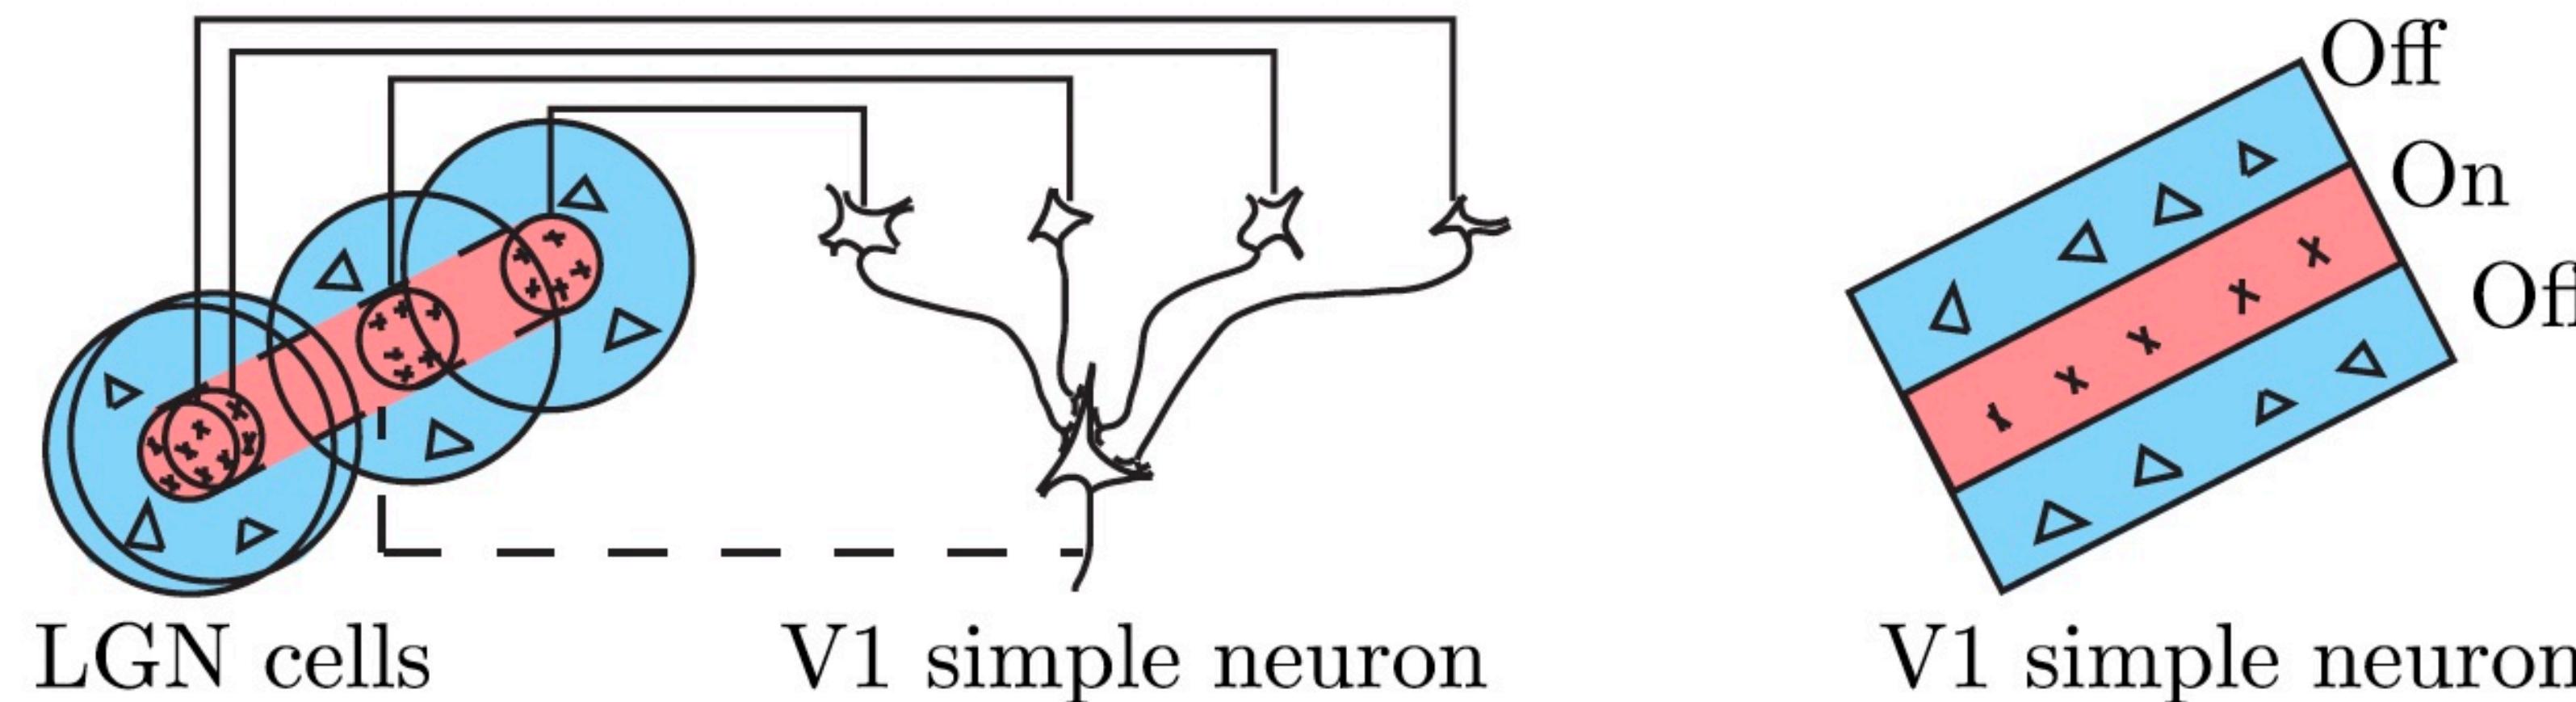
Biological inspiration: hierarchical processing

- From the retina to visual cortex, neurons mainly have a bigger “receptive field”.



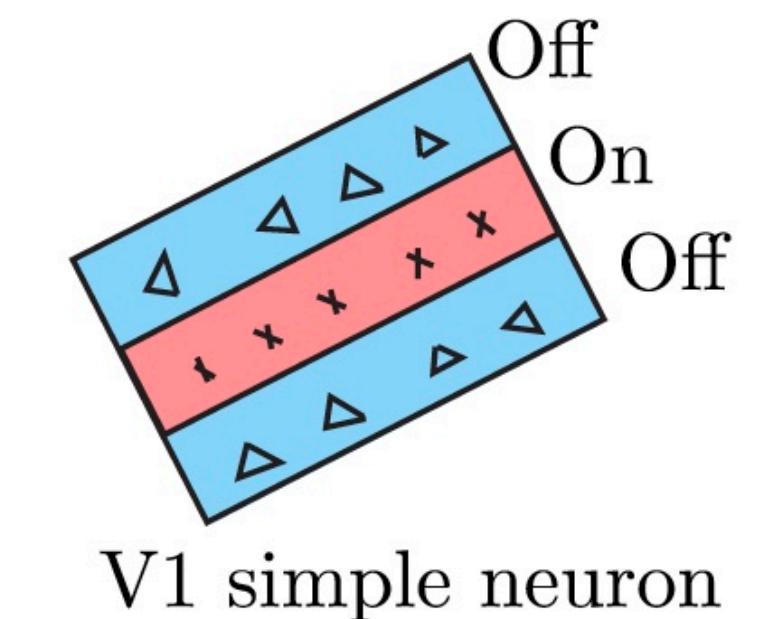
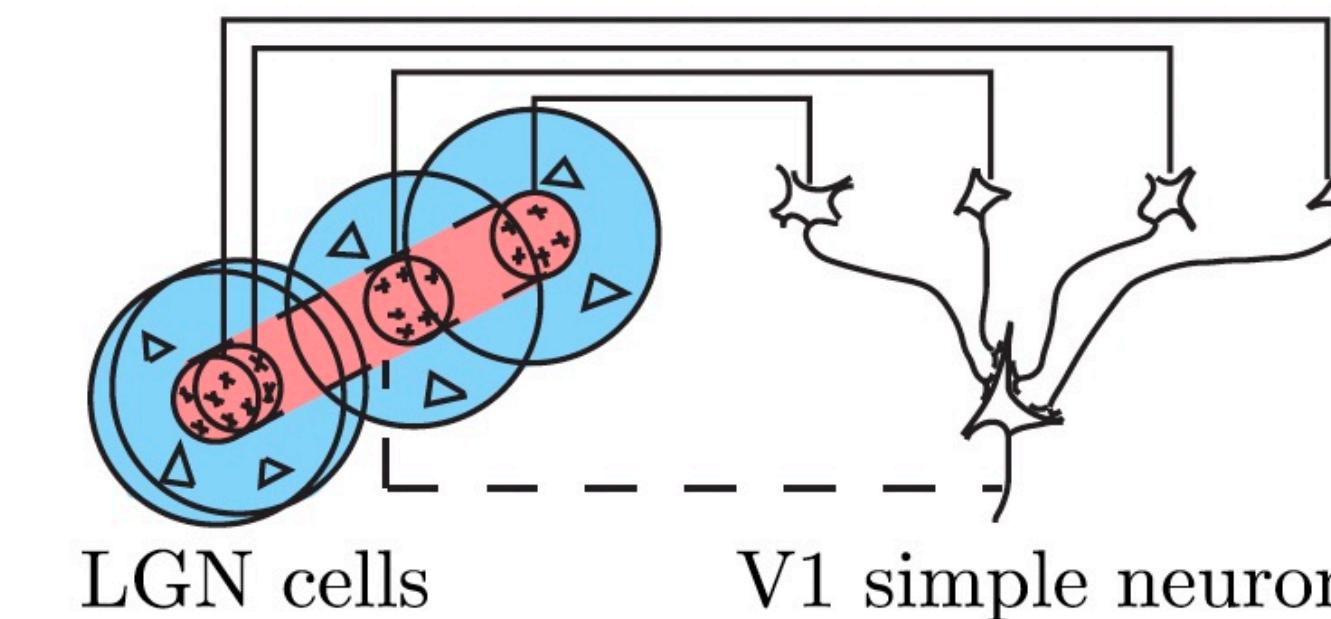
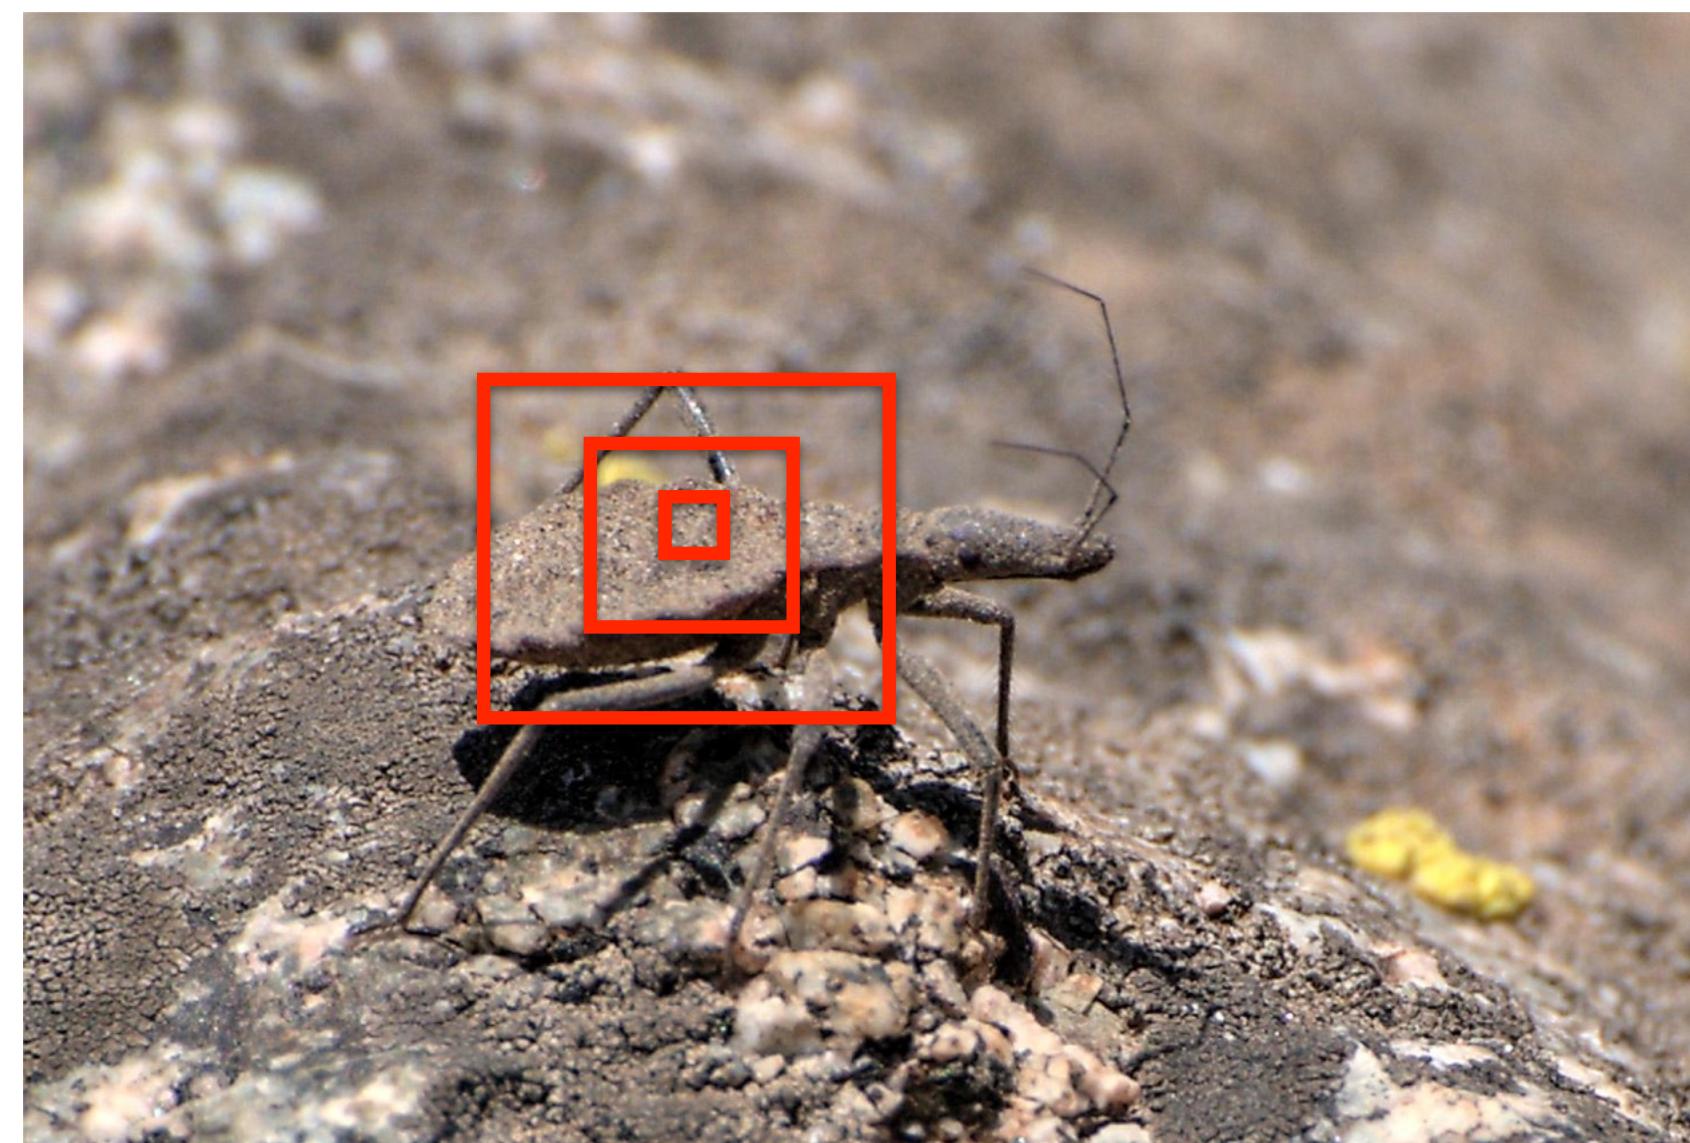
Biological inspiration: hierarchical processing

- At Retina level neurons are sensitive to “contrast” in a specific part of the visual field.
- A precise set of contrast sensitive neurons, can detect a line (V1).
- A precise set of line sensitive neurons can detect a grating pattern (V1)



Biological inspiration: hierarchical processing

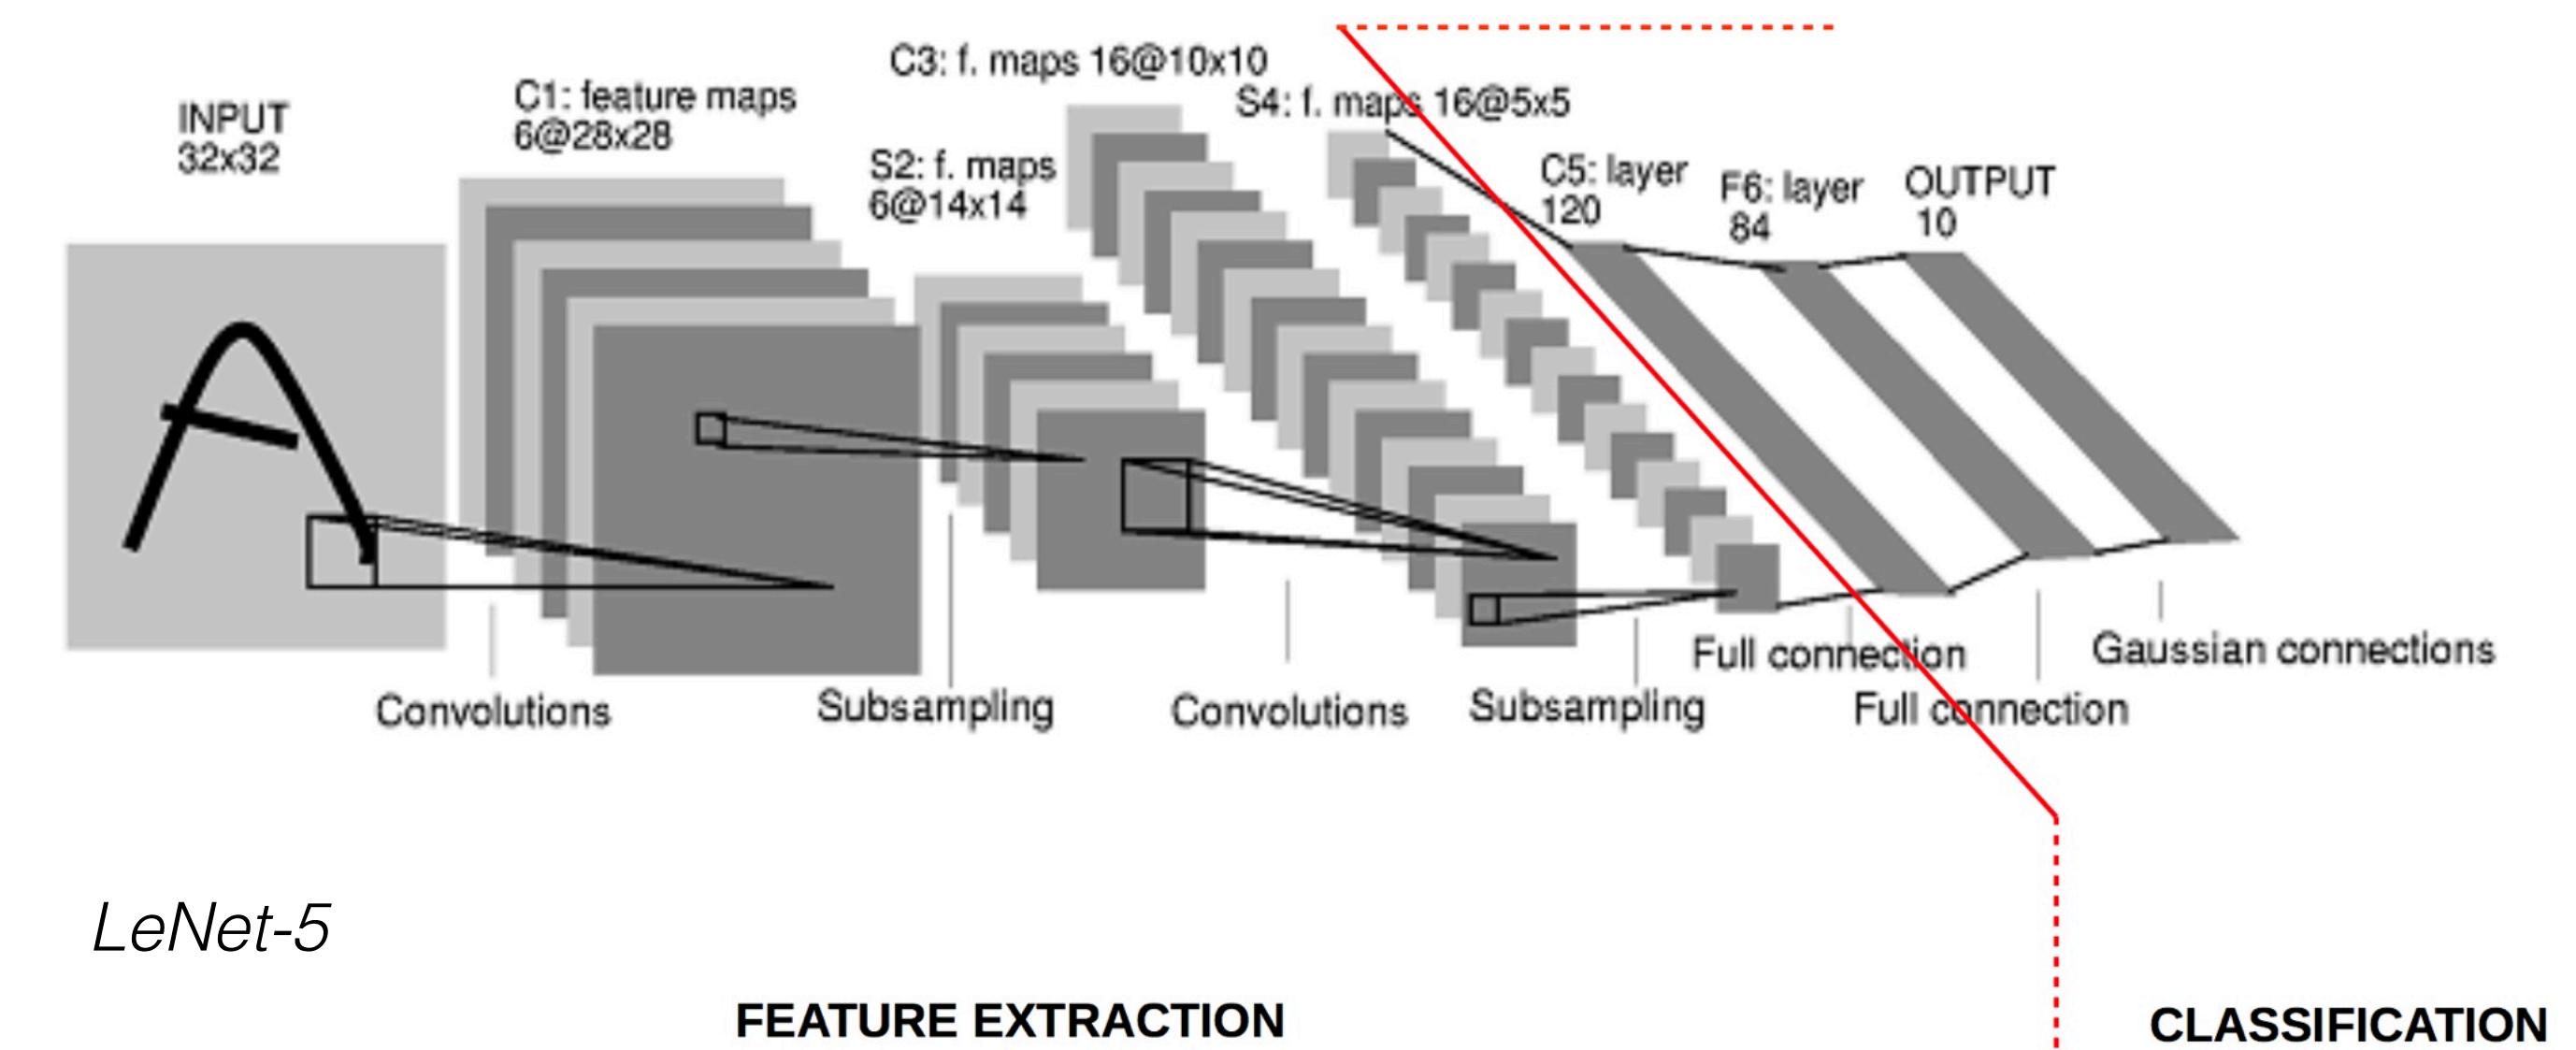
- The receptive field increases as processing gets more complex.



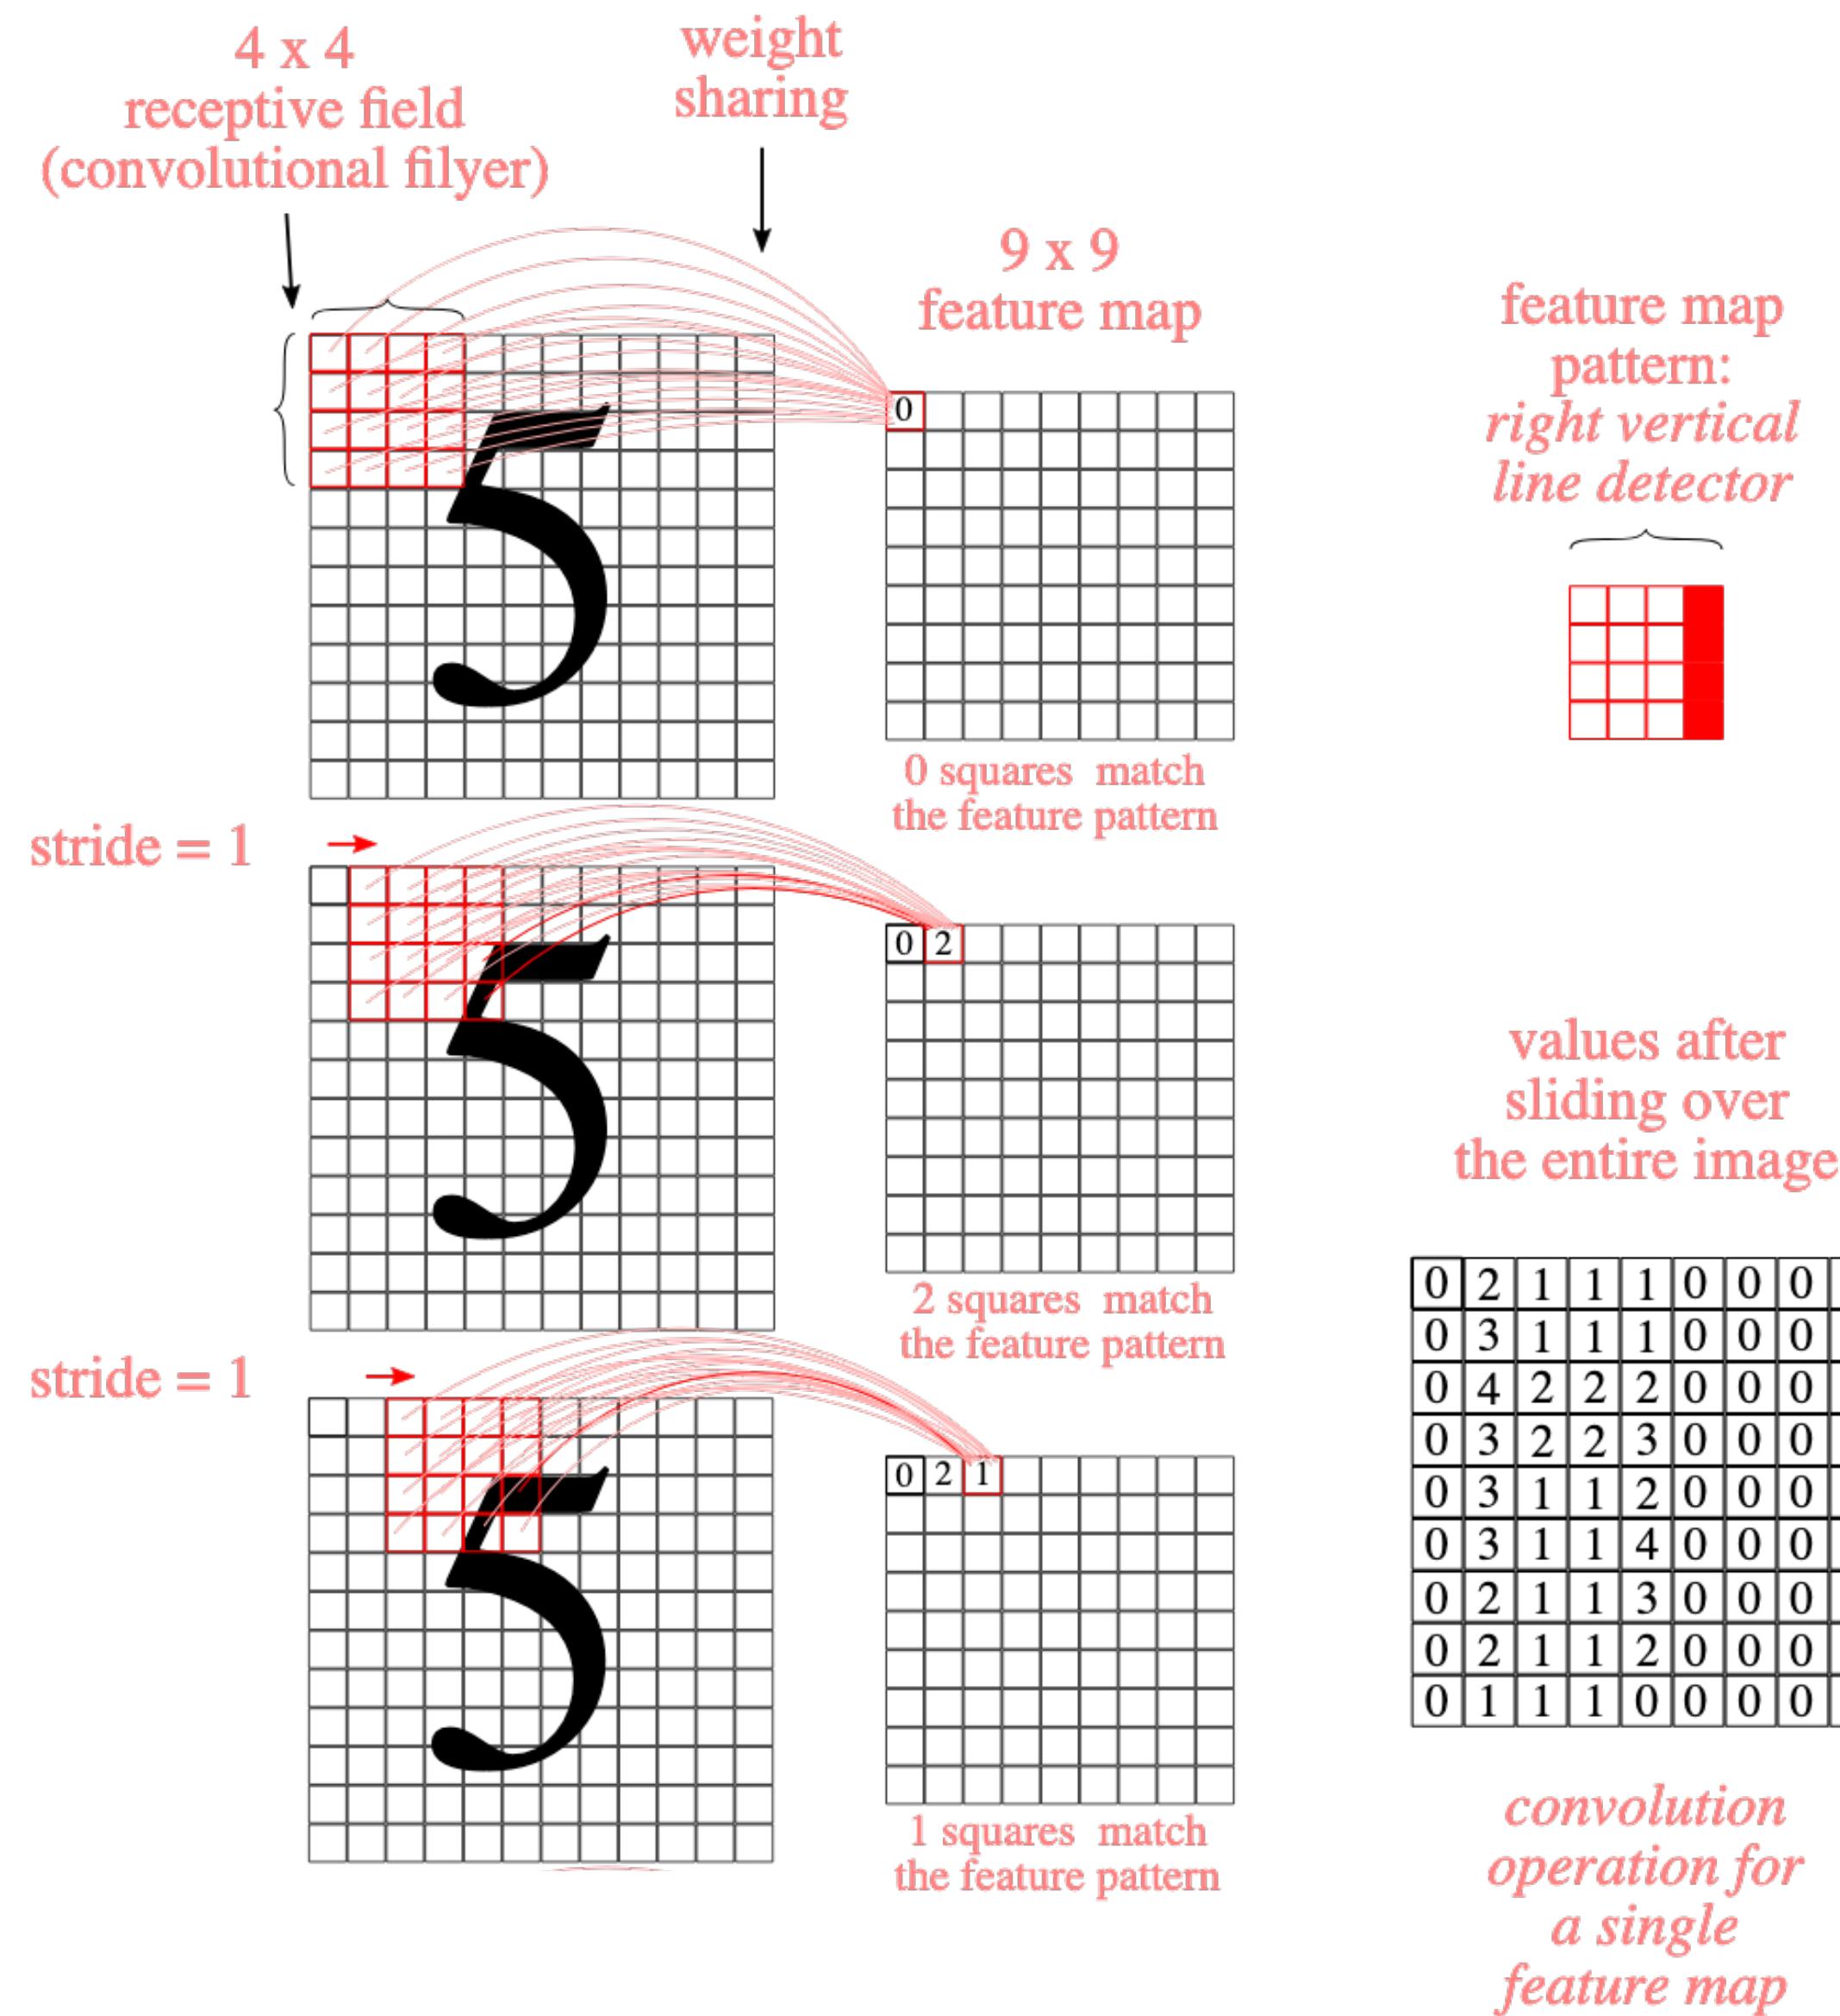
(BUT, 80% of connection are feedback in areas like LGN!)

Convolutional Neural Networks (CNN)

- Set of layers: convolutions, subsampling, fully connected.
- Given an architecture, learn **convolutions** (filters)
- Early version proposed in 90'.
- Take ideas from biological models (neocognitron)

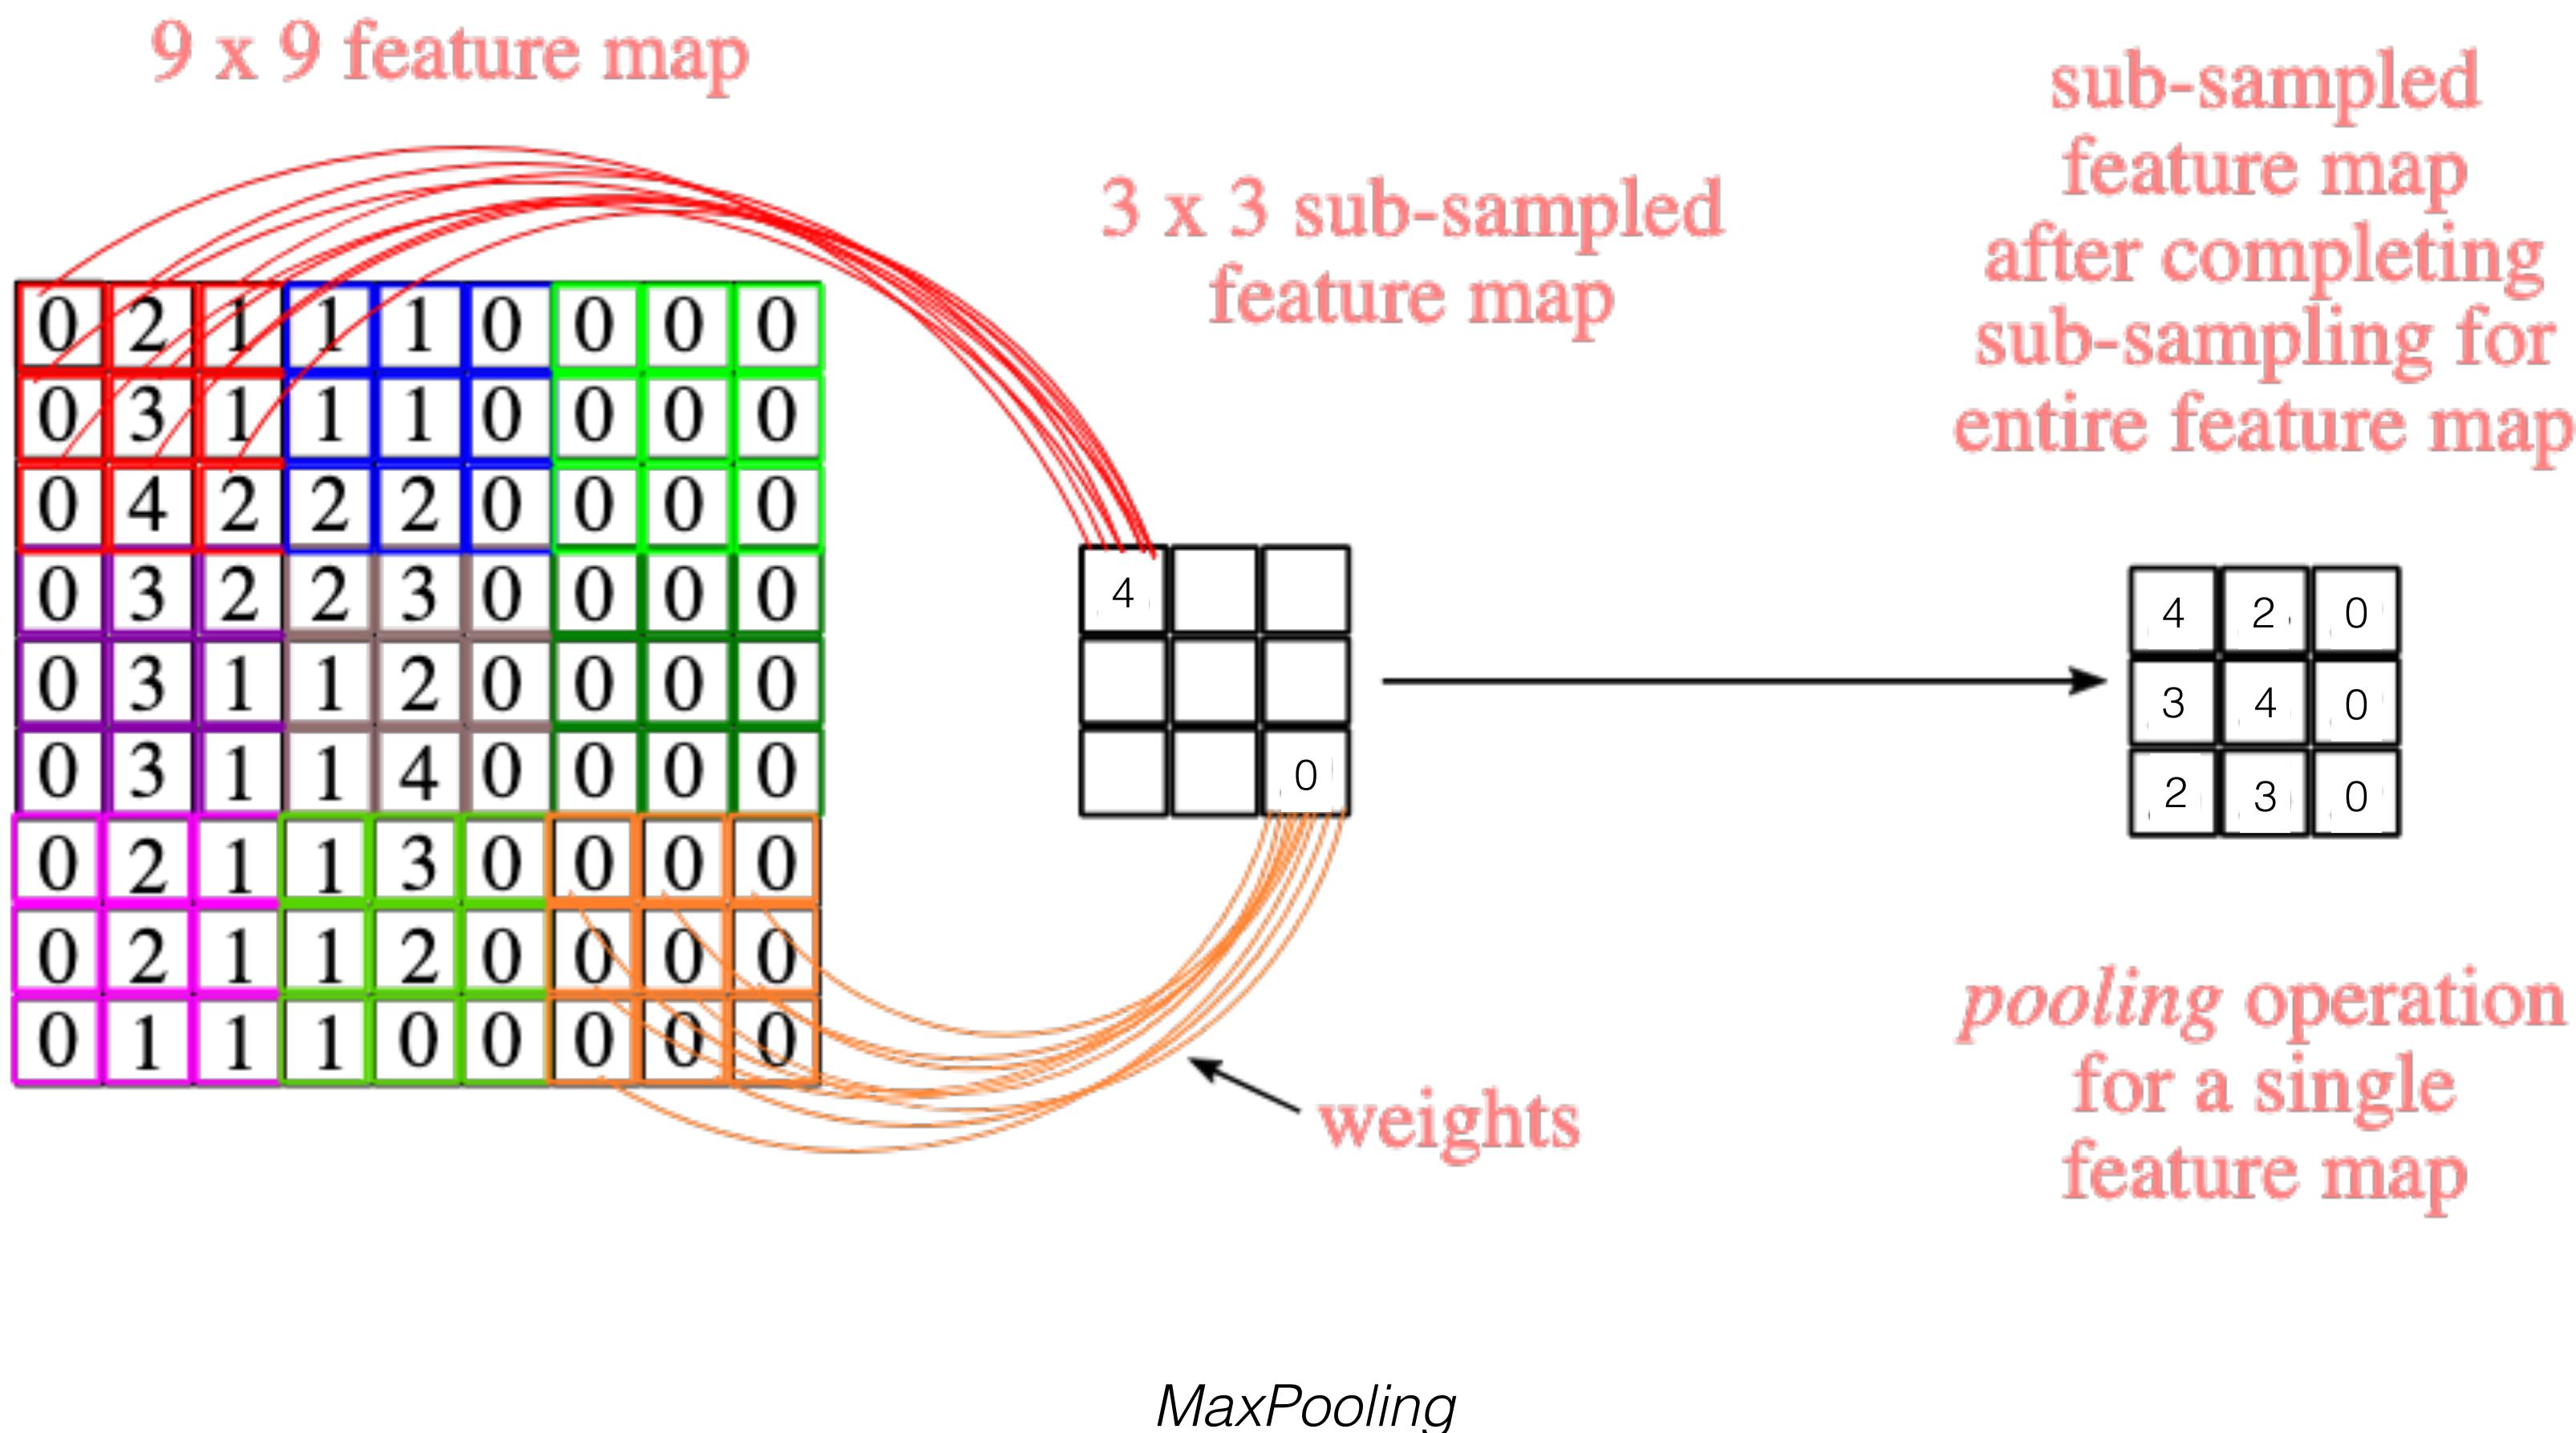


CNN: convolutional layer



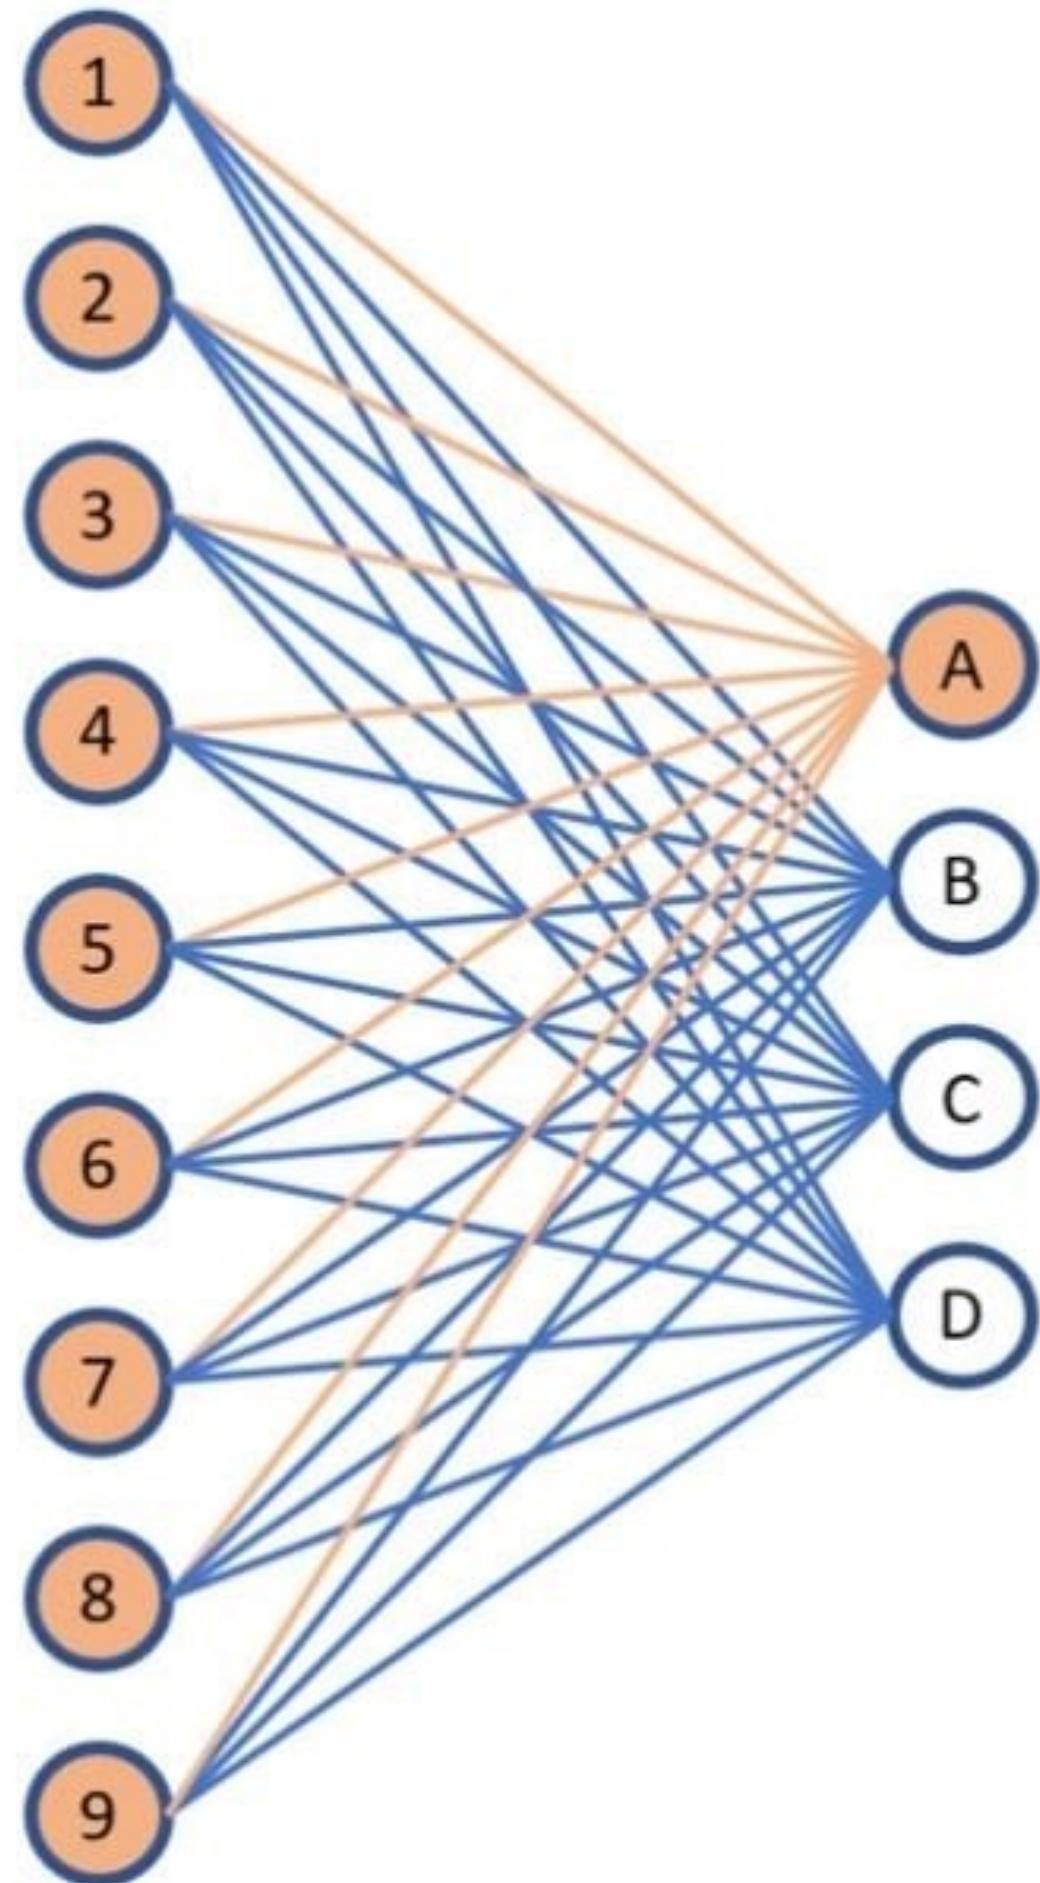
- Learn relevant features.
- Weights must be tuned.
- Increase dimension.

CNN: pooling



- No weights
- Reduce dimension.
- + Spatial invariance.

CNN: fully connected (FC)

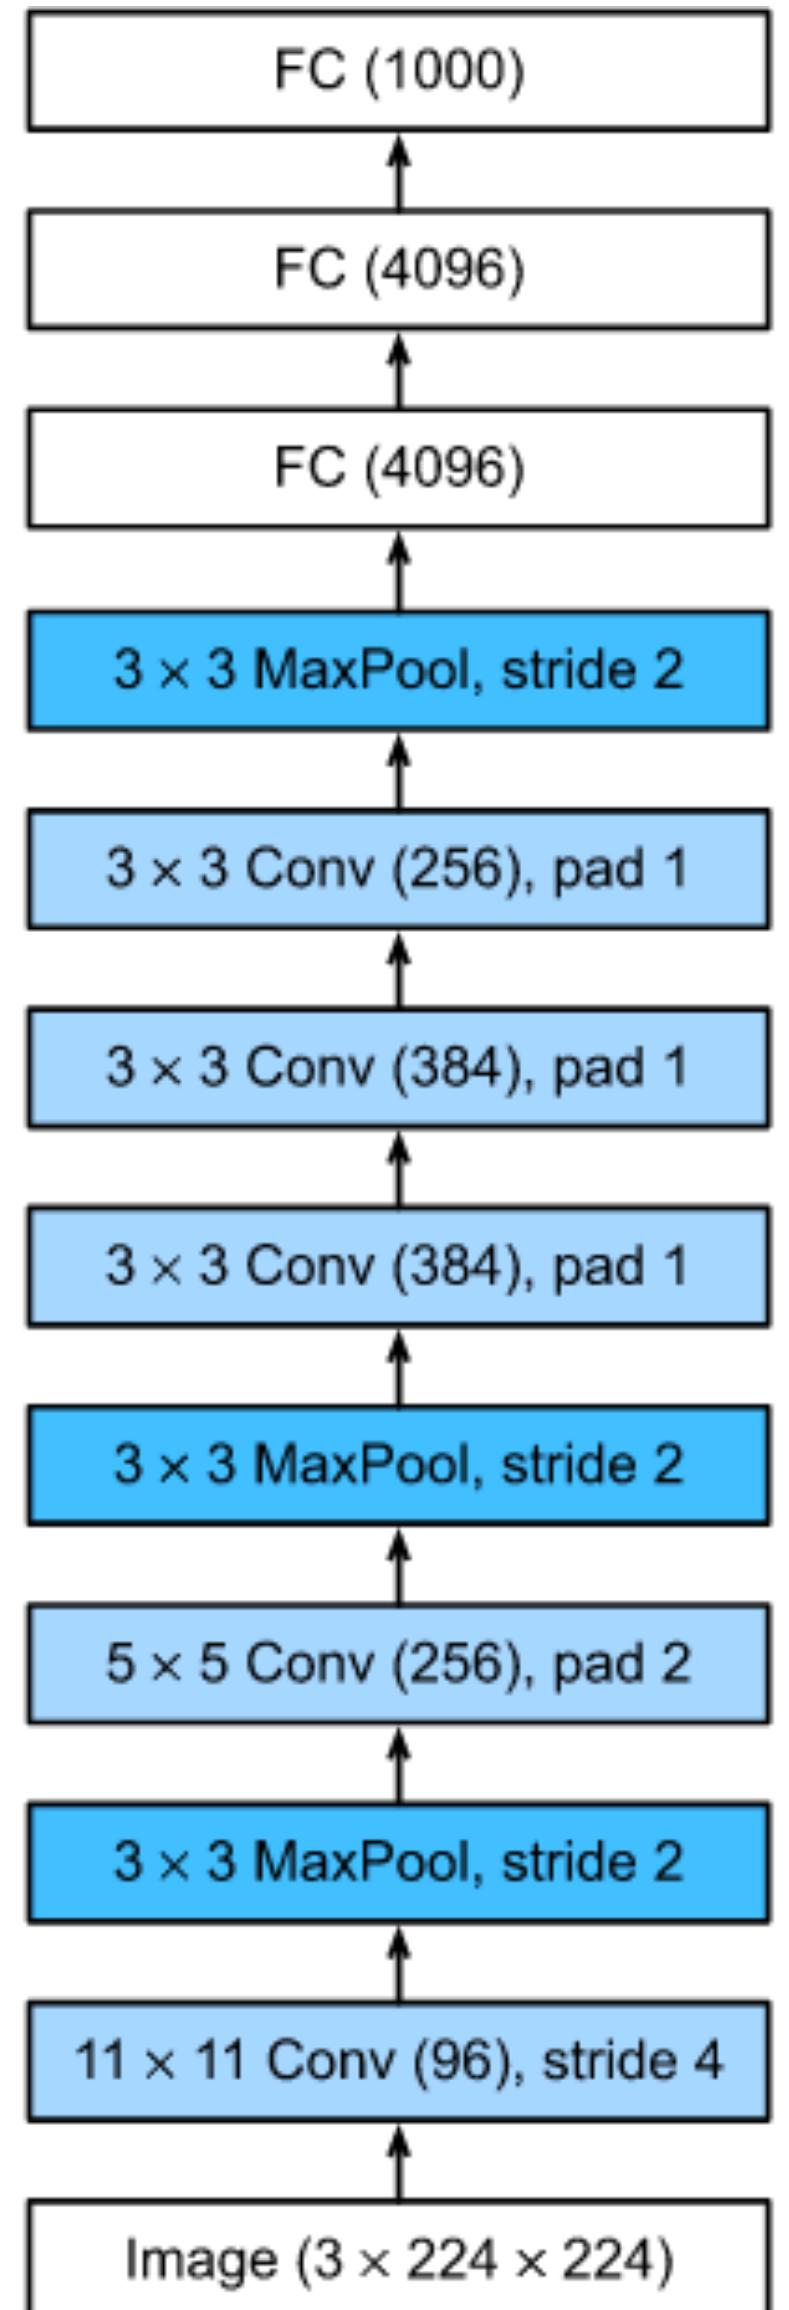


- Learn flexible features.
- +Weights to be tuned.

CNN: AlexNet main features

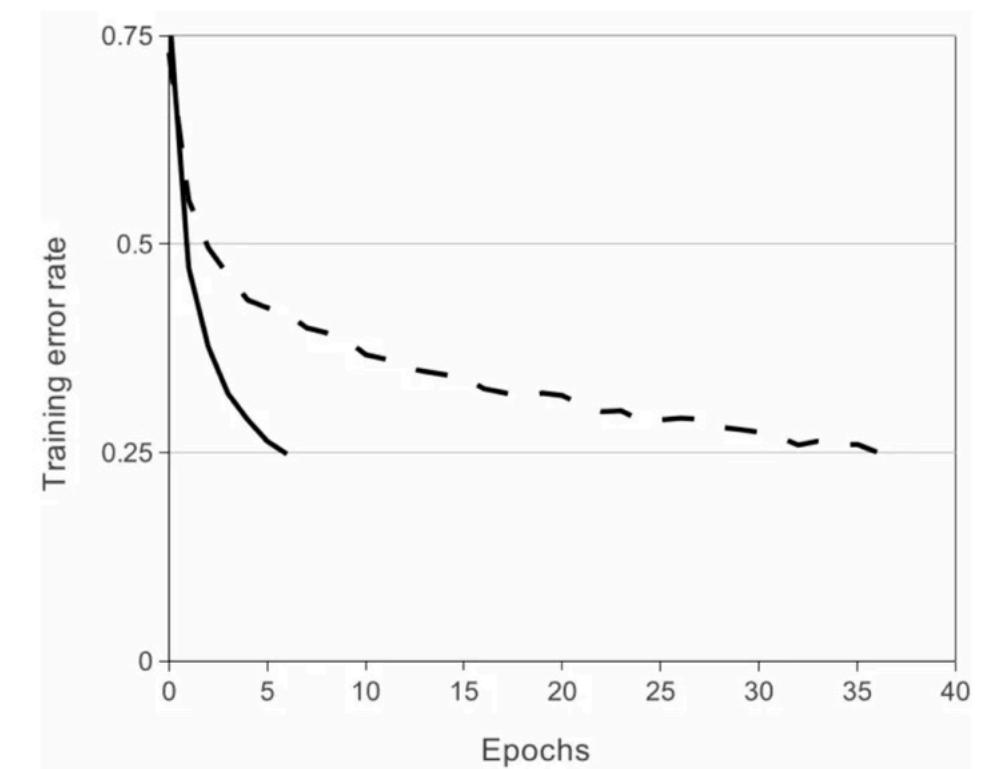
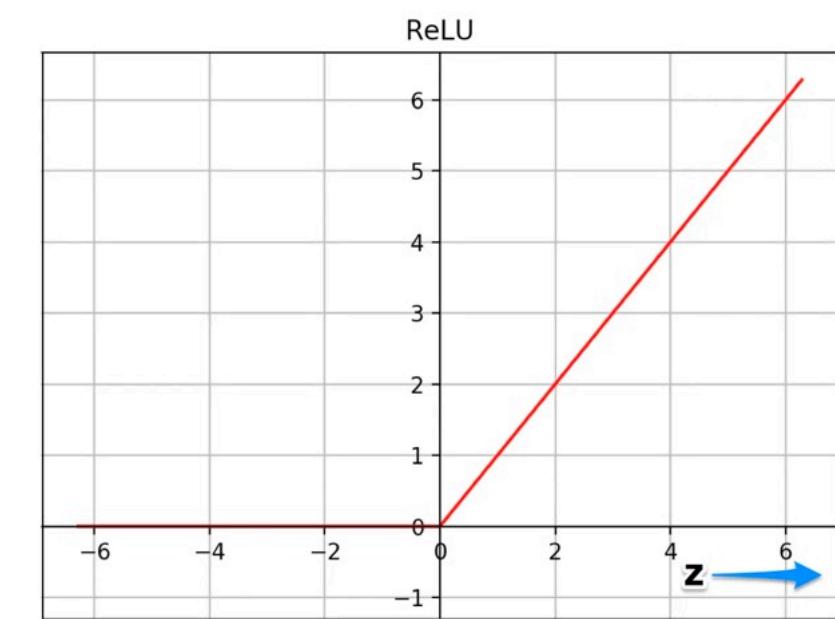
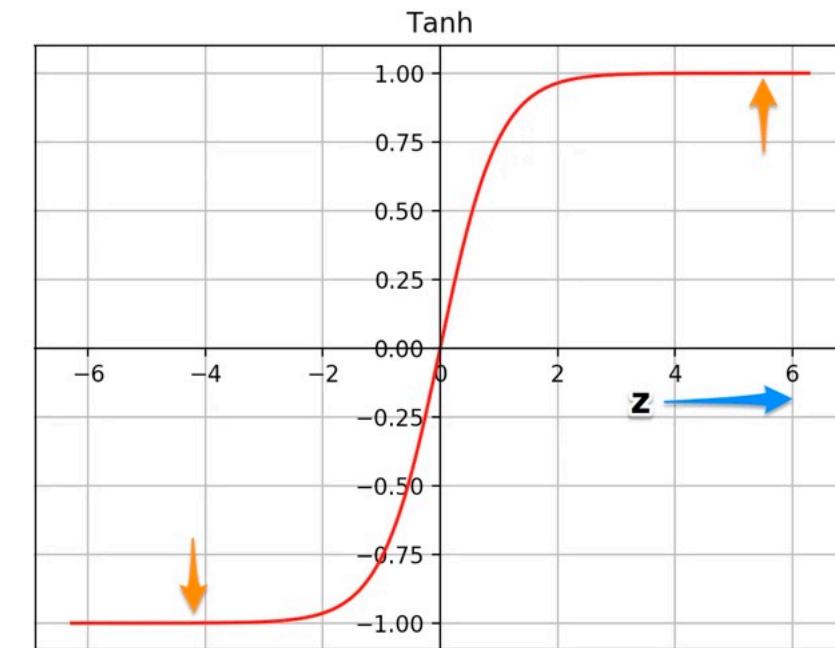


- 8 layers (5 convolutional, 3 FC).
- conv layer + **RELU**.
- **Dropout** in first 2 layers.
- Max-pooling in layers 1, 2, 5.
- GPU implementation.
- Large Scale Visual Recognition Challenge 2012 winner.



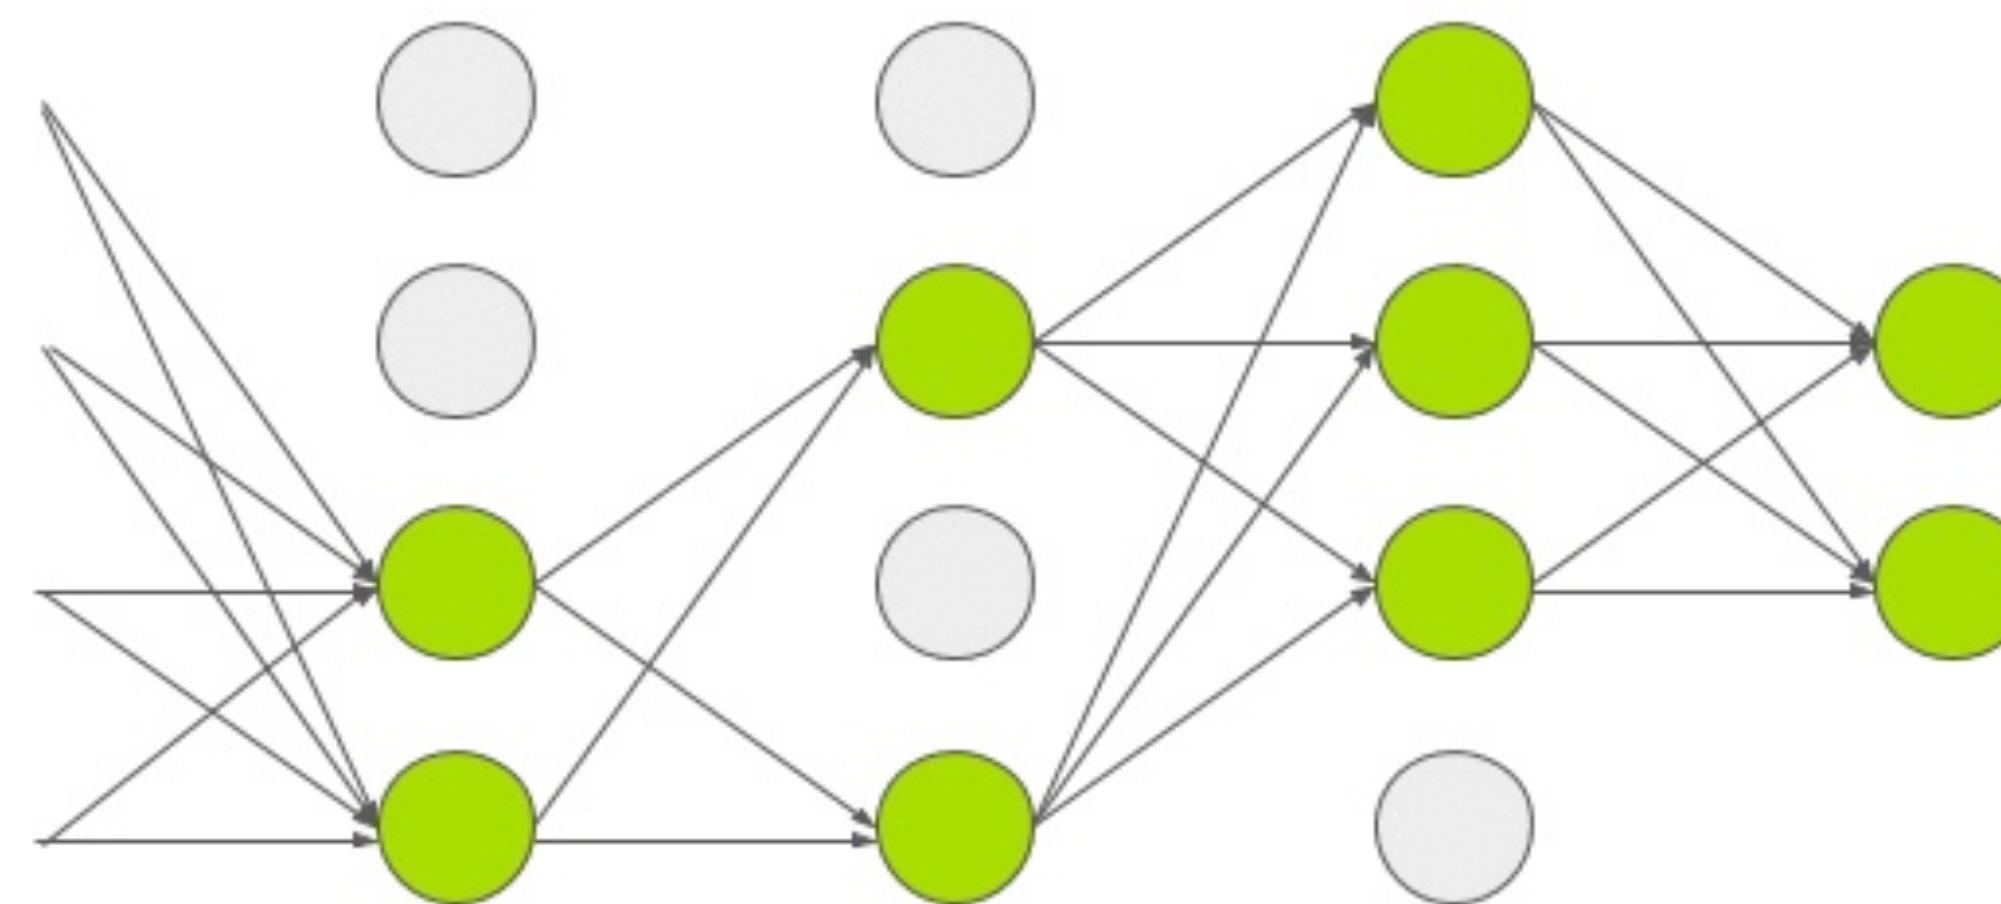
CNN: AlexNet activation function

- Until 2012 tanh / sigmoid were most common activation function.
- Tan / sigmoid saturate, making gradient descend slower.
- Experimentally RELU allowed for faster training.



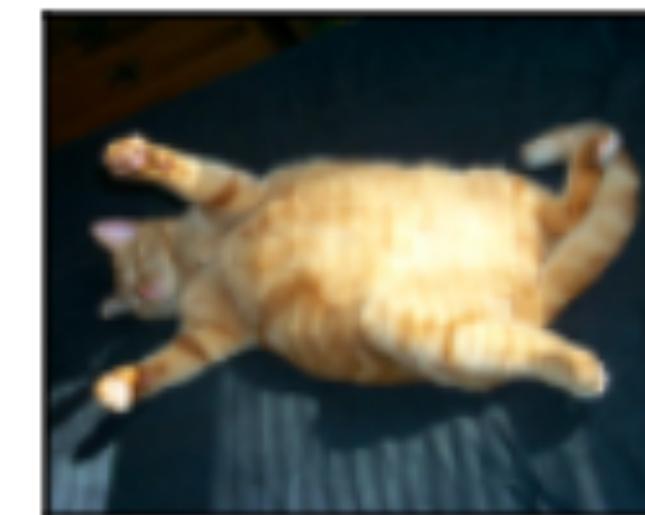
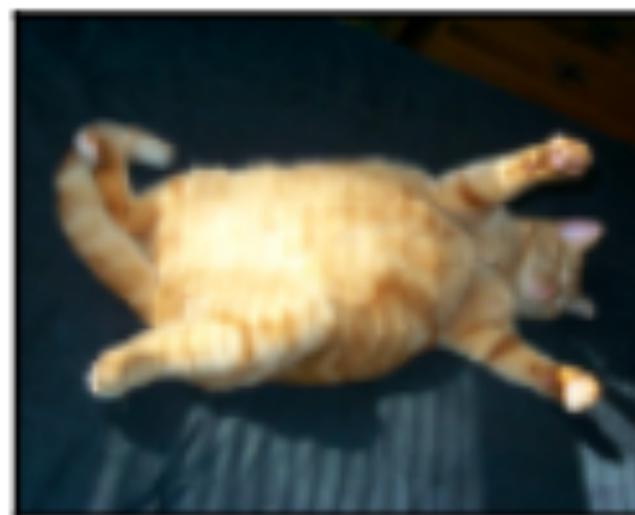
CNN: AlexNet and overfitting

- To reduce overfitting dropout was proposed.
- Each neuron can be disabled during learning with a certain probability.
- (Slower training, requires bigger datasets)

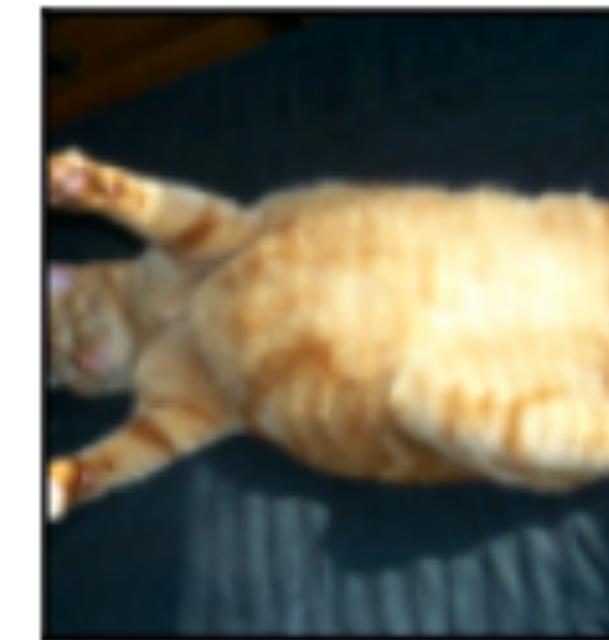


CNN: AlexNet and overfitting

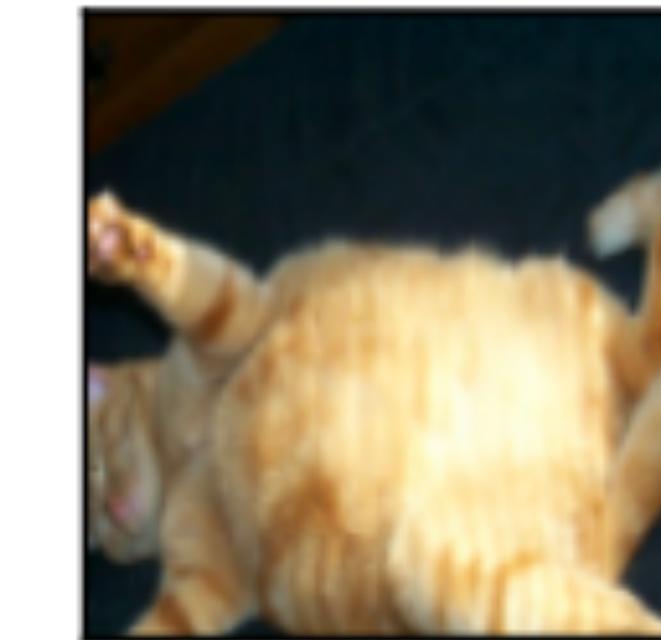
- To generate image variants that the network must be robust.
- Rotations, translations, “noisy image”.
- Idea used today!



Flip horizontal



ZoomIn + crop



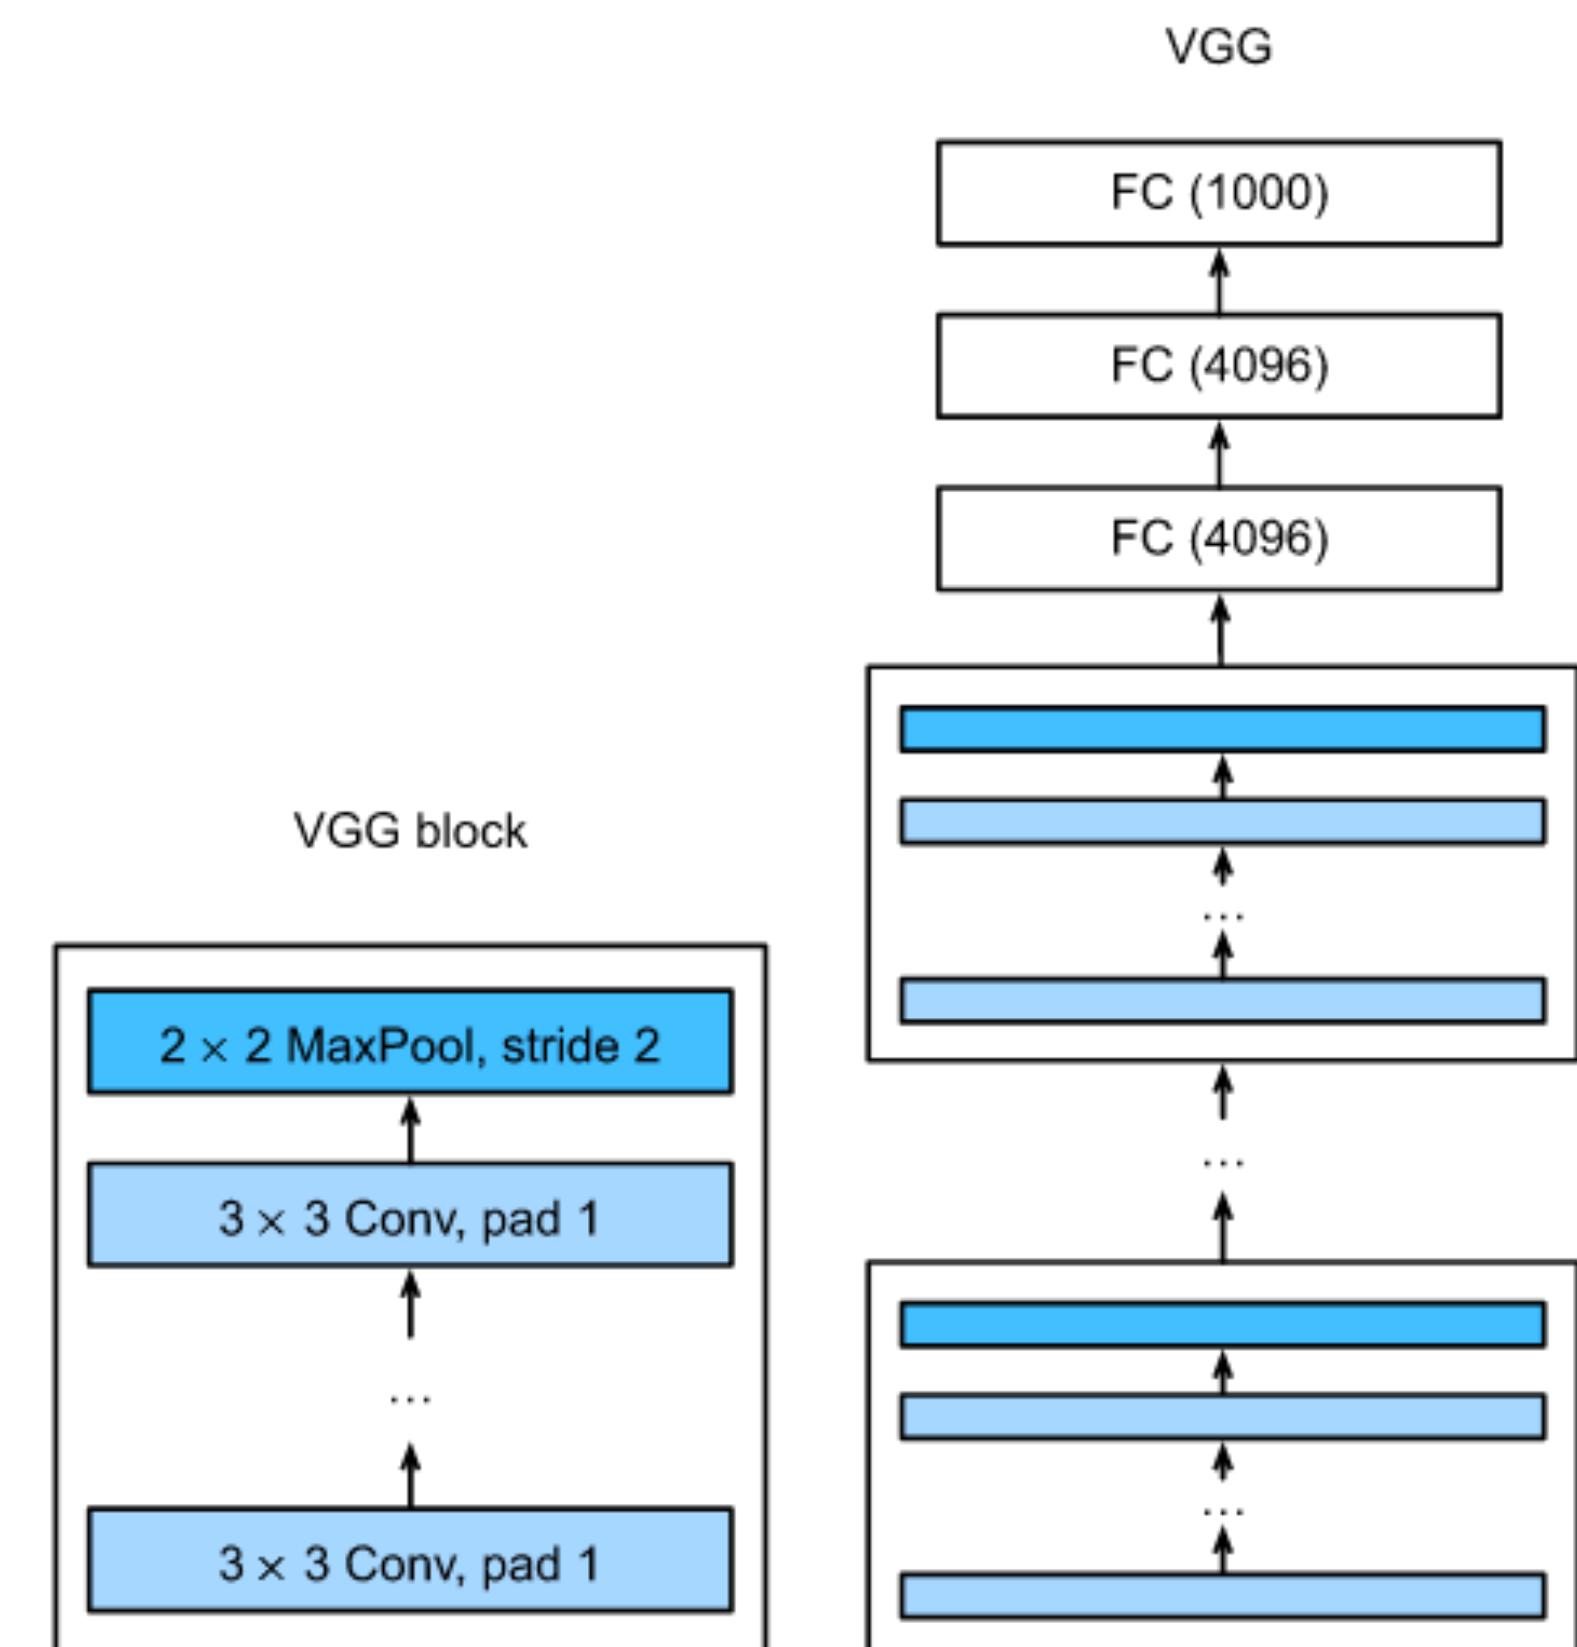
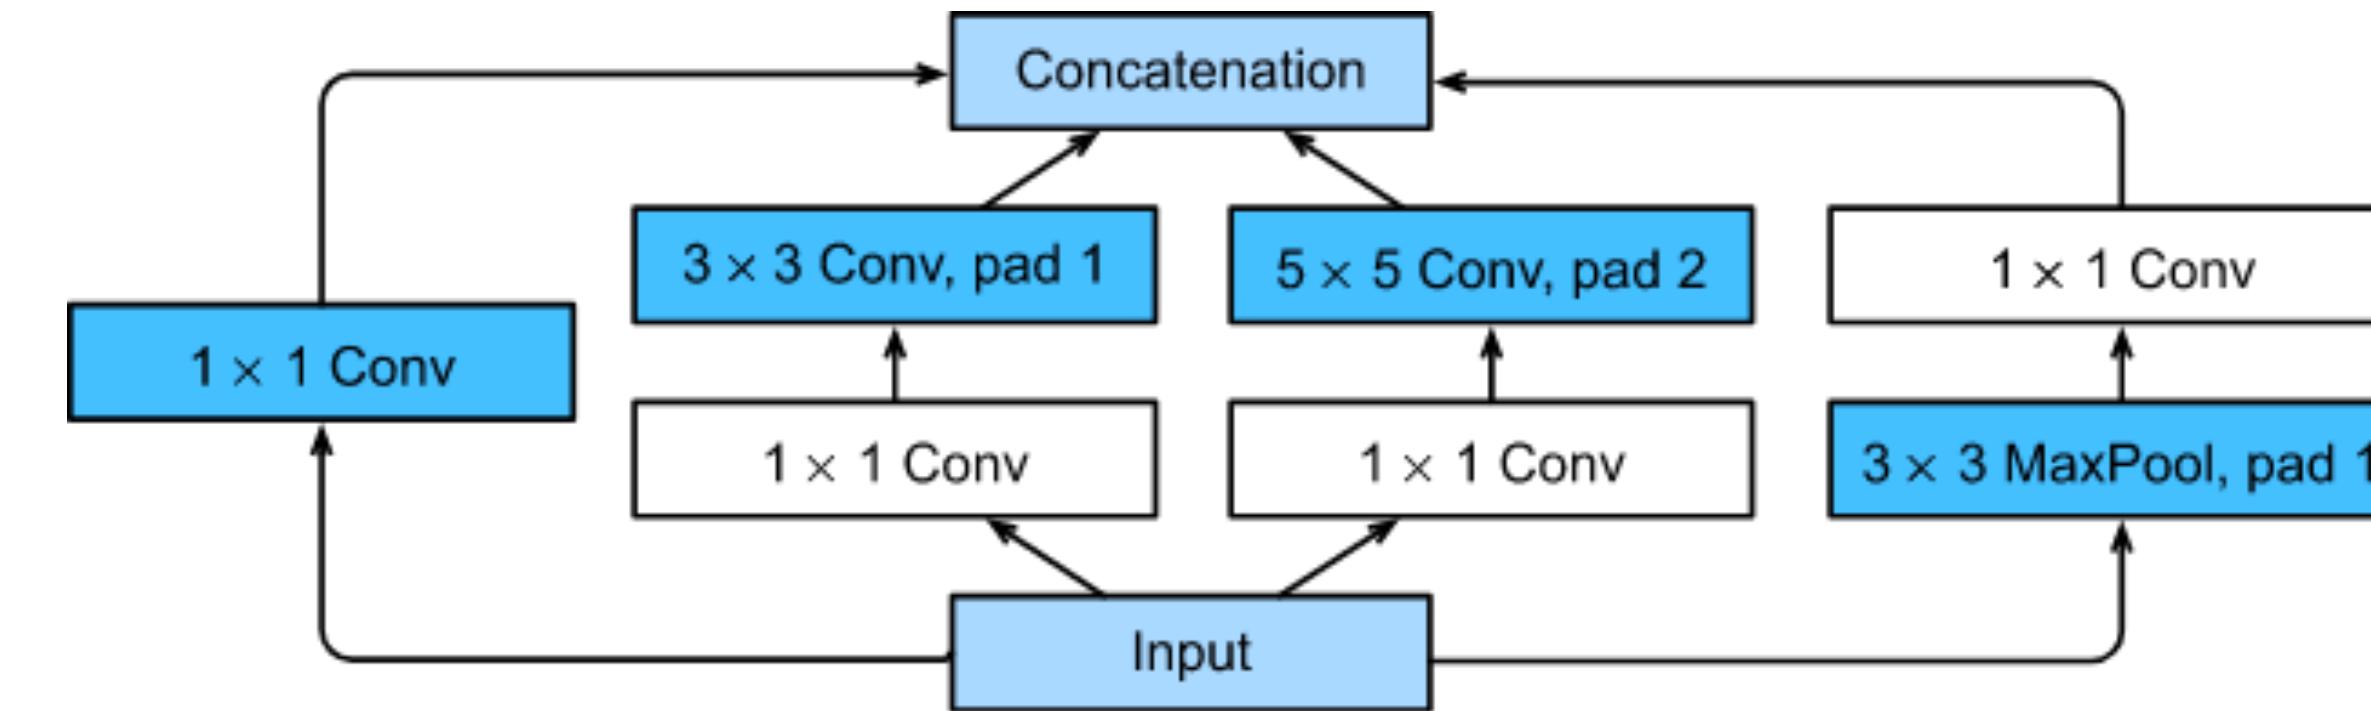
Variación de color



CNN: Visual Geometry Group (VGG)

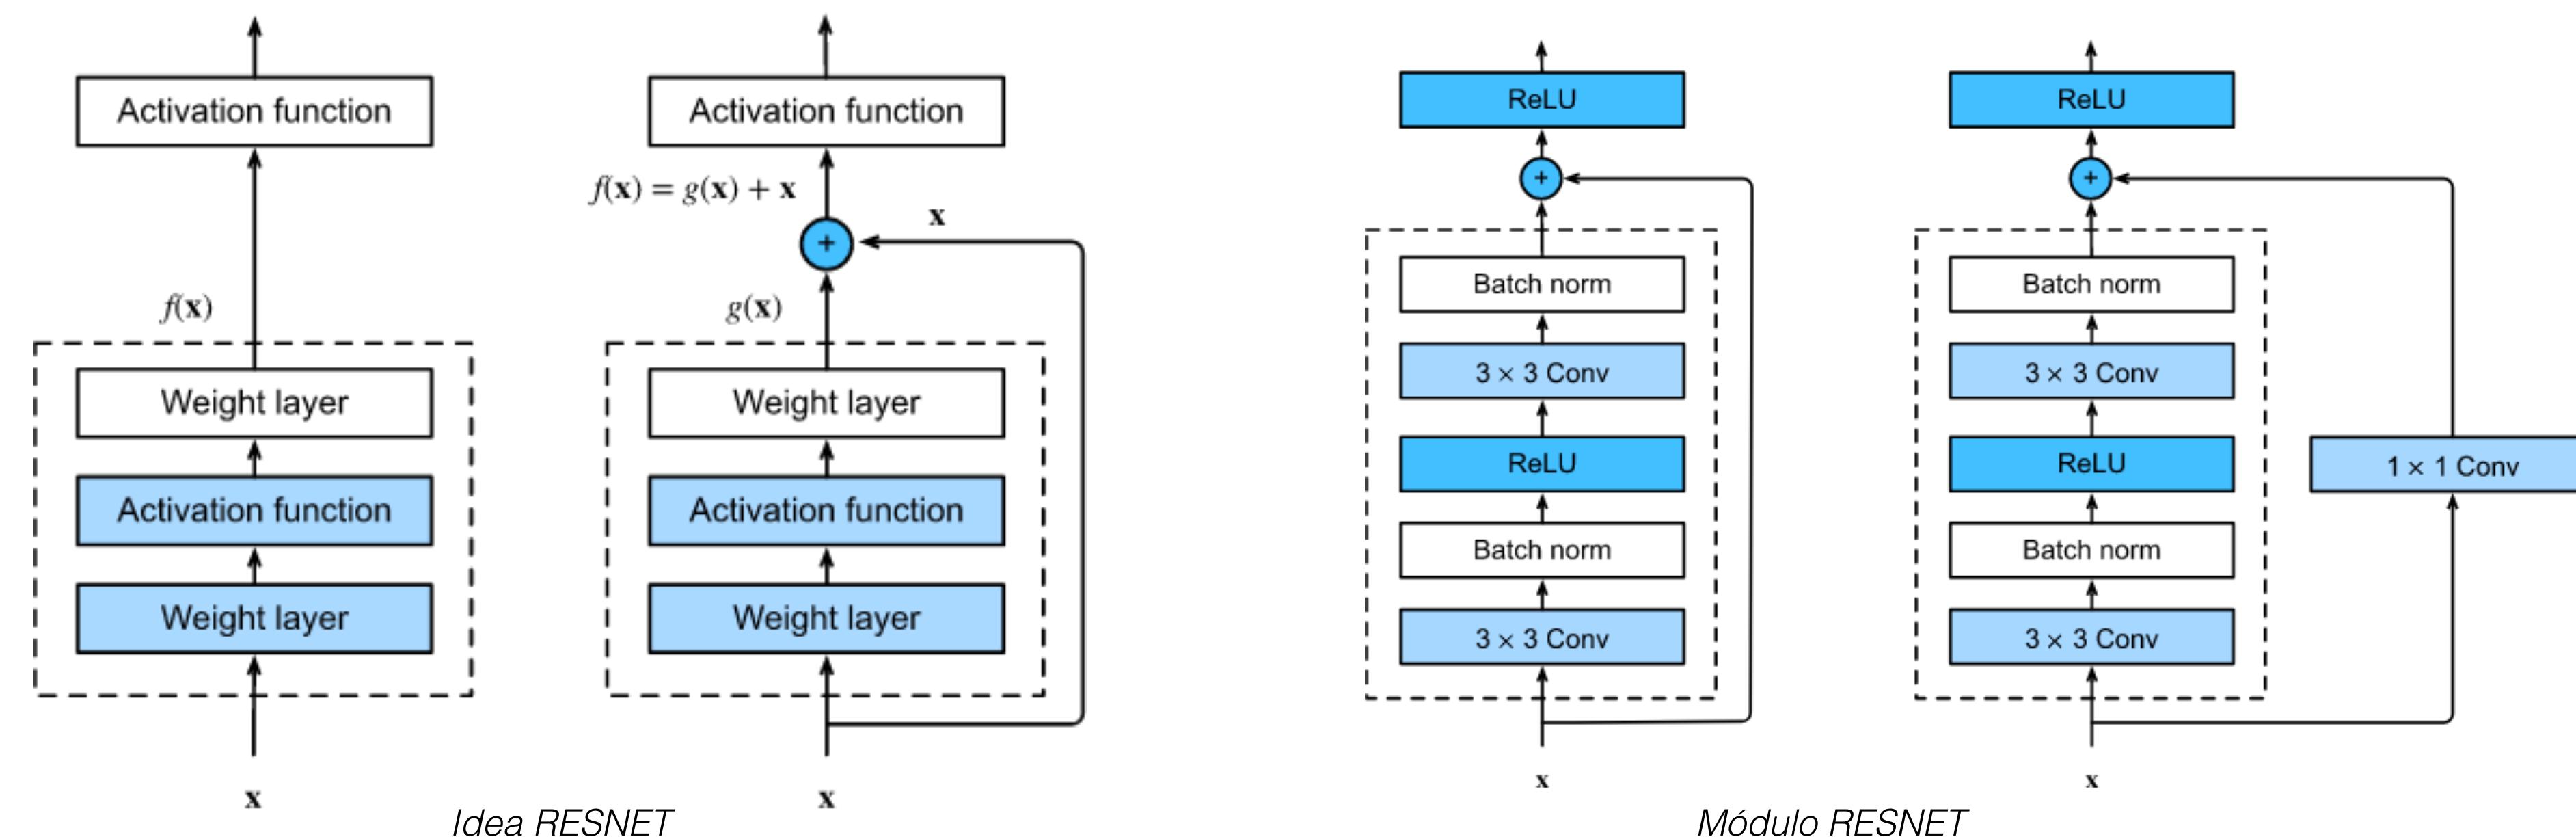


- Conv -> Relu -> pooling quickly lose resolution
- VGG showed that multiple 3x3 conv are equivalent to a 5x5 conv.
- Introduces the “block” idea to configure according to problem/resources.



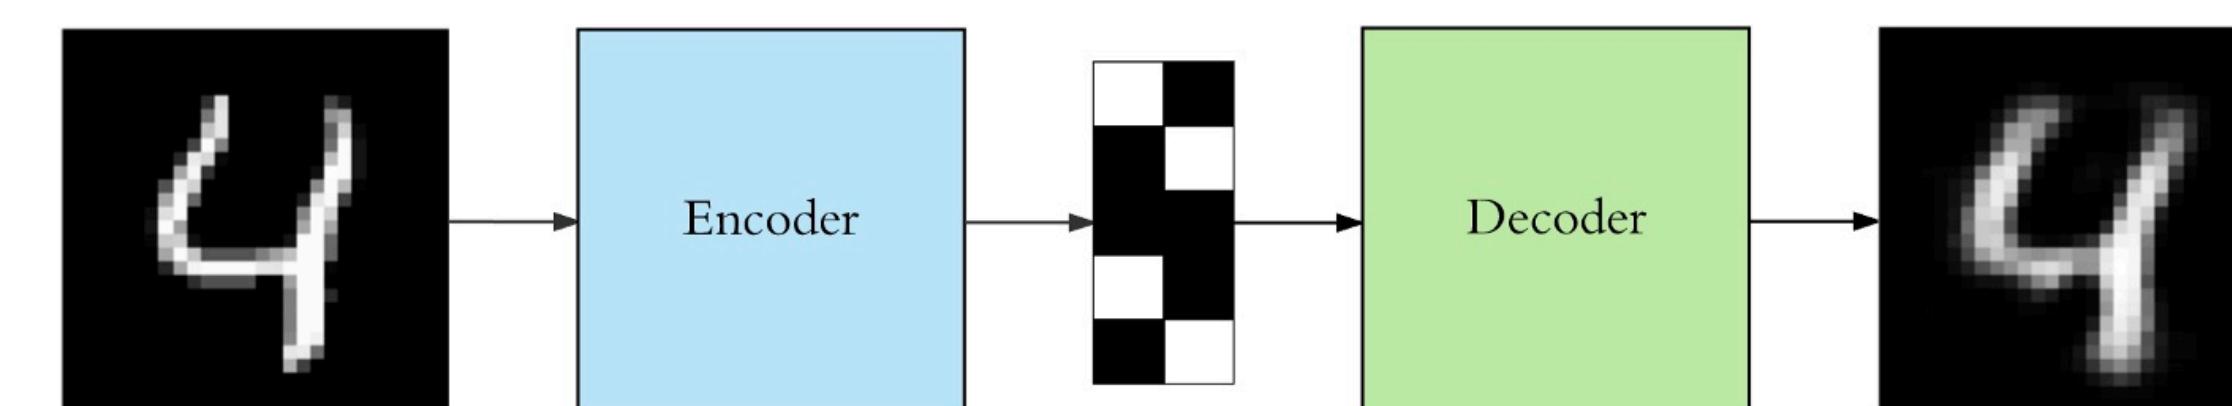
CNN: ResNet

- If we want the network learns $f(x)$ we want a module to learn from residual $g(x)$.
- If $f(x)=x$, $g(x)=0$, we always have the option blocks does nothing (subset).

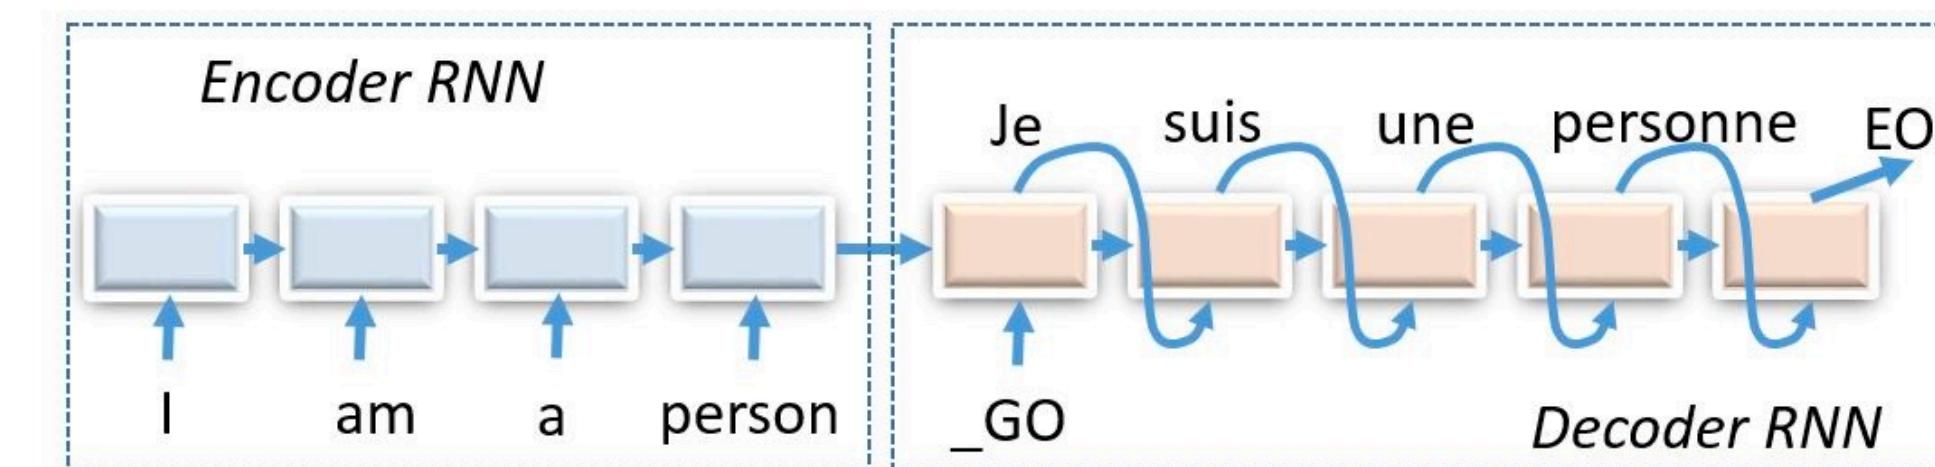


CNN: Encoder/decoder

- Strategy to learn “representations” in lower dimension.
- Allows to work with problem where input/output have the same size (finally segmentation!).

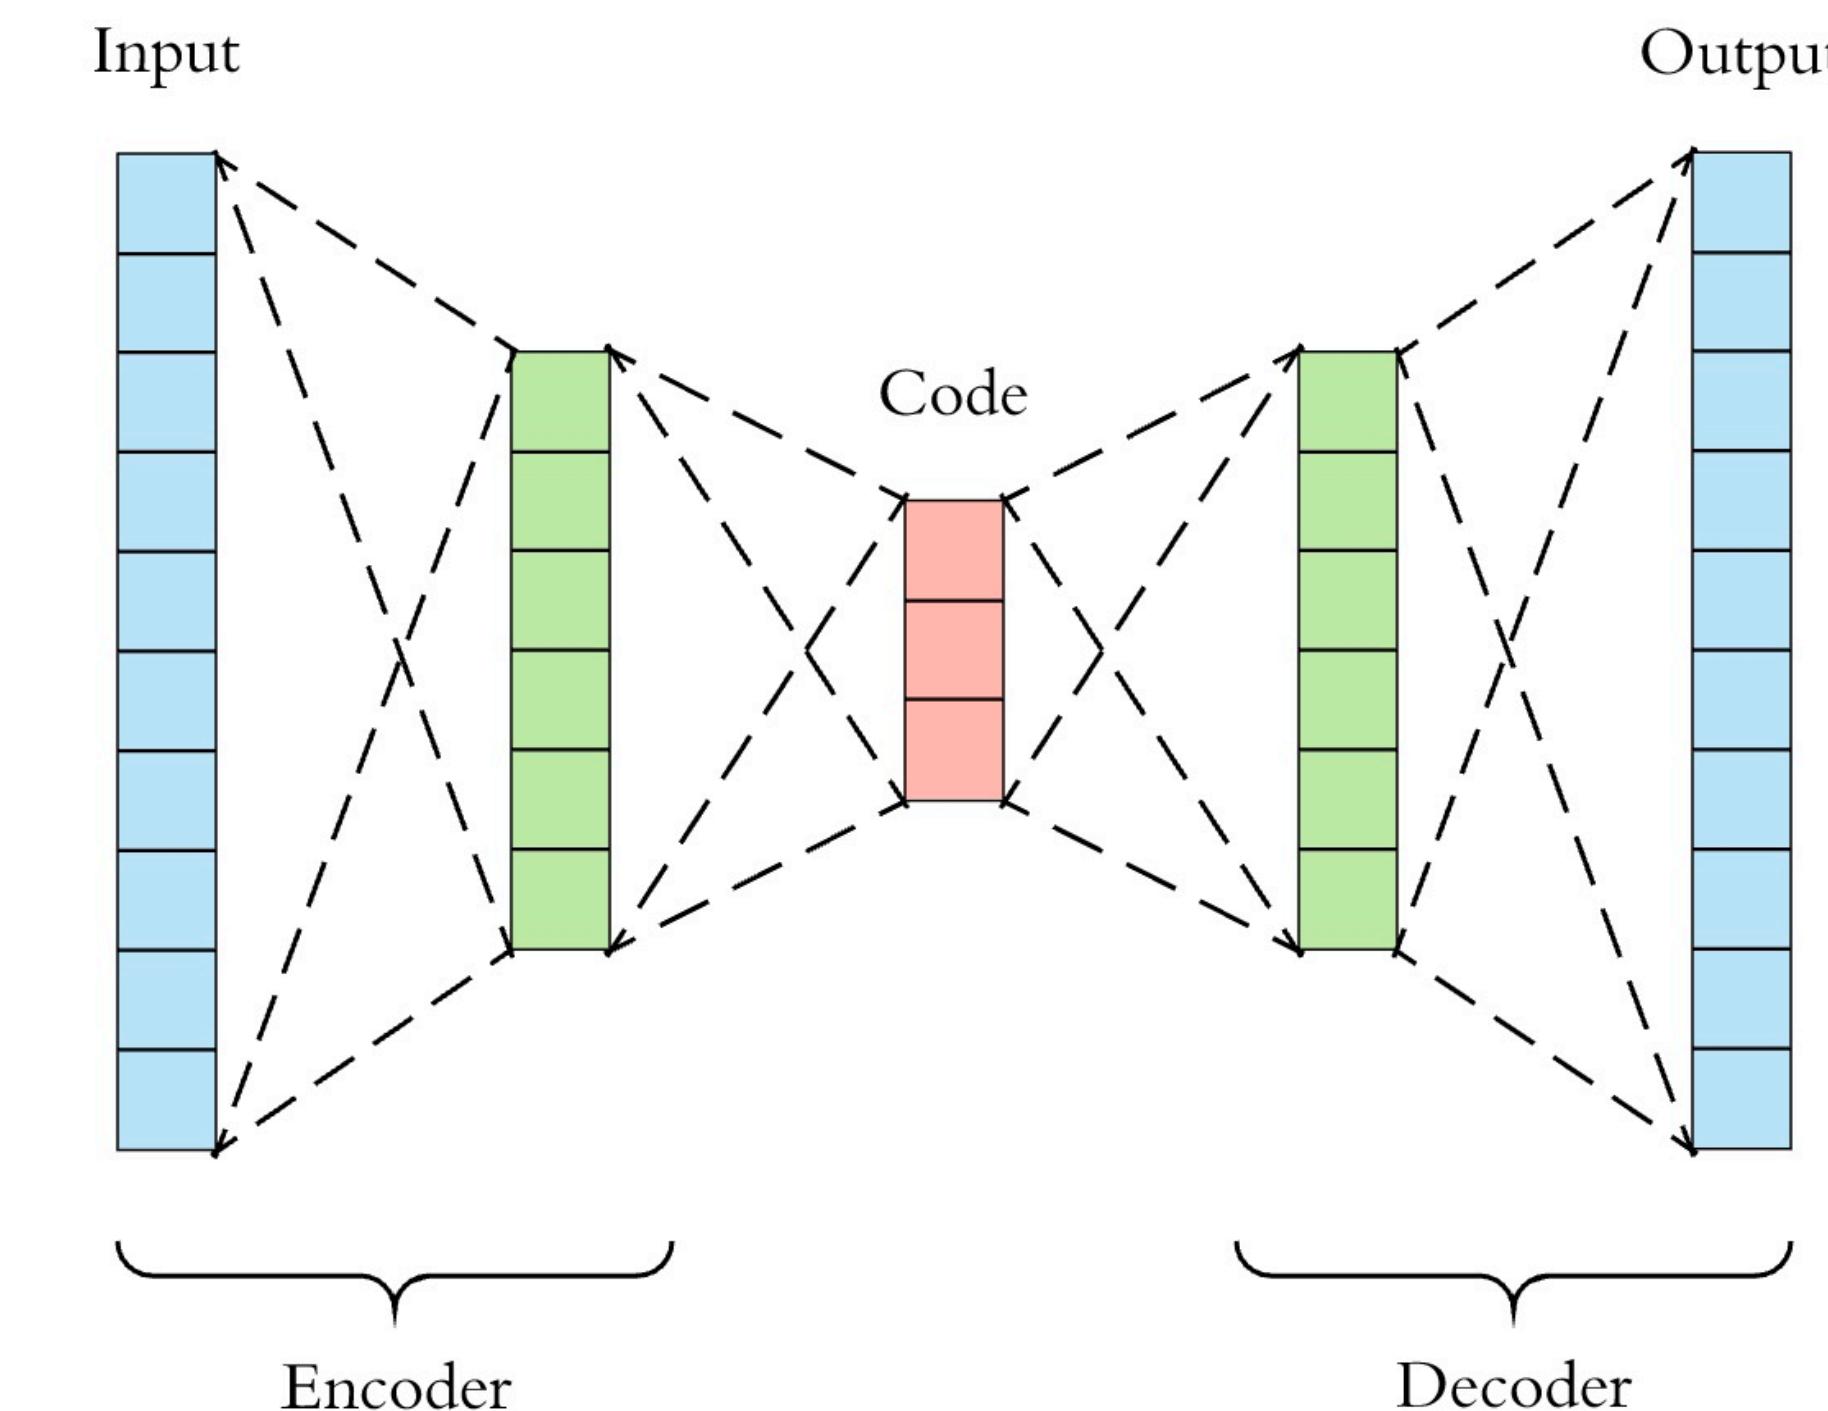
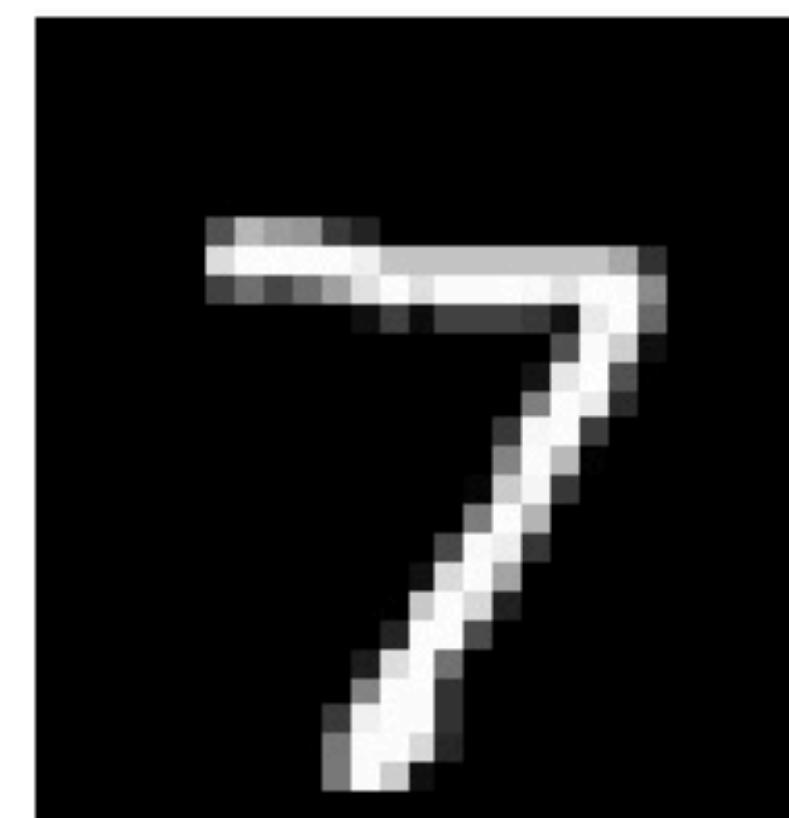


Input Code Output



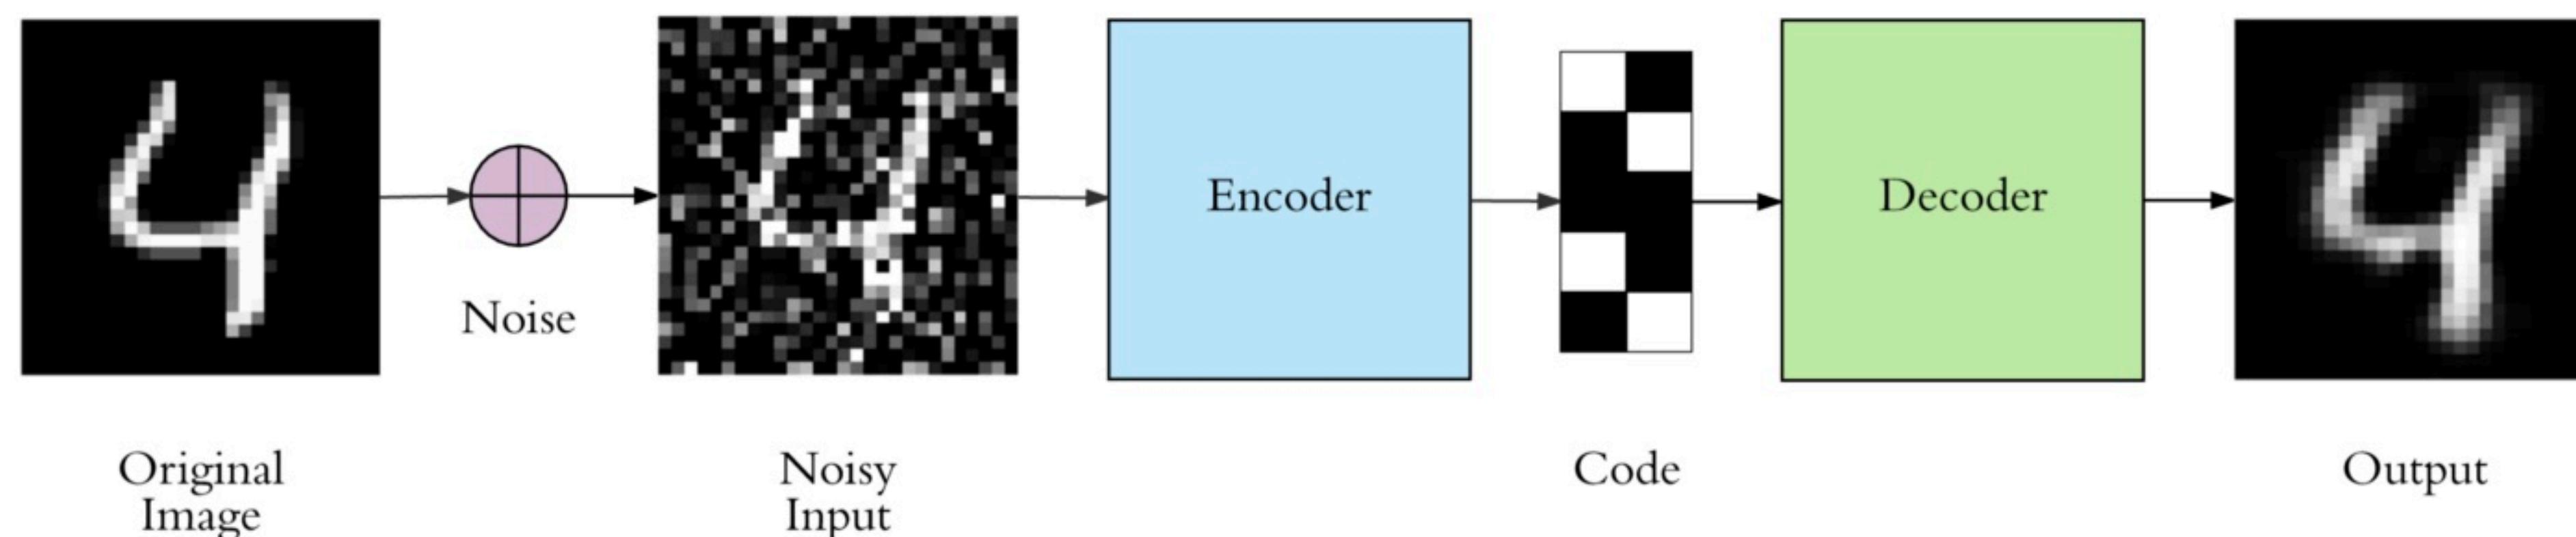
CNN: Encoder/decoder trainning

- A simple Encoder/decoder with FC layers.
- We want to recognize handwritten digits.
- How to train?



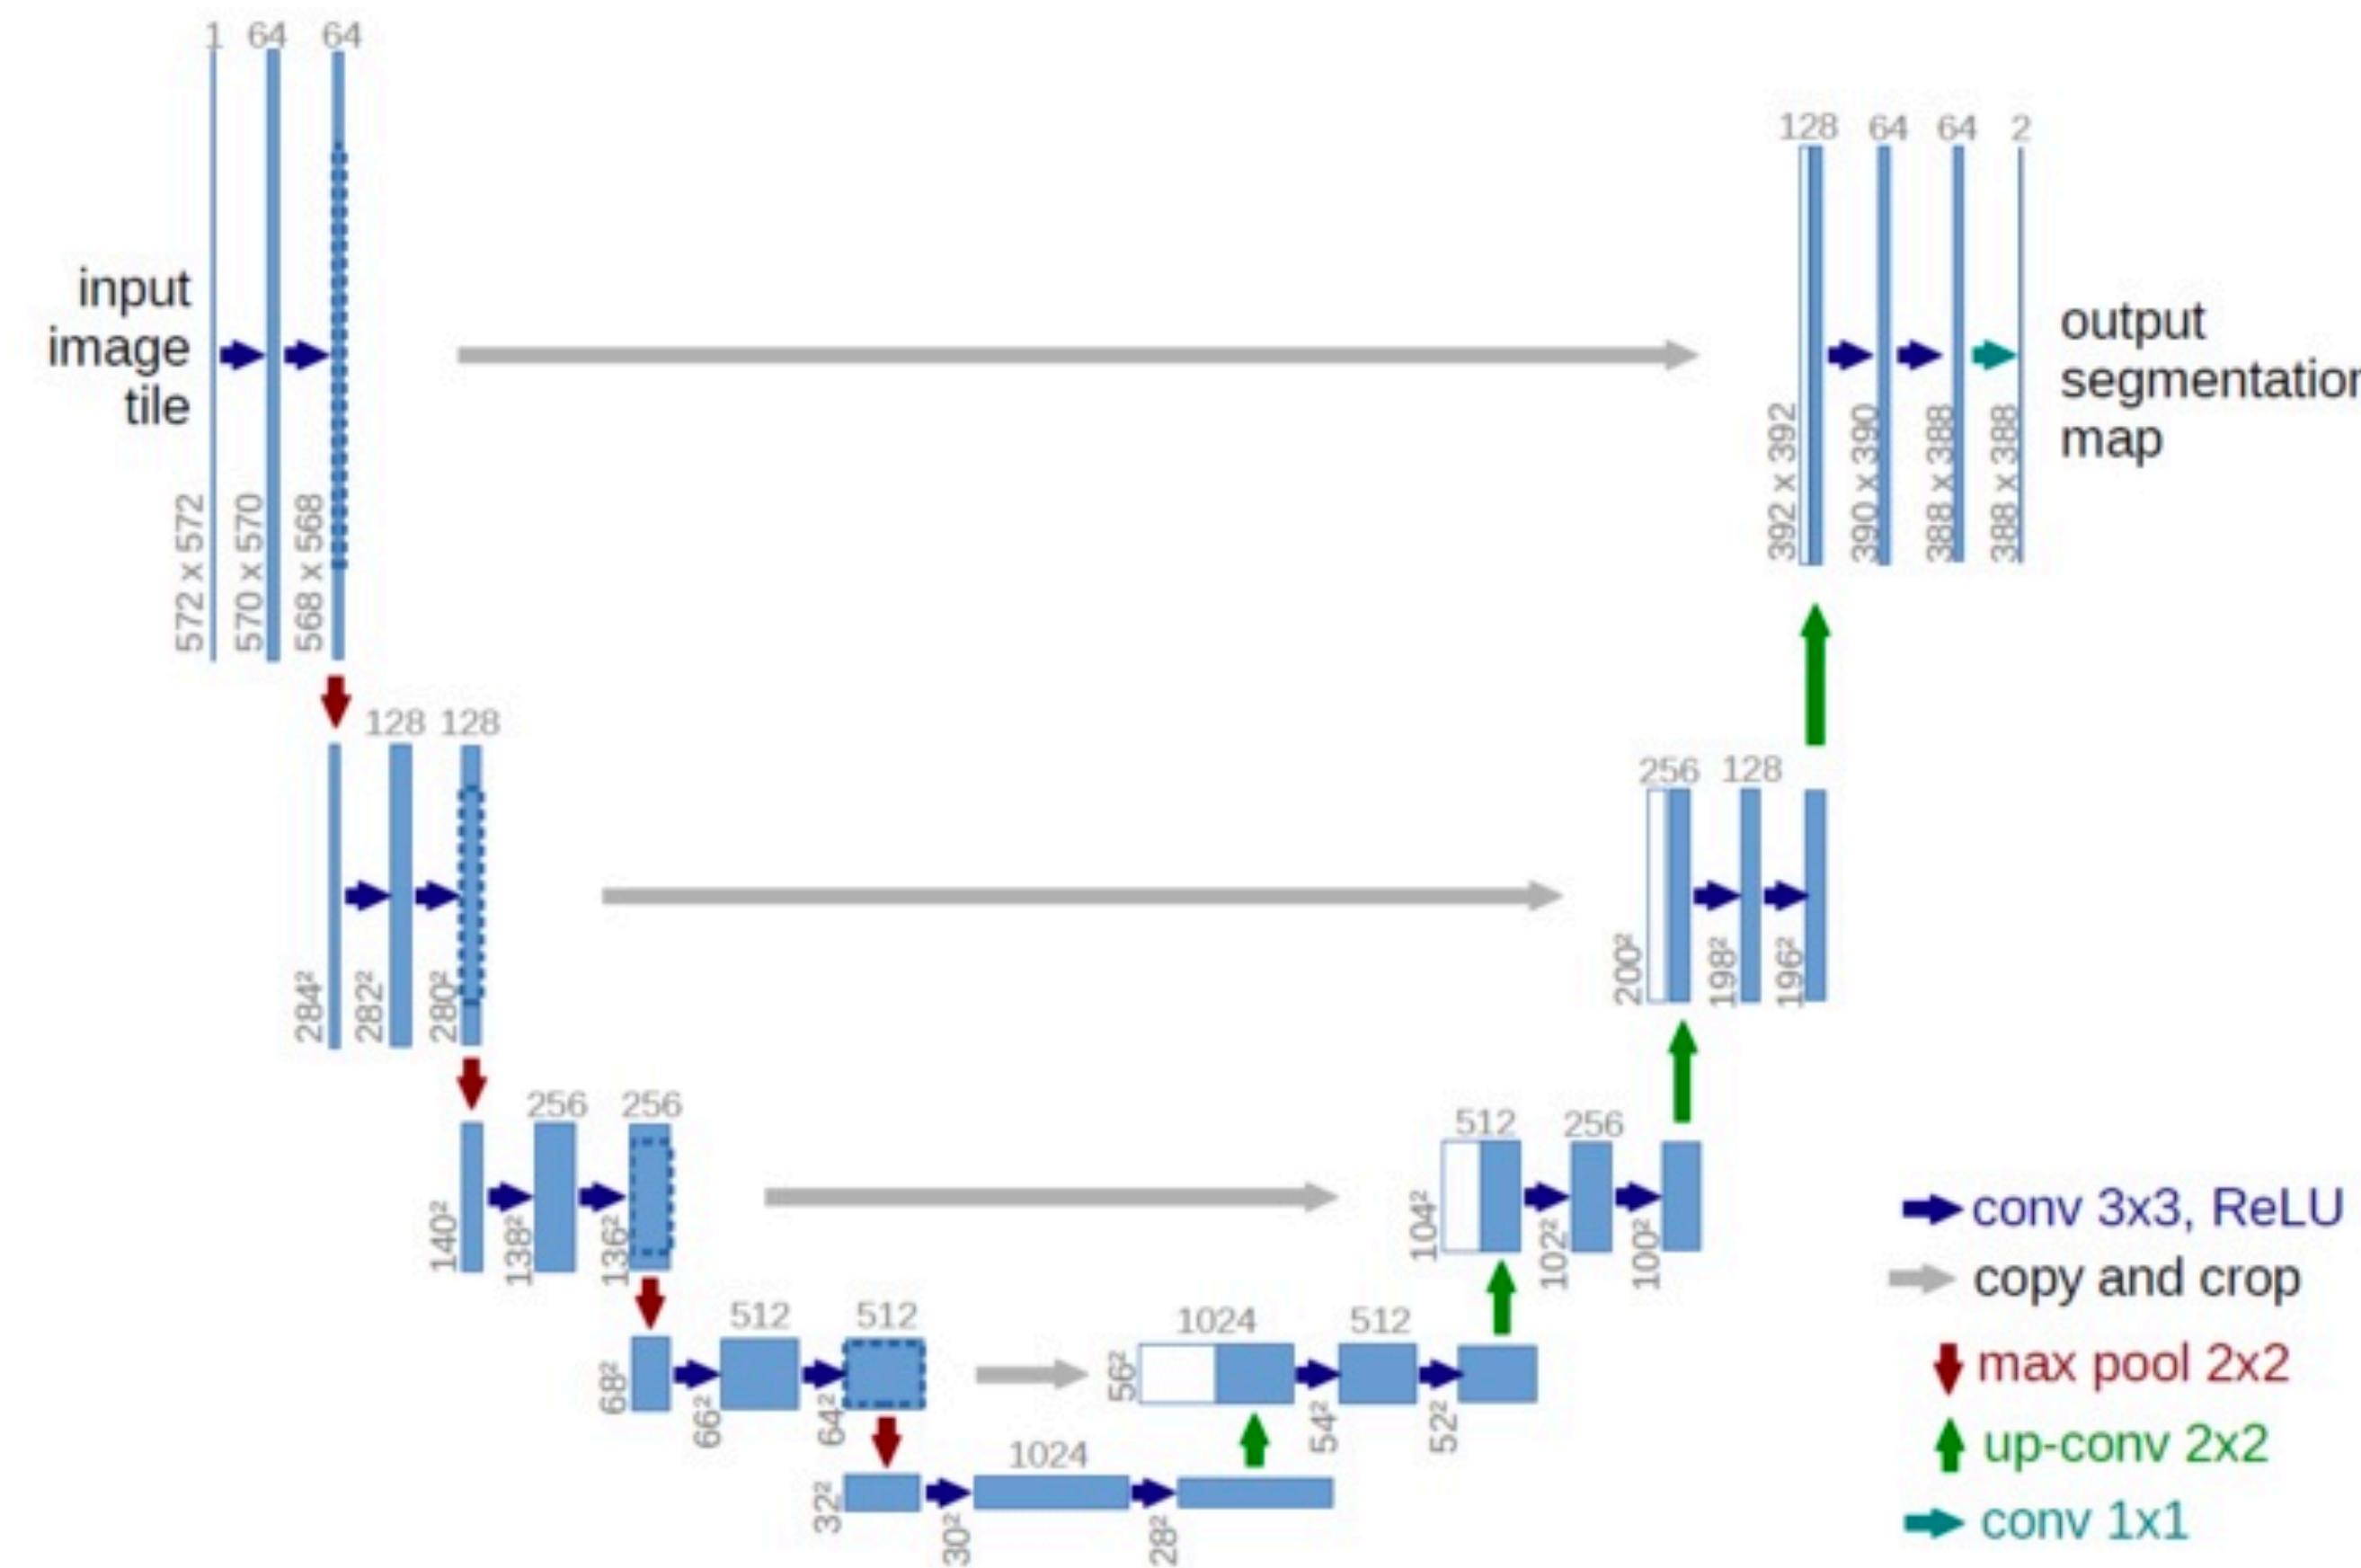
CNN: Encoder/decoder segmentation

- After training, we understand concepts and decode.



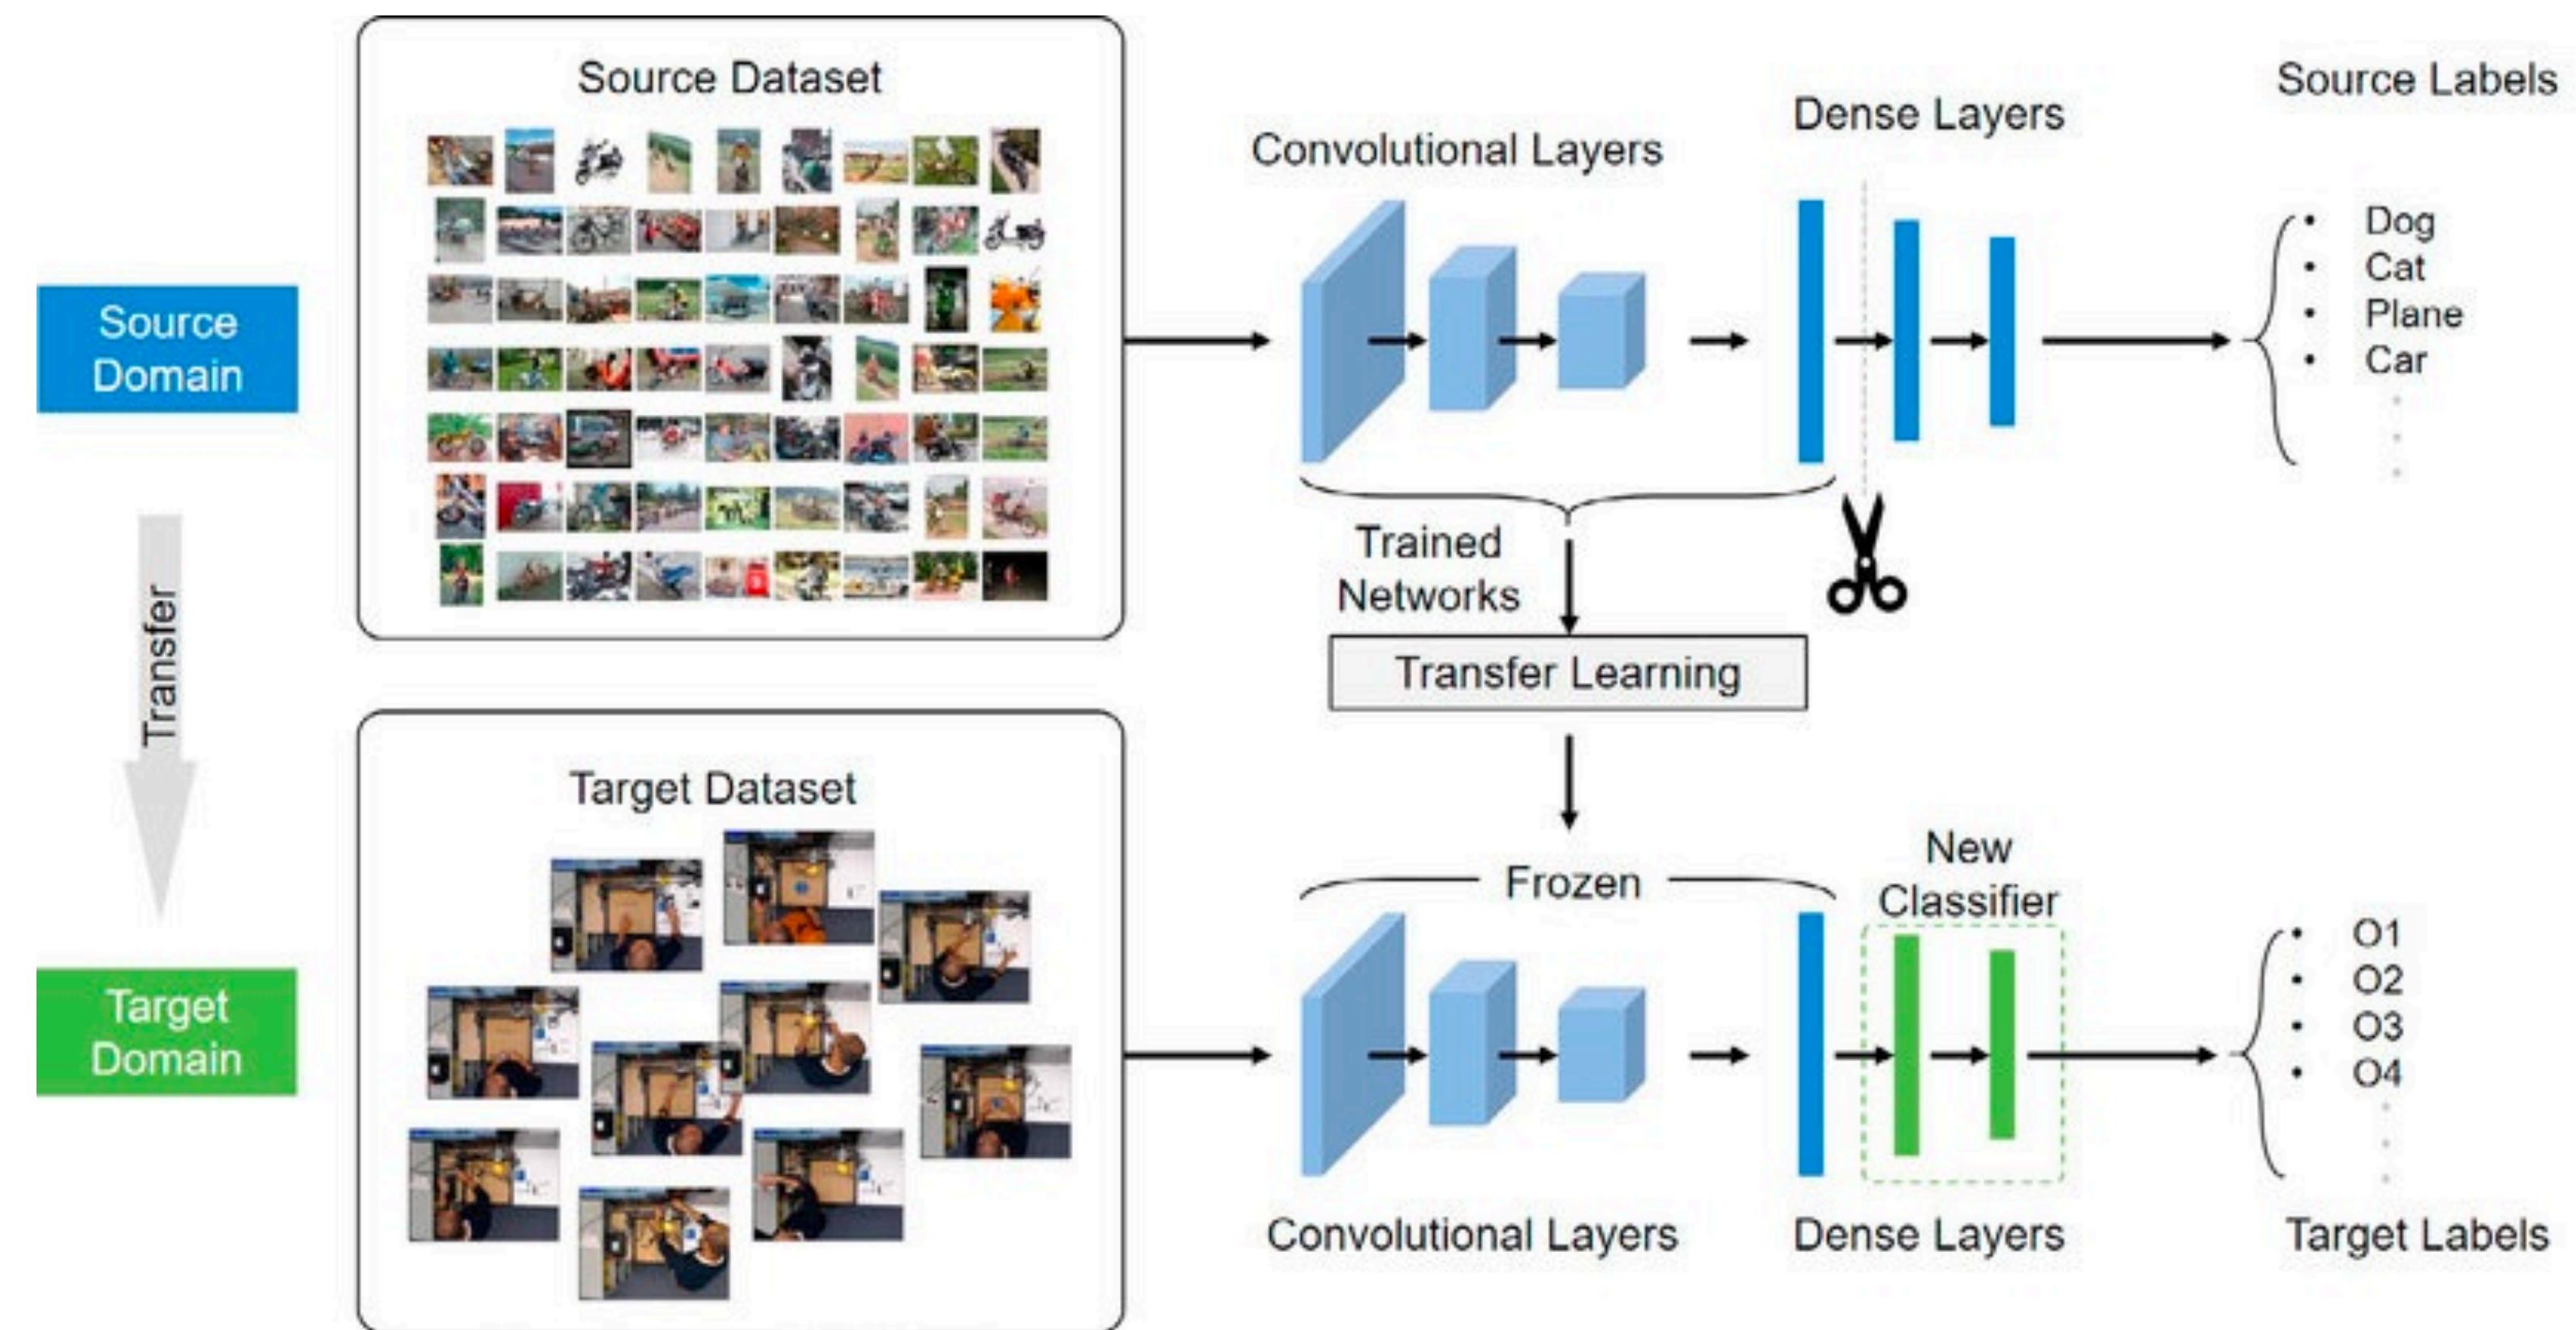
CNN: Encoder/decoder segmentation

- *Ronneberger et al 2015* proposed encoder/decoder for segmentation using convolutional layers (U-Net).
 - Very popular in biomedical applications.



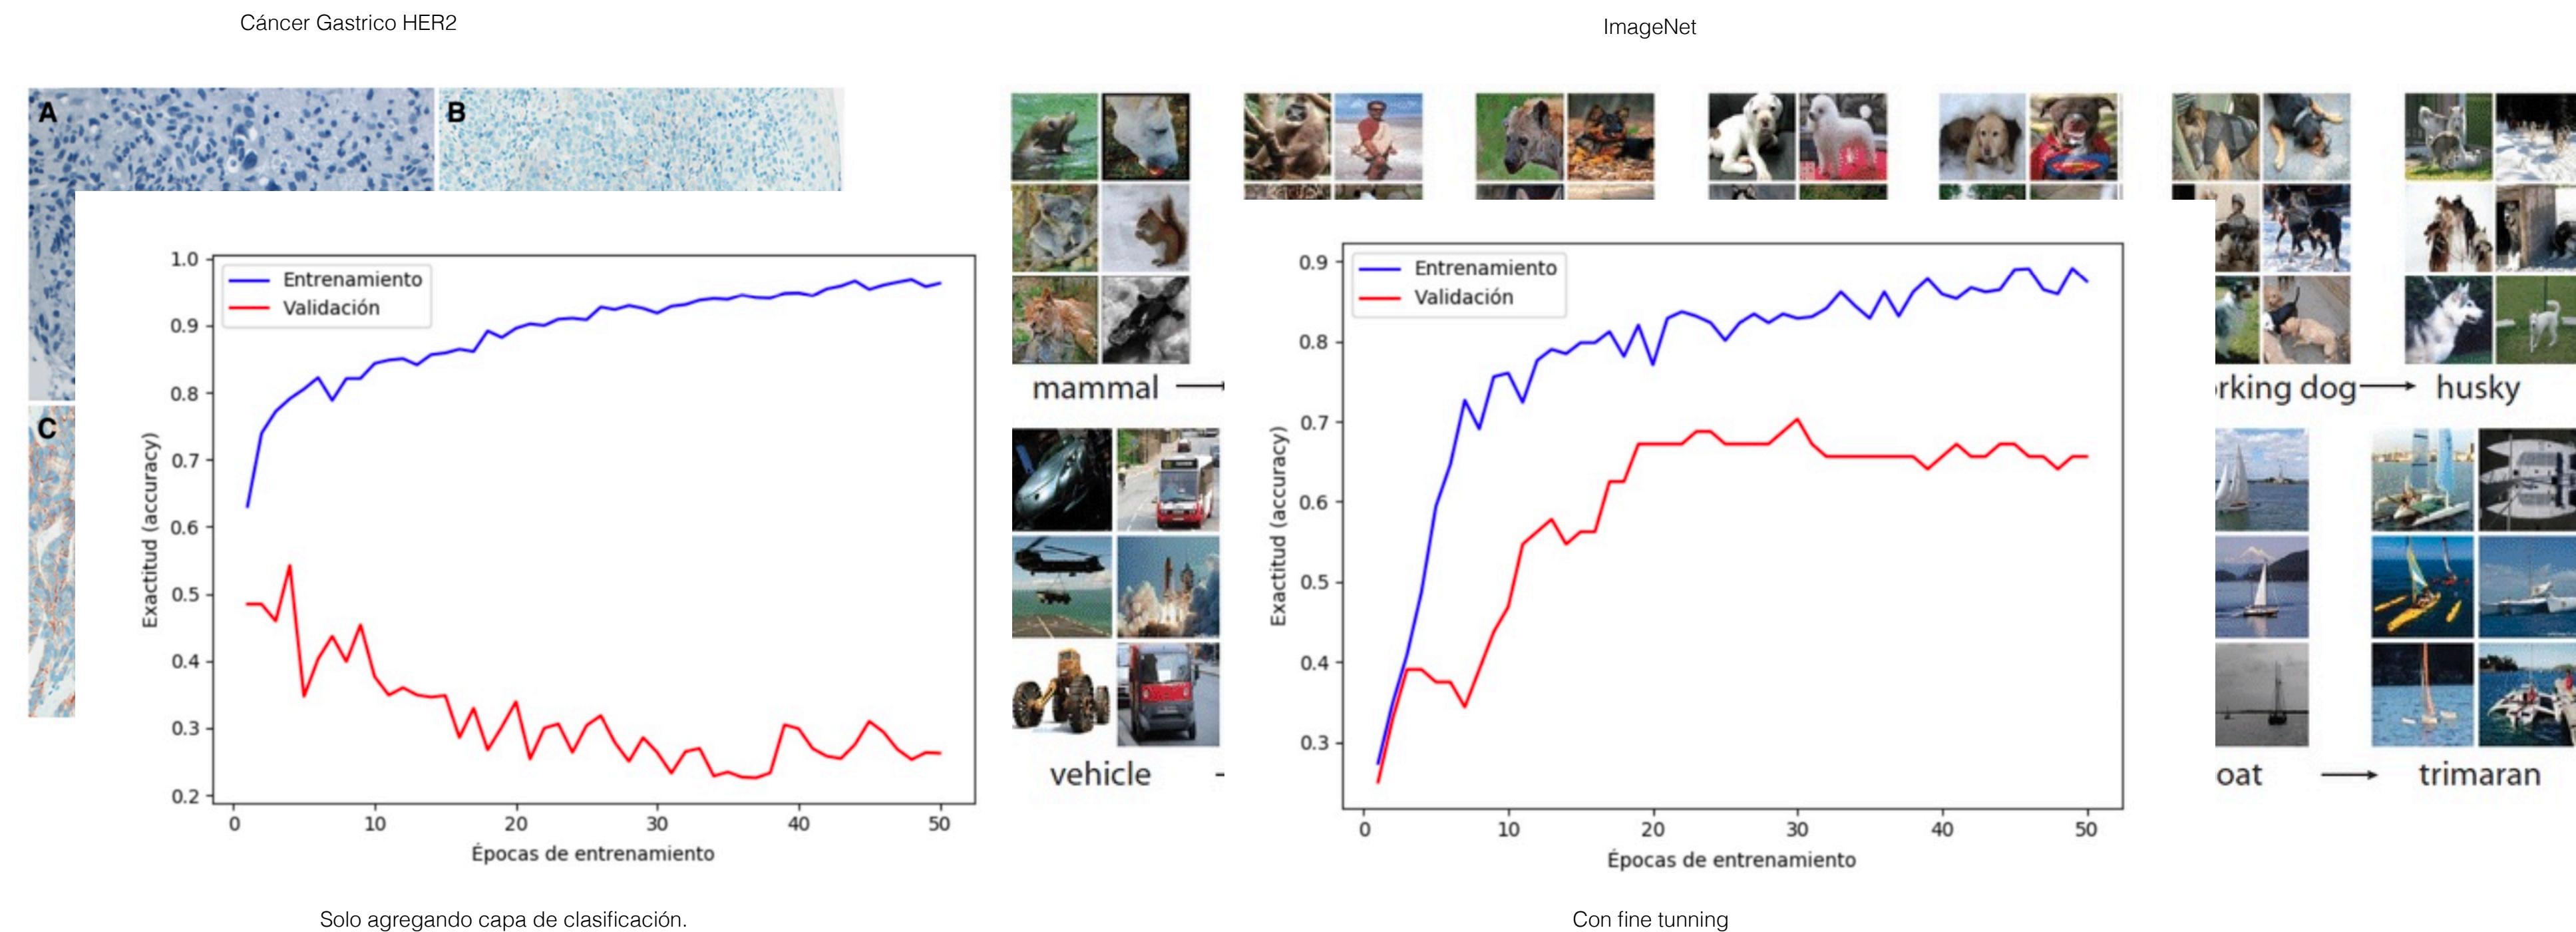
CNN: Transfer learning

- Training a CNN requires many images ($\sim 10^6$).
- To transfer learning implies to train in one domain and to partially train into another.



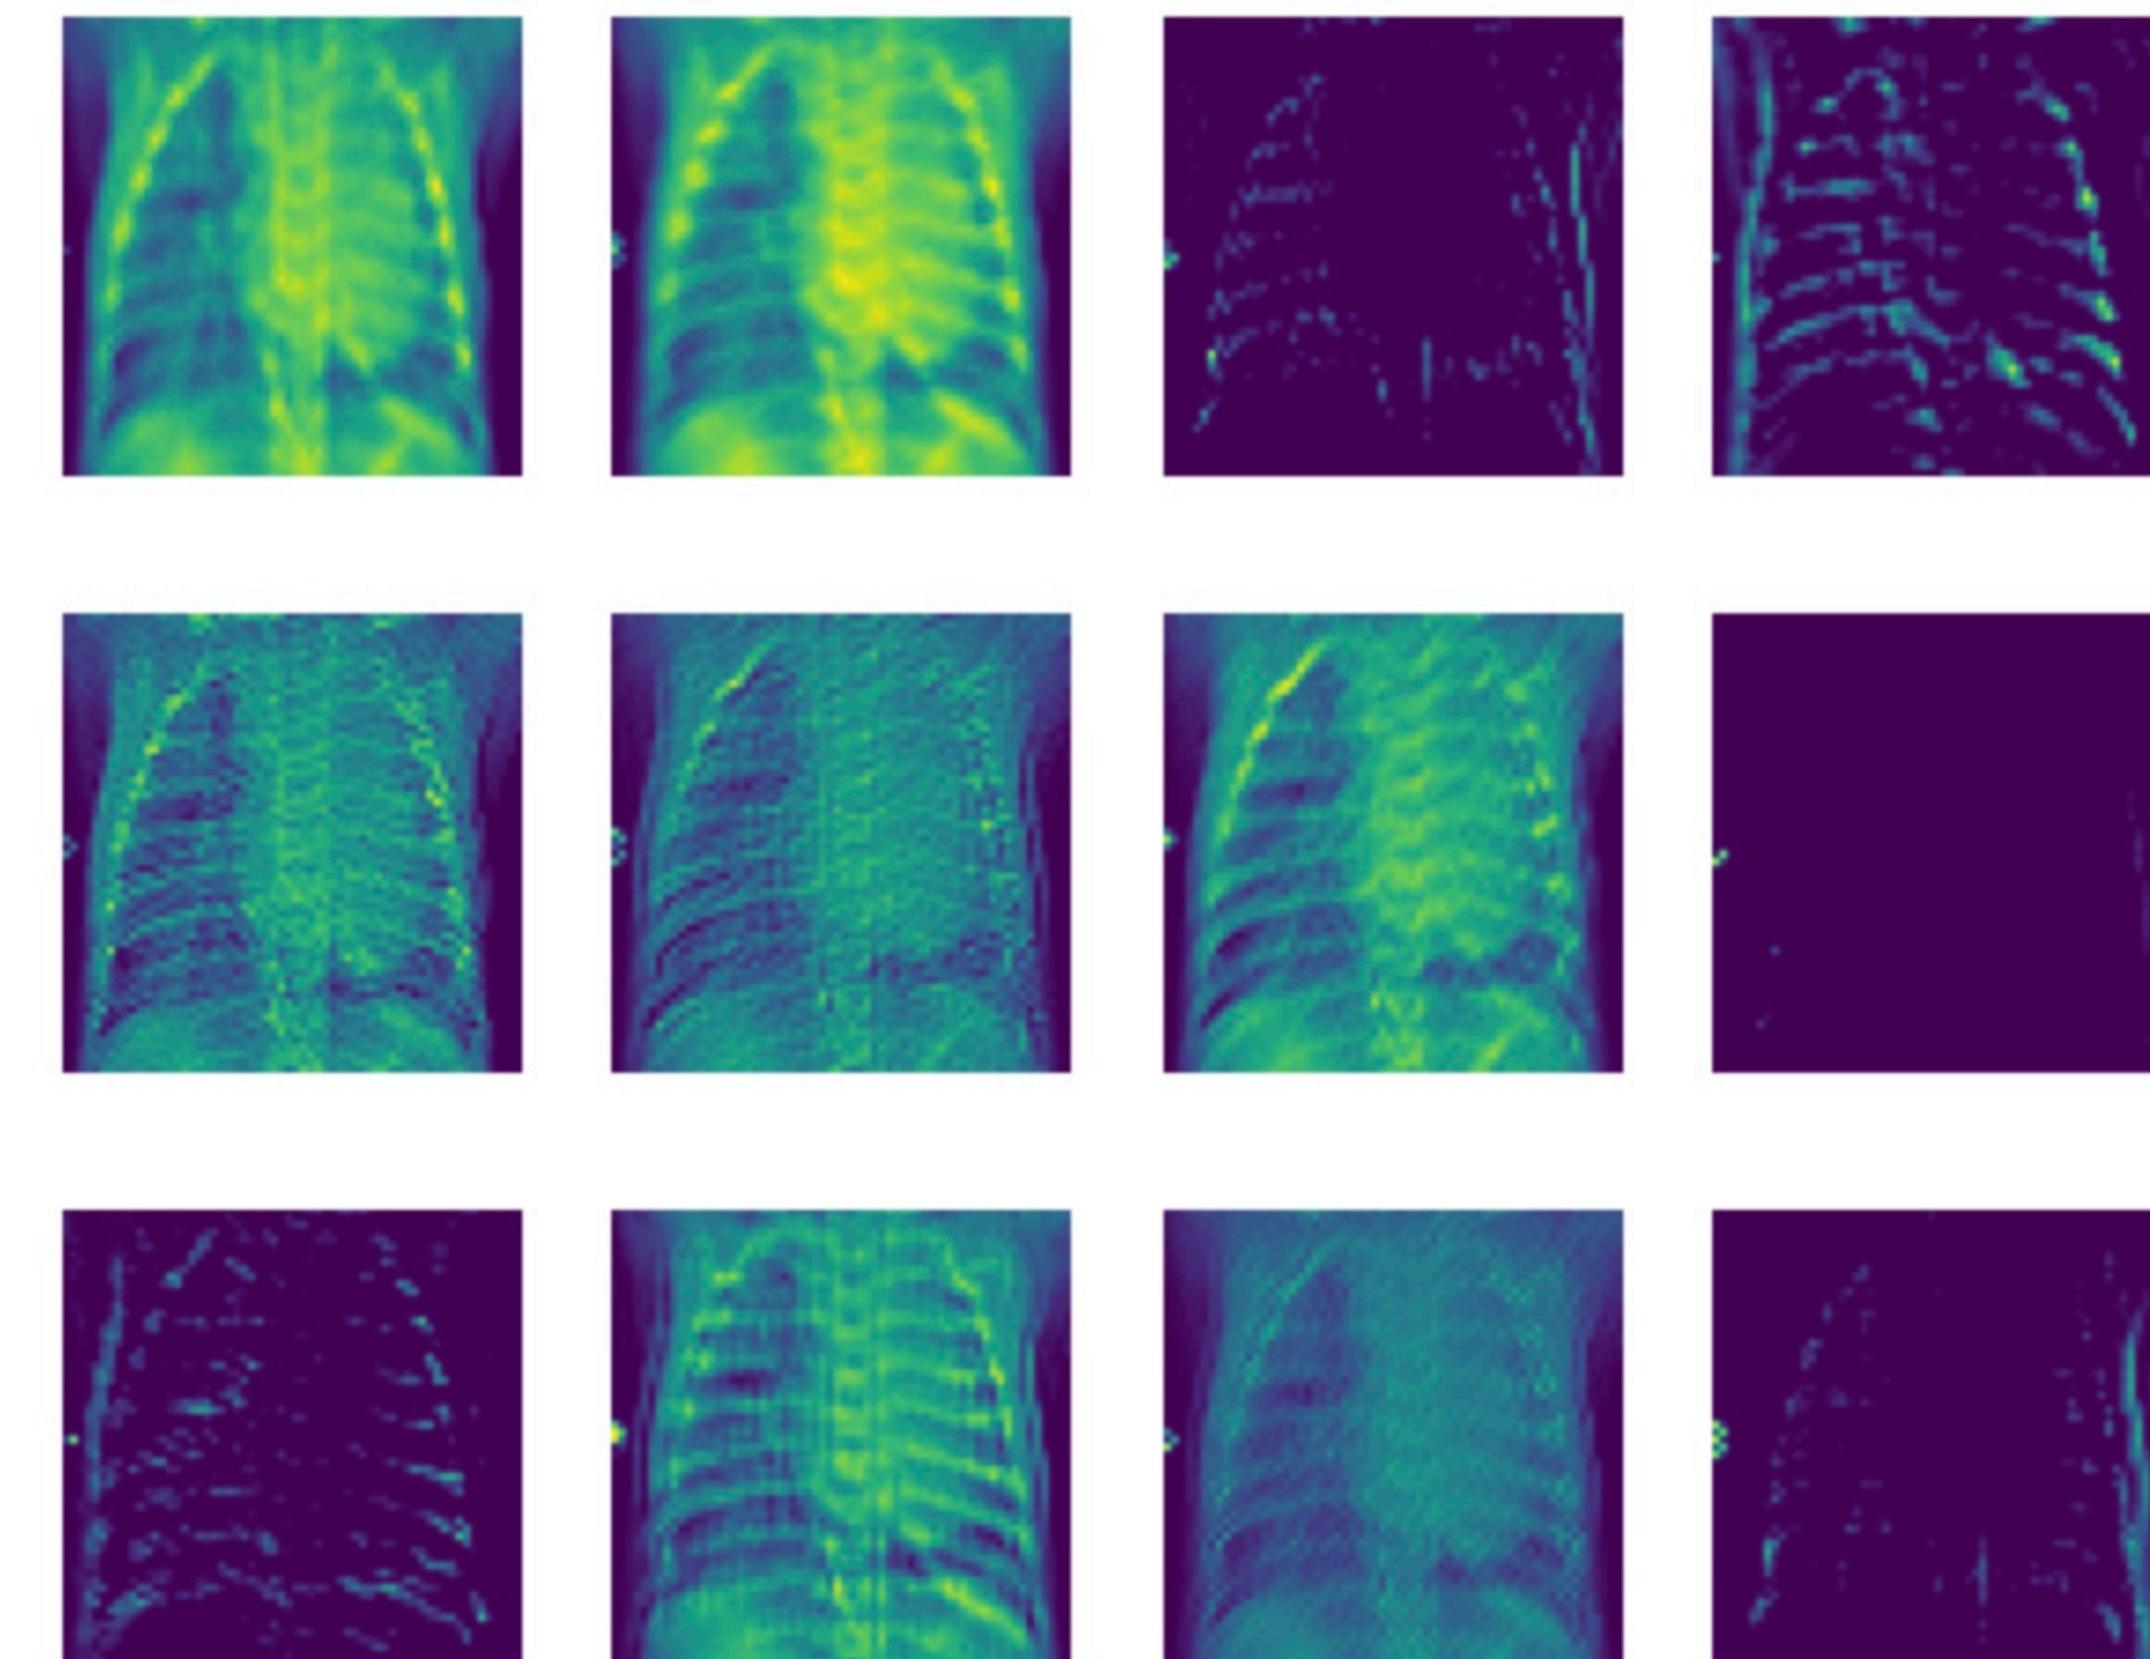
CNN: Transfer learning

- Carefully evaluate distance between domains to define: what to transfer, what to refine.
- Example. Gastric Cancer using InceptionV3



CNN visualization: activation

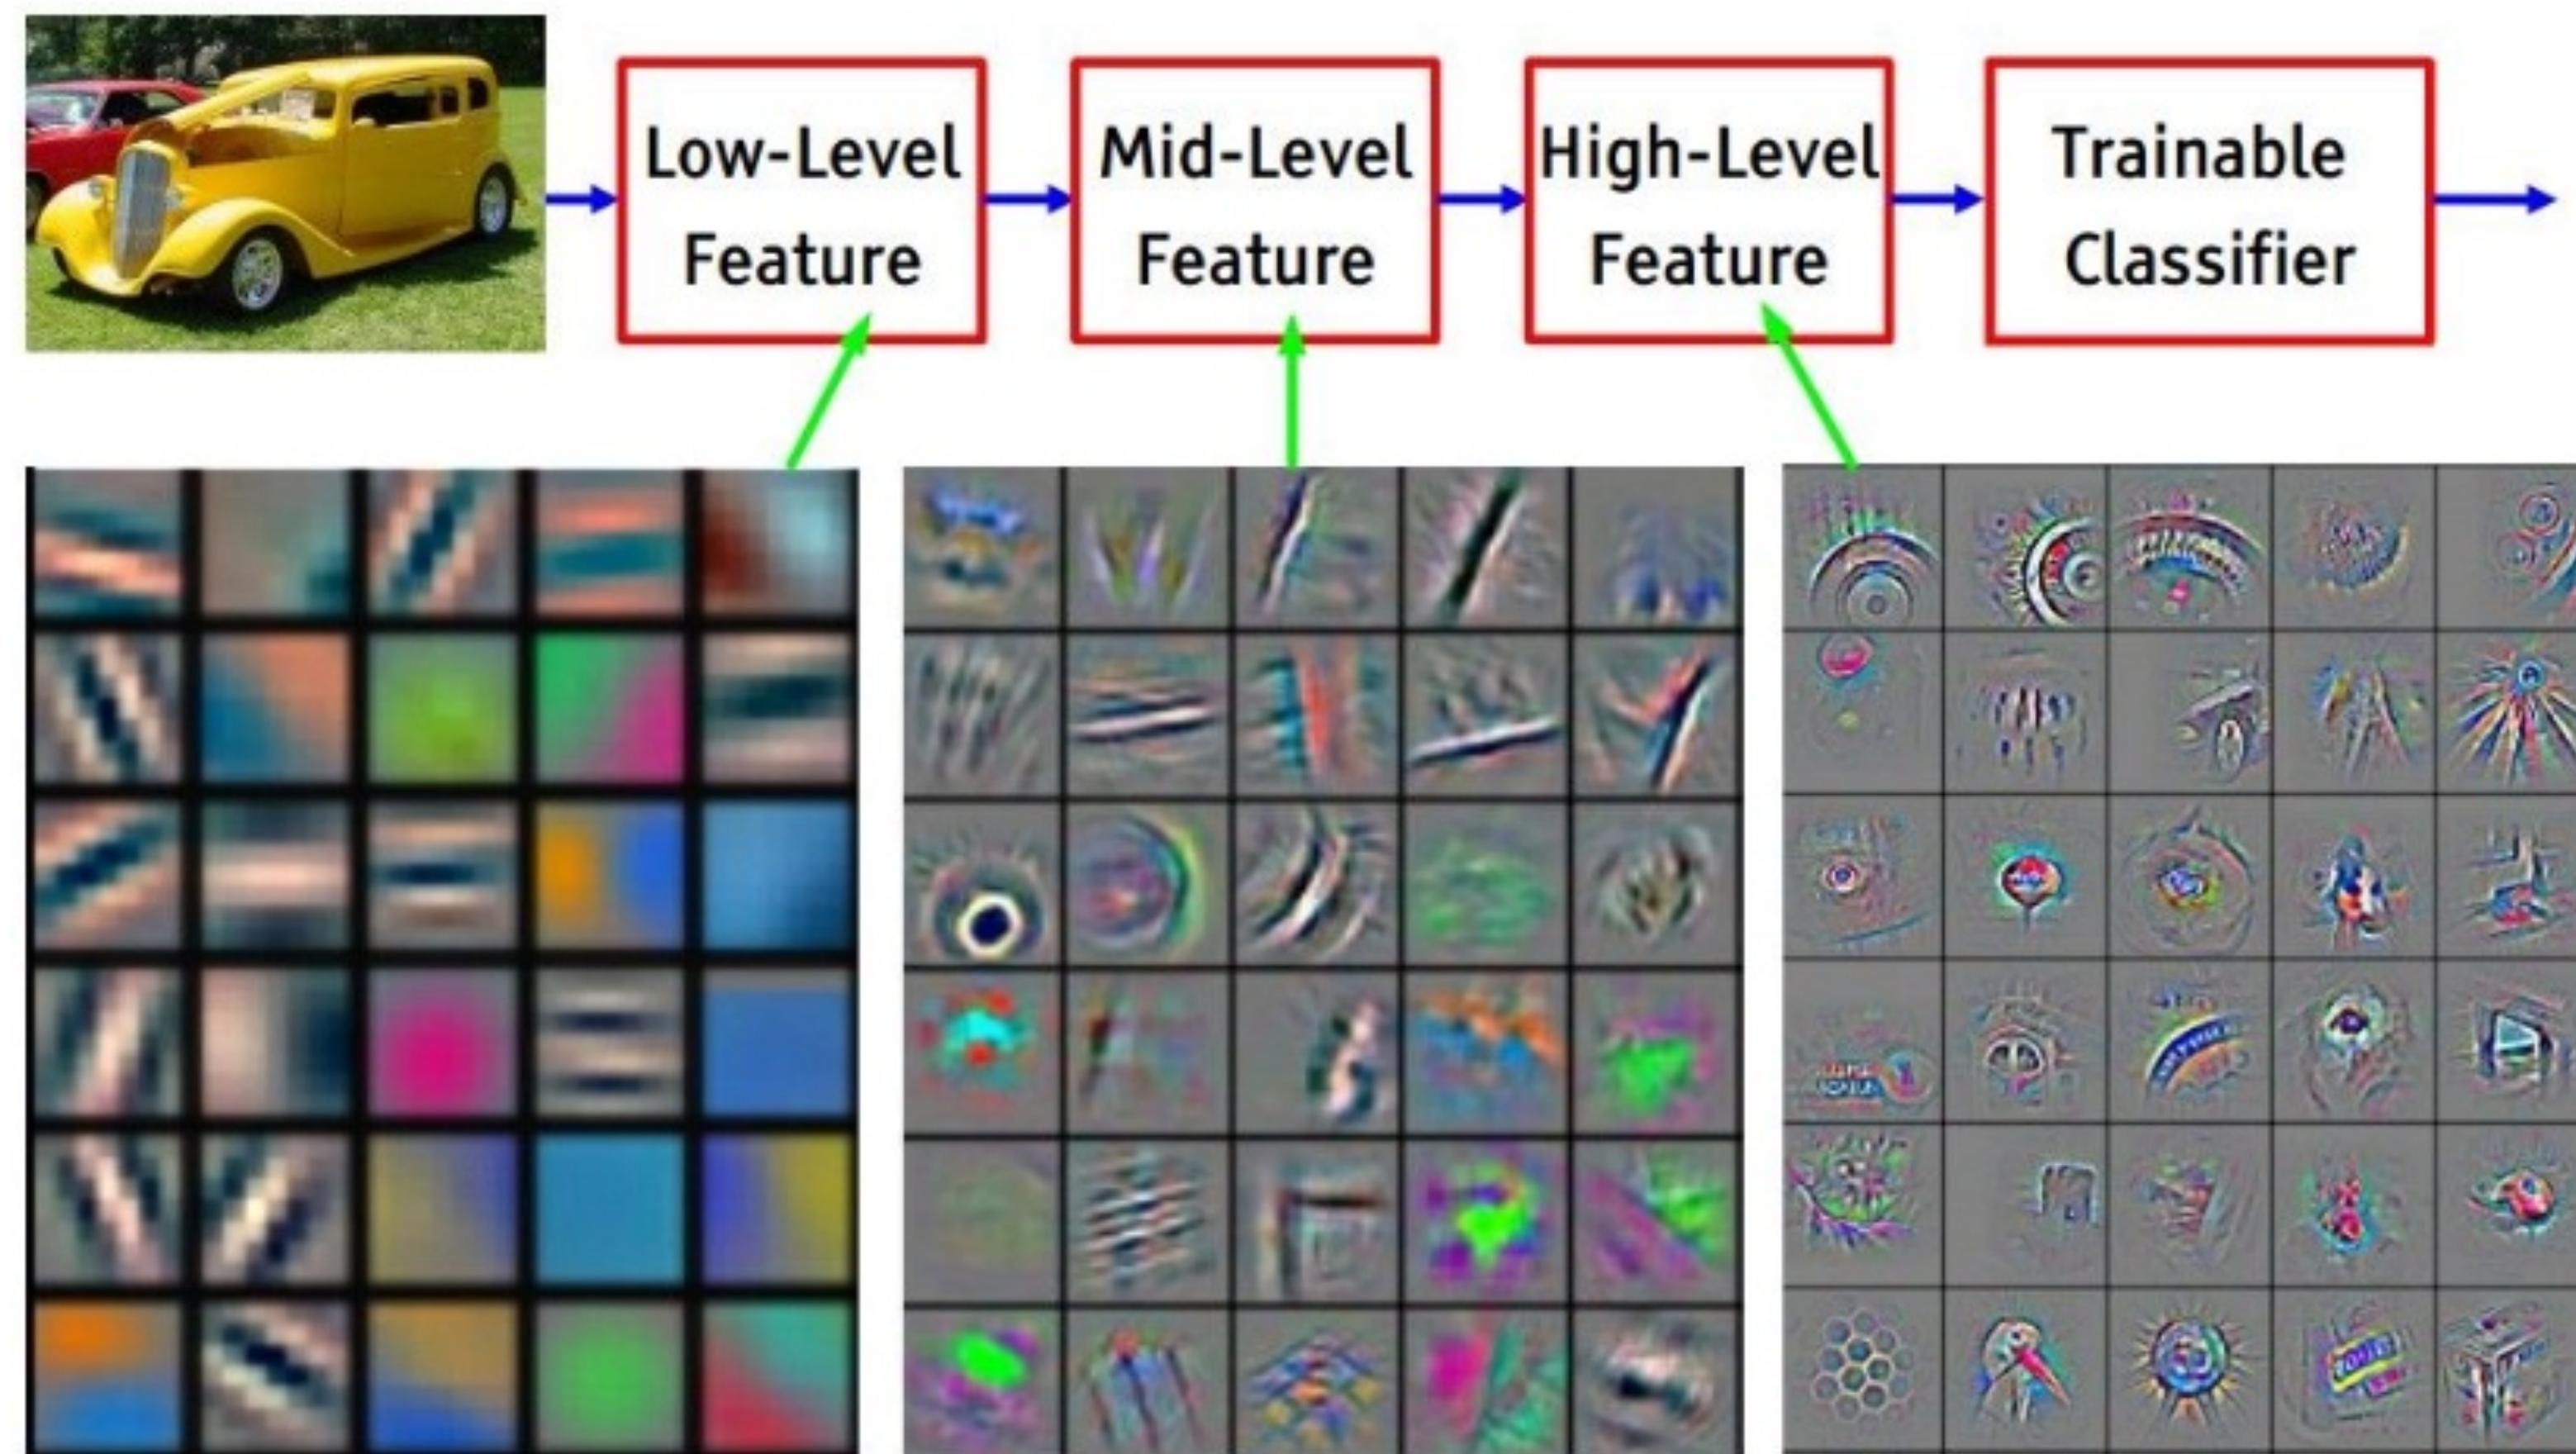
- How the network activates during inference?
- Remember we can visualize as filters!



Set of convolutions in test architecture (first layer)

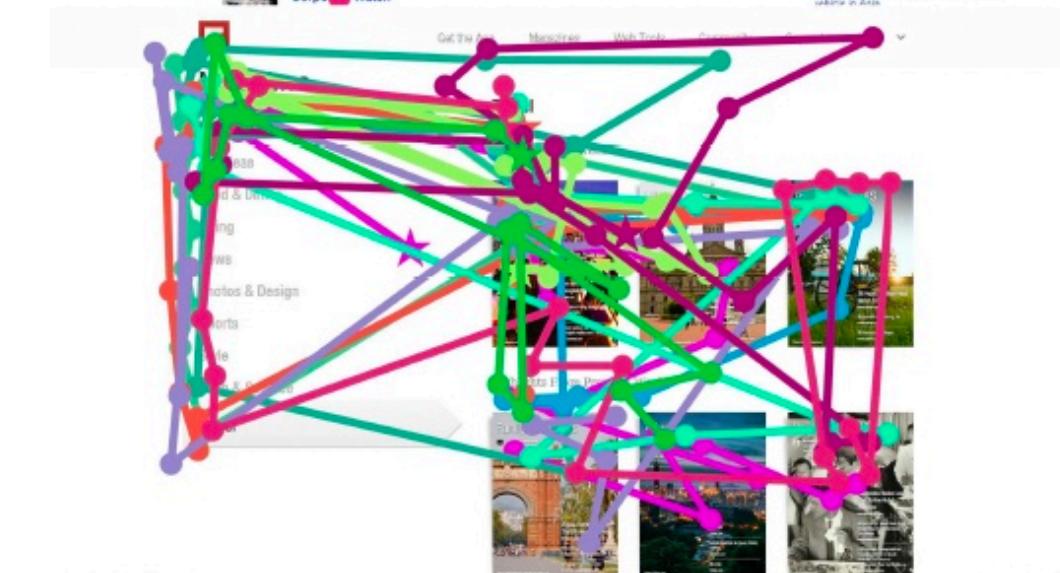
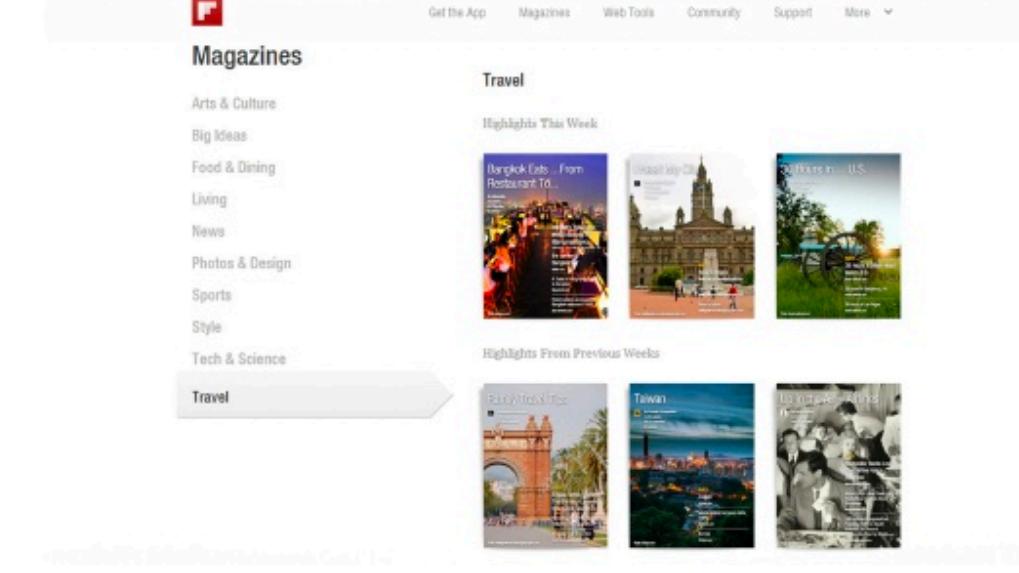
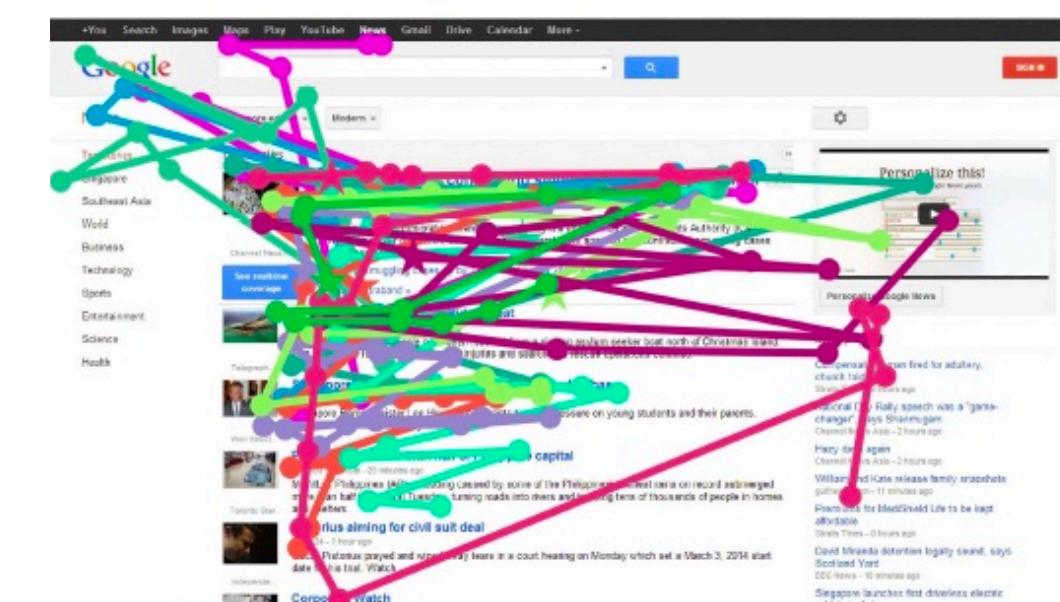
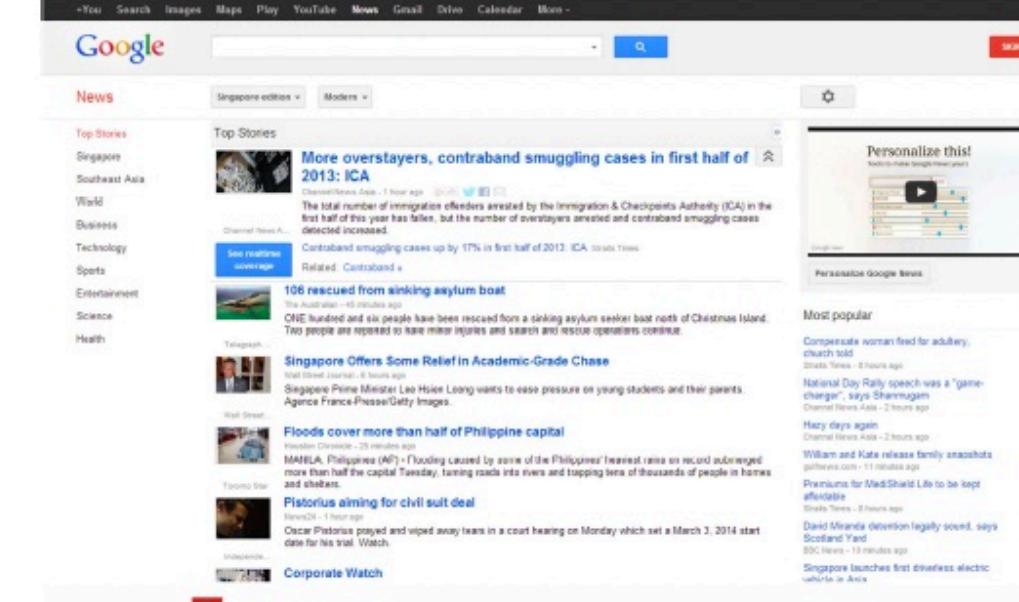
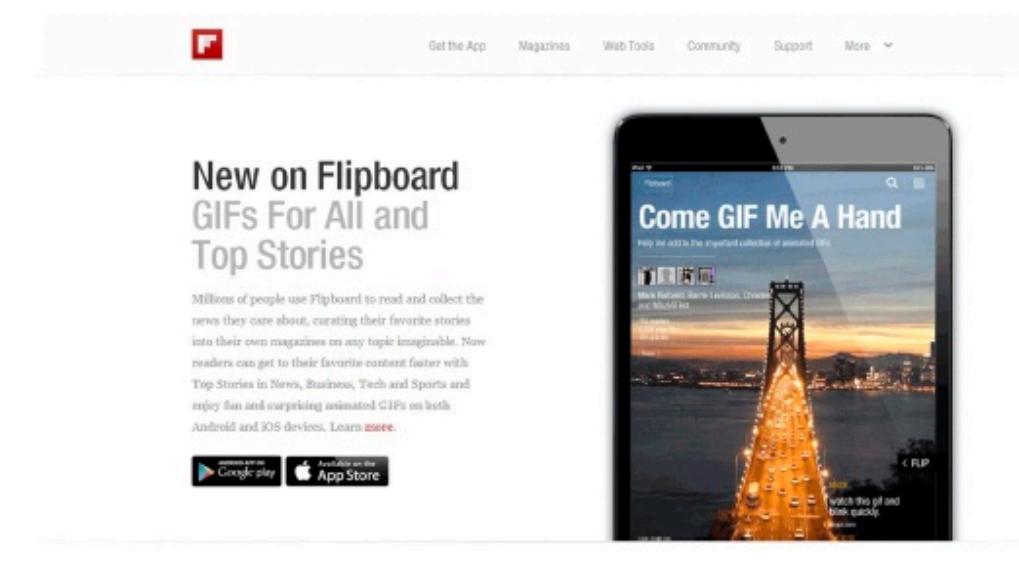
CNN visualization: kernels

- How are the kernels the network learns? We can show kernels.



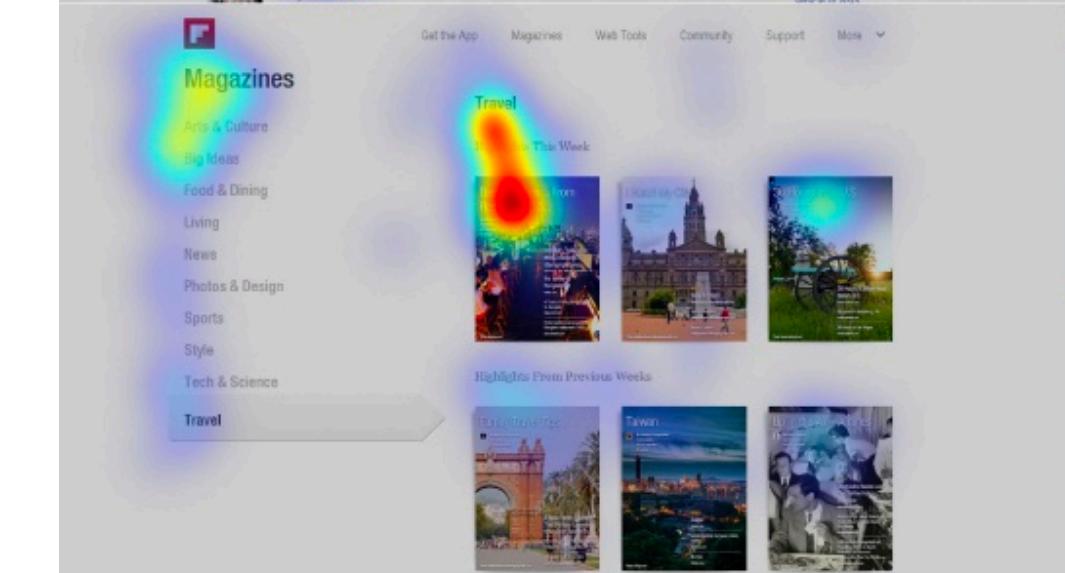
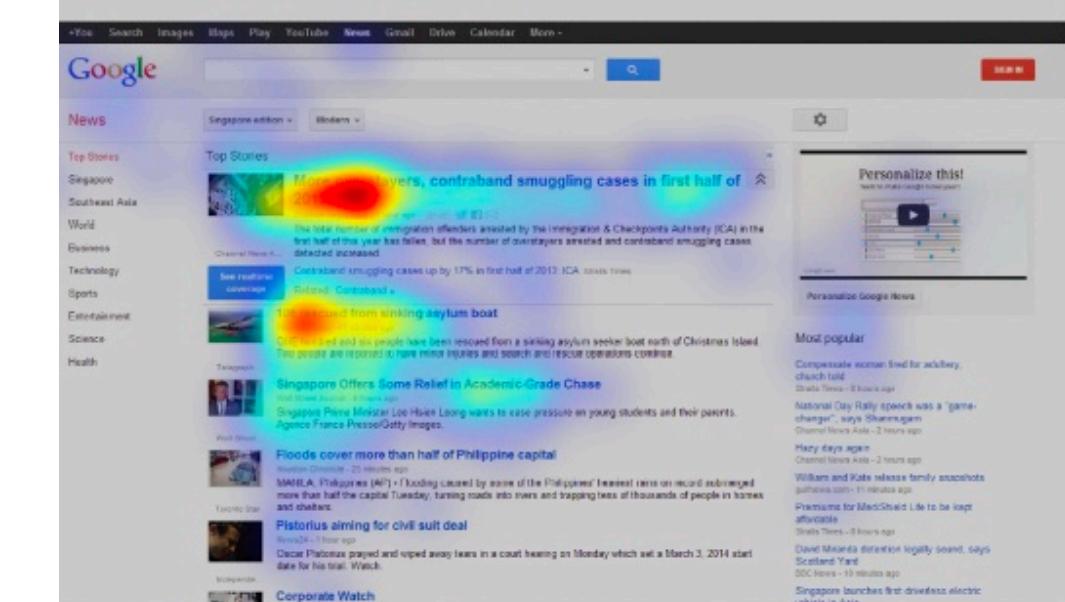
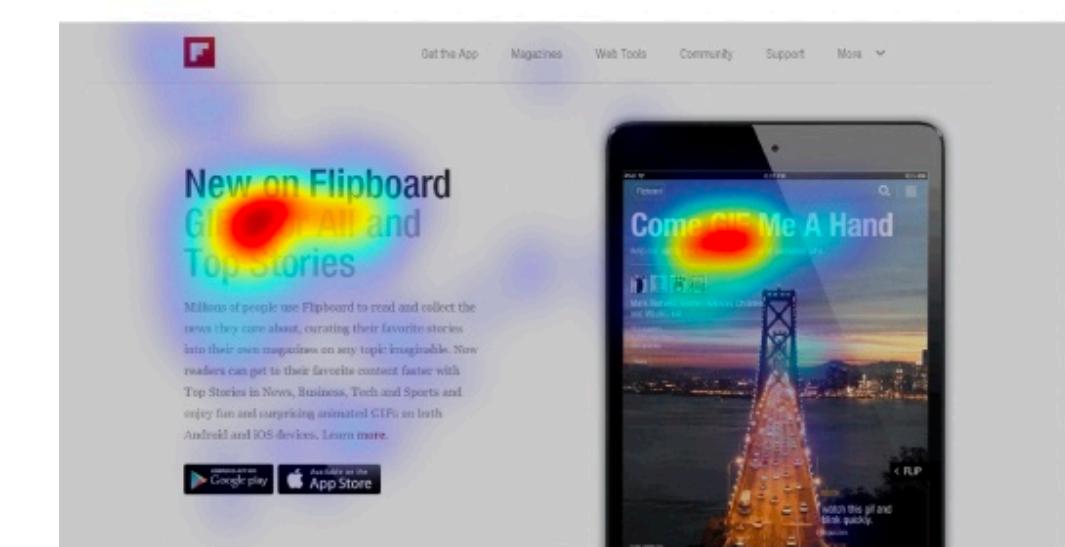
CNN visualization: saliency map

- In image processing models can predict where we will look at :)



Input

Saccadic patterns



Saccadic patterns heatmap

CNN visualization: saliency map



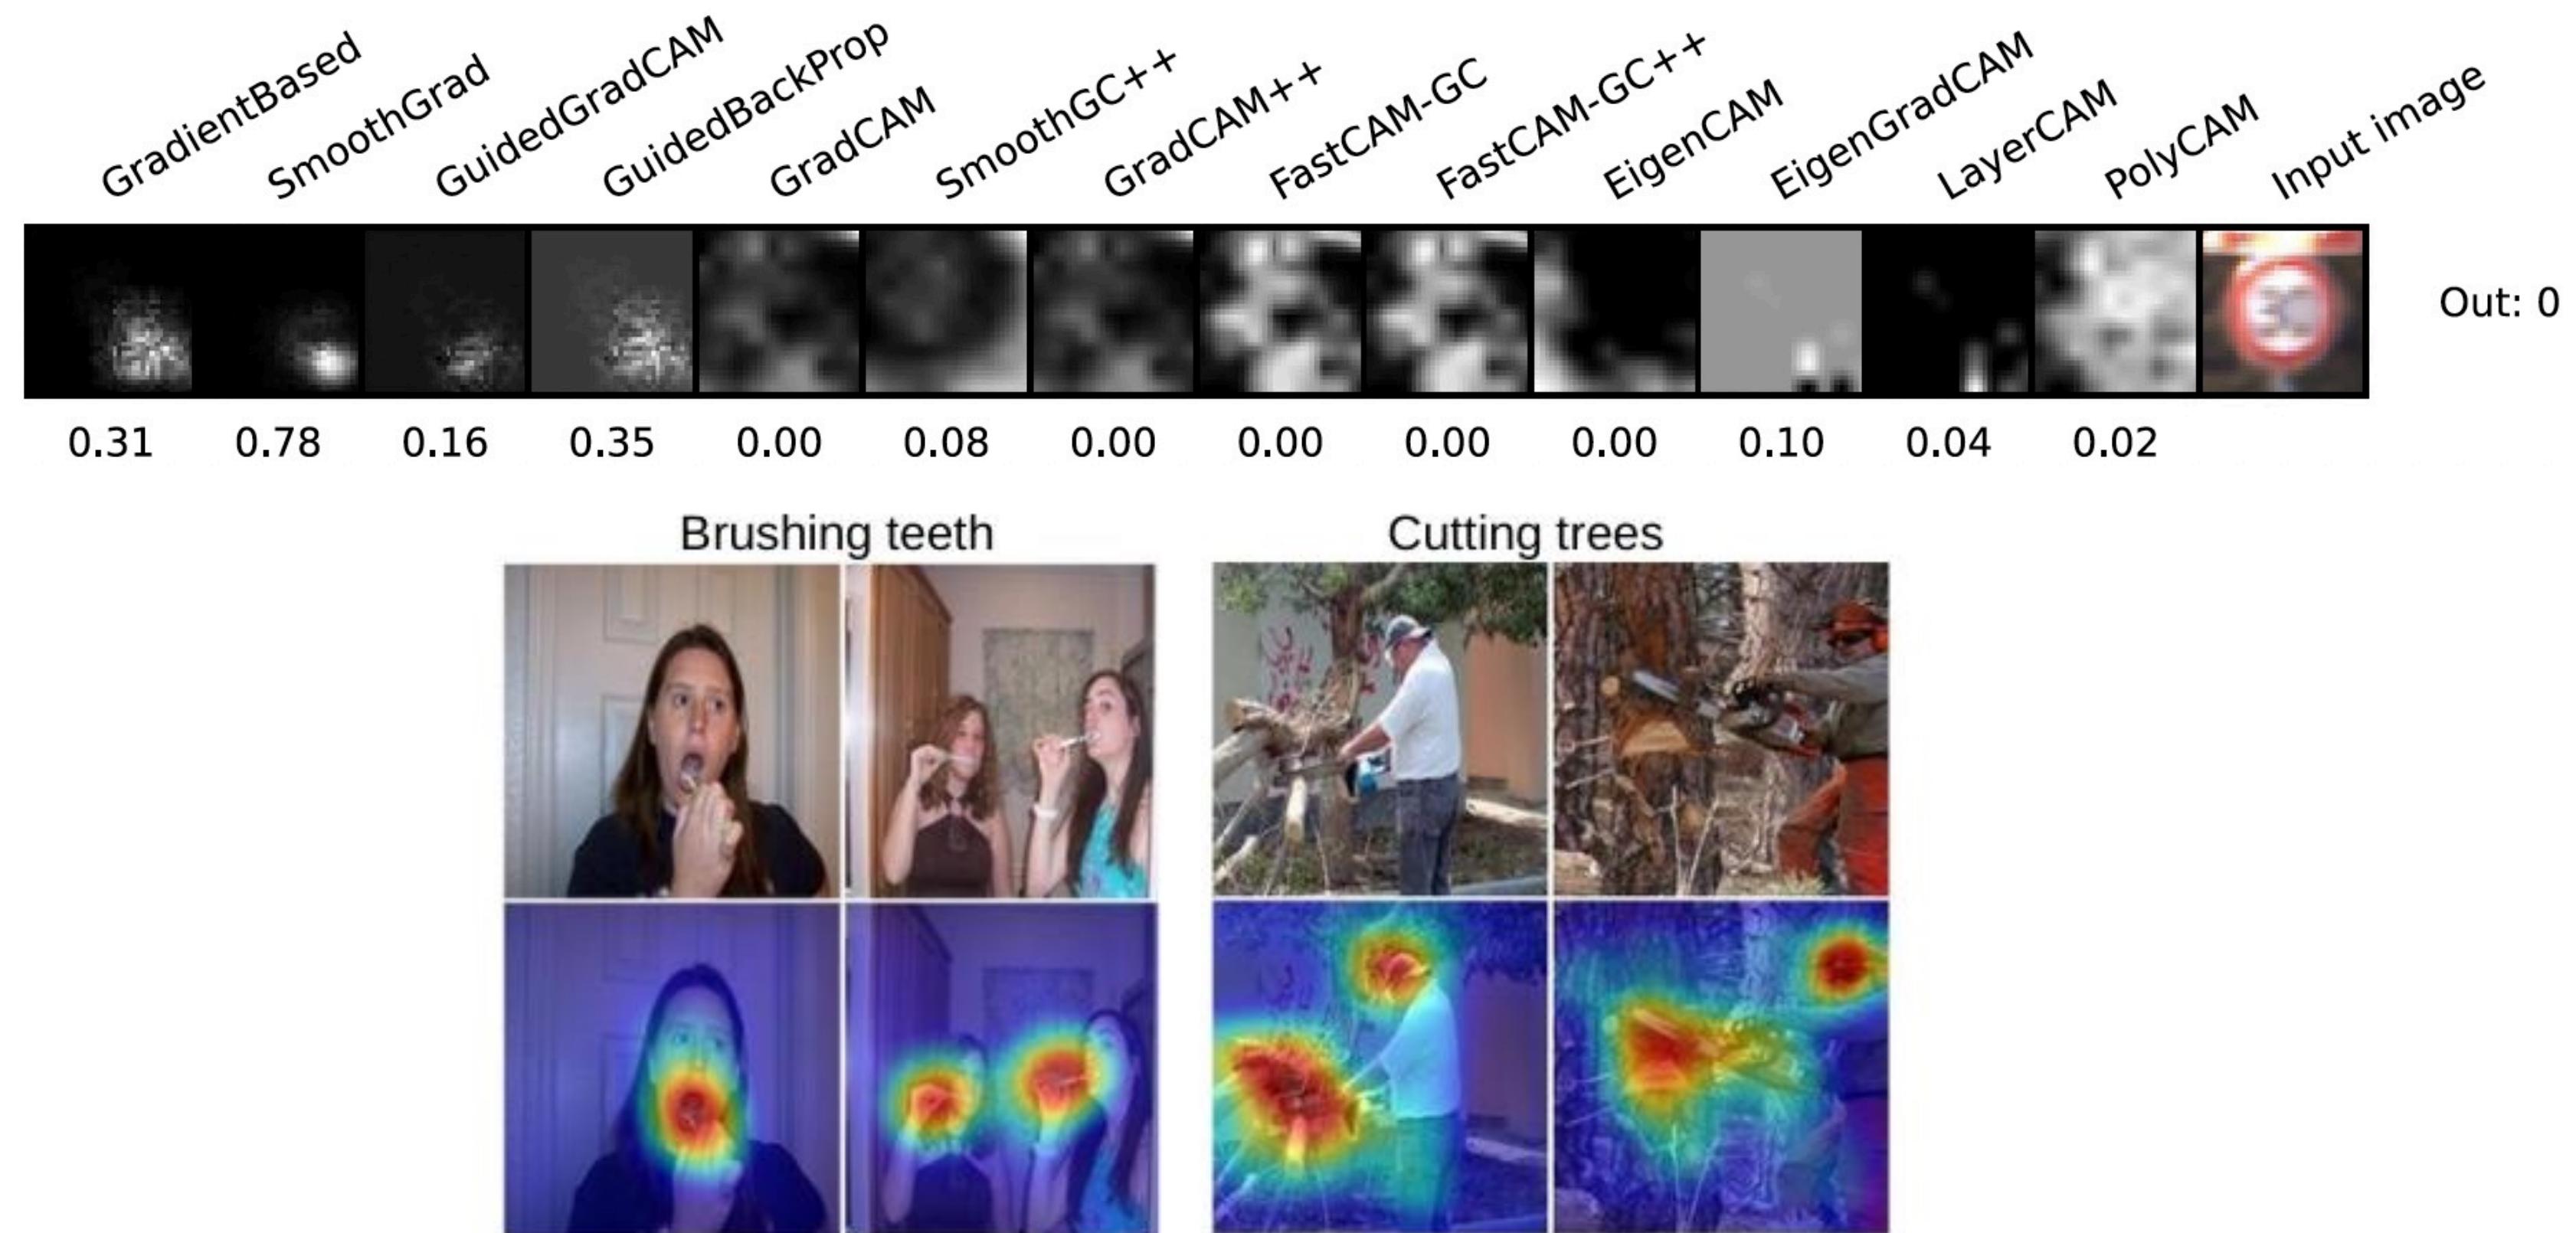
- For a model (AlexNet, VGG, etc), image I , and a class (cell or not cell), which pixels were more relevant to classify?
- If your network output is $S(I)$, you can approximate as:

$$S(I) \approx w^T I + b \qquad w = \frac{\partial S}{\partial I}$$

- So, we can quantify pixel relevance as how much gradient in back propagation it receives.

CNN visualization: saliency map

- Multiple algorithms compute saliency maps: gradient based, deconvolution, class activation, grad-cam.



Exercises



- **Try** the imageProcessing_3_CNN.ipynb notebook
- Add a layer to the CNN, and visualize activations/kernels